

# FOUNDATION-0.2a: Core Database Schema - PART 1 (Code)

## 🎯 CONTEXT

**Phase:** FOUNDATION (Week 1 - Day 1 Morning)  
**Component:** PostgreSQL Core Schema + Alembic Migration  
**Estimated Time:** 30 min AI execution + 10 min verification  
**Complexity:** MEDIUM  
**Risk Level:** LOW (foundational work, well-tested patterns)  
**Files:** Part 1 of 2 (Code implementation)

---

## 📦 DEPENDENCIES

### Must Complete First:

- **PILOT-01:** Bootstrap project structure  COMPLETED
- **PILOT-02:** Orchestrator skeleton  COMPLETED
- **PILOT-03:** Cost Agent skeleton  COMPLETED
- **PILOT-04:** LangGraph integration  COMPLETED
- **PILOT-05:** Spot migration workflow  COMPLETED

### Required Services Running:



```

# Verify infrastructure
cd ~/optiinfra
make verify

# Expected output:
# PostgreSQL... ✓ HEALTHY
# ClickHouse... ✓ HEALTHY
# Qdrant... ✓ HEALTHY
# Redis... ✓ HEALTHY

# Verify PostgreSQL connection
psql postgresql://optiinfra:password@localhost:5432/optiinfra -c "SELECT version();"

# Expected: PostgreSQL version output

```

## Required Environment:



bash

```

# Verify Python environment
cd ~/optiinfra
python --version # Python 3.11+

# Check if shared directory exists
ls shared/database/
# Should exist from PILOT-01 bootstrap

```

## 🎯 OBJECTIVE

Create a **complete PostgreSQL database schema** that provides the foundation for all OptiInfra agents. This schema will store:

1. **Customers** - Customer accounts and metadata
2. **Agents** - Registered agents (Cost, Performance, Resource, Application)
3. **Events** - System events and audit logs
4. **Recommendations** - Optimization recommendations from agents
5. **Approvals** - Customer approvals for recommendations
6. **Optimizations** - Executed optimizations and results

## Success Criteria:

- Complete SQLAlchemy models created (6 tables)
- Alembic migration working (001\_core\_schema.py)
- Database tables created successfully
- Seed data inserted (test customers, agents) - **SEE PART 2**
- All relationships defined (foreign keys)
- Indexes created (for performance)
- Validation tests pass (15+ tests) - **SEE PART 2**
- Can query all tables successfully
- Documentation complete

## Failure Signs:

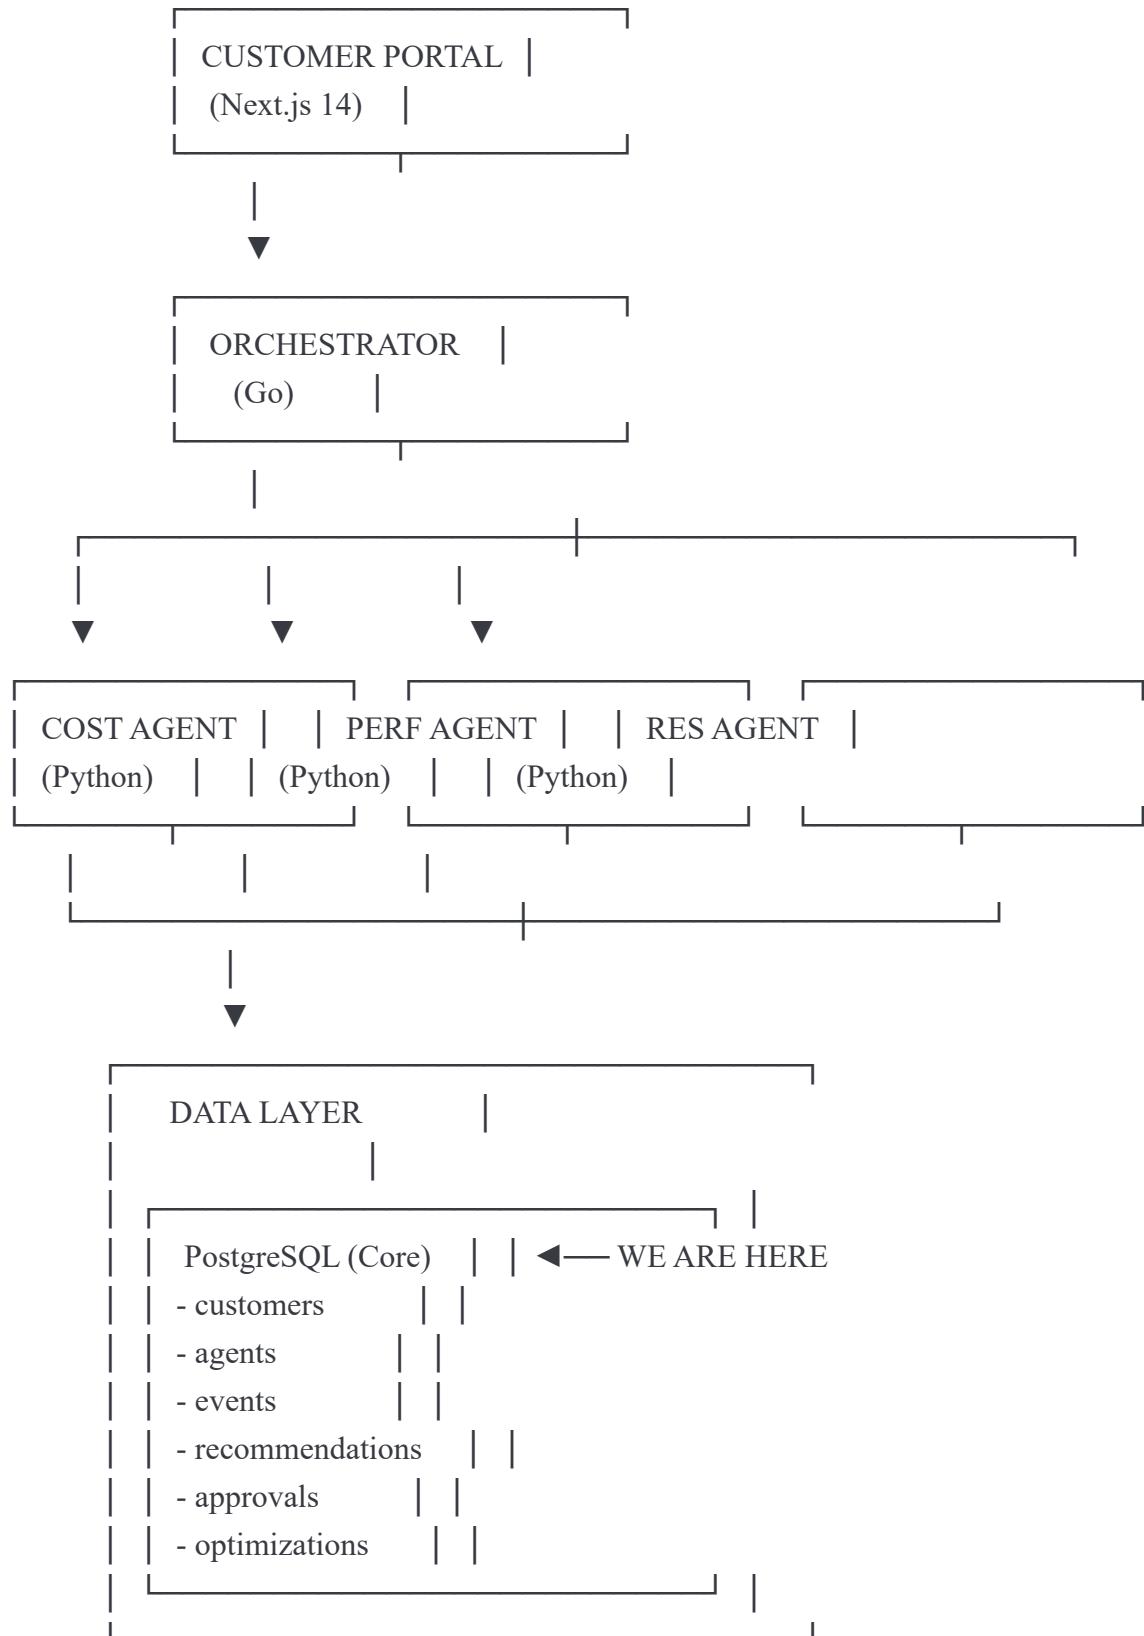
- Migration fails to run
- Tables not created
- Foreign key constraints broken
- Cannot insert test data
- Queries fail

---

## TECHNICAL SPECIFICATION

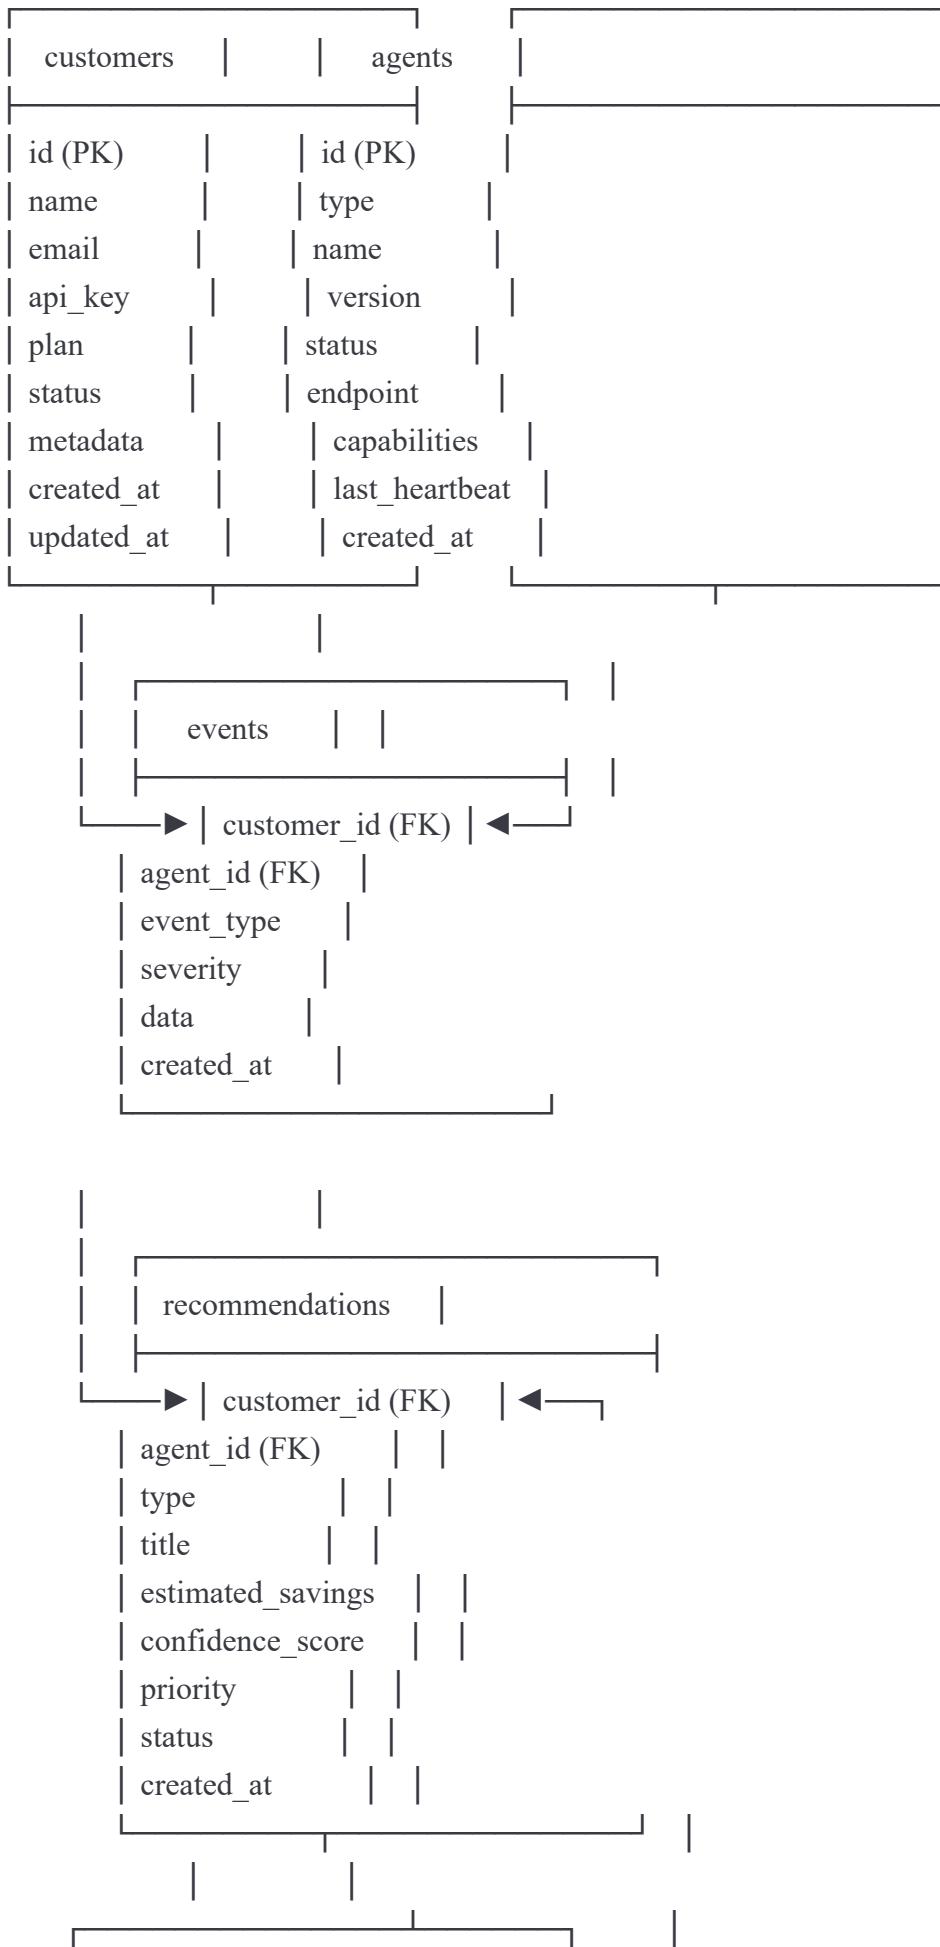
### Architecture Context

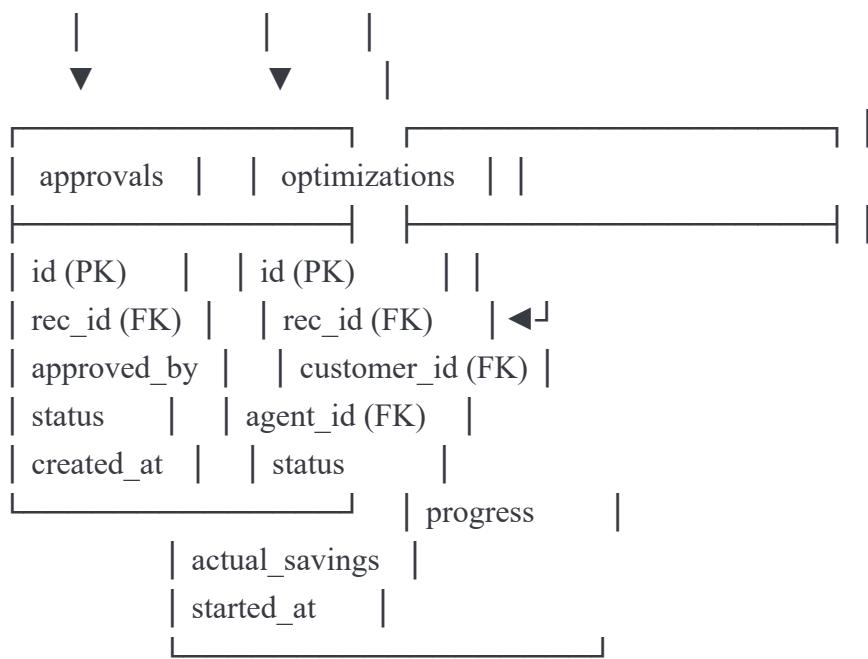




## Database Schema Design







## File Structure to Create (Part 1)



```

shared/
└── database/
    └── models/
        ├── __init__.py      # MODIFY: Export core models ★
        └── core.py          # CREATE: SQLAlchemy models ★
    └── migrations/
        └── versions/
            └── 001_core_schema.py # CREATE: Alembic migration ★

```

**Note:** Seeds, tests, and validation are in **PART 2**

## IMPLEMENTATION - PART 1

### FILE 1: CREATE shared/database/models/core.py

This is the most important file - contains ALL SQLAlchemy models.

Copy this ENTIRE code to Windsurf:



python

"""

Core database models for OptiInfra.

Provides the foundation for all agents.

"""

```
import uuid
from datetime import datetime
from typing import Optional
from enum import Enum as PyEnum

from sqlalchemy import (
    Column, String, Integer, Float, Boolean, DateTime, Text,
    ForeignKey, Enum, Index, UniqueConstraint, CheckConstraint,
    DECIMAL, JSON
)
from sqlalchemy.dialects.postgresql import UUID, JSONB
from sqlalchemy.orm import relationship, declarative_base
from sqlalchemy.sql import func
```

# Base class for all models

```
Base = declarative_base()
```

# =====

# ENUMS

# =====

```
class CustomerPlan(str, PyEnum):
    """Customer subscription plans"""
    FREE = "free"
    STARTUP = "startup"
    ENTERPRISE = "enterprise"
```

```
class CustomerStatus(str, PyEnum):
    """Customer account status"""
    ACTIVE = "active"
    SUSPENDED = "suspended"
    CHURNED = "churned"
```

```
class AgentType(str, PyEnum):
```

```
"""Types of agents in the system"""
ORCHESTRATOR = "orchestrator"
COST = "cost"
PERFORMANCE = "performance"
RESOURCE = "resource"
APPLICATION = "application"
```

```
class AgentStatus(str, PyEnum):
    """Agent health status"""
    STARTING = "starting"
    HEALTHY = "healthy"
    DEGRADED = "degraded"
    FAILED = "failed"
    STOPPED = "stopped"
```

```
class EventSeverity(str, PyEnum):
    """Event severity levels"""
    INFO = "info"
    WARNING = "warning"
    ERROR = "error"
    CRITICAL = "critical"
```

```
class RecommendationPriority(str, PyEnum):
    """Priority of recommendations"""
    LOW = "low"
    MEDIUM = "medium"
    HIGH = "high"
    CRITICAL = "critical"
```

```
class RecommendationStatus(str, PyEnum):
    """Status of recommendations"""
    PENDING = "pending"
    APPROVED = "approved"
    REJECTED = "rejected"
    EXECUTING = "executing"
    COMPLETED = "completed"
    FAILED = "failed"
```

```
ROLLED_BACK = "rolled_back"
```

```
class ApprovalStatus(str, PyEnum):
```

```
    """Approval status"""
    PENDING = "pending"
    APPROVED = "approved"
    REJECTED = "rejected"
```

```
class OptimizationStatus(str, PyEnum):
```

```
    """Optimization execution status"""
    QUEUED = "queued"
    EXECUTING = "executing"
    COMPLETED = "completed"
    FAILED = "failed"
    ROLLED_BACK = "rolled_back"
```

```
# =====
```

```
# MODELS
```

```
# =====
```

```
class Customer(Base):
```

```
    """
    Customer accounts.
```

Represents companies/users using OptiInfra.

Each customer has multiple agents optimizing their infrastructure.

```
"""

 = "customers"
```

```
id = Column(
```

```
    UUID(as_uuid=True),
    primary_key=True,
    default=uuid.uuid4,
    comment="Unique customer identifier"
```

```
)
```

```
name = Column(
```

```
    String(255),
    nullable=False,
```

```
comment="Customer/company name"
)
email = Column(
    String(255),
    nullable=False,
    unique=True,
    index=True,
    comment="Primary contact email"
)
api_key = Column(
    String(64),
    nullable=False,
    unique=True,
    index=True,
    default=lambda: str(uuid.uuid4()).replace("-", ""),
    comment="API authentication key"
)
plan = Column(
    Enum(CustomerPlan),
    nullable=False,
    default=CustomerPlan.FREE,
    comment="Subscription plan"
)
status = Column(
    Enum(CustomerStatus),
    nullable=False,
    default=CustomerStatus.ACTIVE,
    index=True,
    comment="Account status"
)
metadata = Column(
    JSONB,
    nullable=True,
    default={},
    comment="Flexible metadata storage"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="Account creation timestamp"
)
```

```

)
updated_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    onupdate=func.now(),
    comment="Last update timestamp"
)

# Relationships
events = relationship("Event", back_populates="customer", cascade="all, delete-orphan")
recommendations = relationship("Recommendation", back_populates="customer", cascade="all, delete-orphan")
optimizations = relationship("Optimization", back_populates="customer", cascade="all, delete-orphan")

# Indexes


```

**class Agent(Base):**

""""

Registered agents in the system.

Tracks all agents (Cost, Performance, Resource, Application, Orchestrator) with their capabilities, status, and health.

""""

tablename = "agents"

id = Column(

UUID(as\_uuid=True),
 primary\_key=True,
 default=uuid.uuid4,
 comment="Unique agent identifier"

)

type = Column(

```
Enum(AgentType),  
nullable=False,  
index=True,  
comment="Agent type (cost, performance, etc.)"  
)  
name = Column(  
    String(255),  
    nullable=False,  
    comment="Agent name/hostname"  
)  
version = Column(  
    String(50),  
    nullable=True,  
    comment="Agent version (semver)"  
)  
status = Column(  
    Enum(AgentStatus),  
    nullable=False,  
    default=AgentStatus.STARTING,  
    index=True,  
    comment="Current health status"  
)  
endpoint = Column(  
    String(255),  
    nullable=True,  
    comment="Agent HTTP endpoint URL"  
)  
capabilities = Column(  
    JSONB,  
    nullable=False,  
    default=[],  
    comment="List of agent capabilities"  
)  
metadata = Column(  
    JSONB,  
    nullable=True,  
    default={},  
    comment="Additional agent metadata"  
)  
last_heartbeat = Column(  
    DateTime(timezone=True),
```

```

nullable=True,
comment="Last heartbeat timestamp"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="Agent registration timestamp"
)
updated_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    onupdate=func.now(),
    comment="Last update timestamp"
)

# Relationships
events = relationship("Event", back_populates="agent", cascade="all, delete-orphan")
recommendations = relationship("Recommendation", back_populates="agent", cascade="all, delete-orphan")
optimizations = relationship("Optimization", back_populates="agent", cascade="all, delete-orphan")

# Indexes
__table_args__ = (
    Index("idx_agent_type", "type"),
    Index("idx_agent_status", "status"),
    Index("idx_agent_heartbeat", "last_heartbeat"),
    UniqueConstraint("type", "name", name="uq_agent_type_name"),
)
def __repr__(self):
    return f"<Agent(id={self.id}, type={self.type.value}, name='{self.name}', status={self.status.value})>"

class Event(Base):
    """
    System events and audit logs.

    Tracks all significant events in the system for auditing,
    debugging, and analytics.
    """

```

```
__tablename__ = "events"

id = Column(
    UUID(as_uuid=True),
    primary_key=True,
    default=uuid.uuid4,
    comment="Unique event identifier"
)
customer_id = Column(
    UUID(as_uuid=True),
    ForeignKey("customers.id", ondelete="CASCADE"),
    nullable=False,
    index=True,
    comment="Associated customer"
)
agent_id = Column(
    UUID(as_uuid=True),
    ForeignKey("agents.id", ondelete="SET NULL"),
    nullable=True,
    index=True,
    comment="Agent that generated event (optional)"
)
event_type = Column(
    String(100),
    nullable=False,
    index=True,
    comment="Event type (e.g., optimization_started)"
)
severity = Column(
    Enum(EventSeverity),
    nullable=False,
    default=EventSeverity.INFO,
    index=True,
    comment="Event severity level"
)
data = Column(
    JSONB,
    nullable=False,
    default={},
    comment="Event payload/details"
)
```

```

created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    index=True,
    comment="Event timestamp"
)

# Relationships
customer = relationship("Customer", back_populates="events")
agent = relationship("Agent", back_populates="events")

# Indexes


```

**class Recommendation(Base):**

""""

Optimization recommendations from agents.

Agents analyze infrastructure and generate recommendations  
for cost savings, performance improvements, etc.

""""

tablename\_ = "recommendations"

```

id = Column(
    UUID(as_uuid=True),
    primary_key=True,
    default=uuid.uuid4,
    comment="Unique recommendation identifier"
)
```

```
customer_id = Column(  
    UUID(as_uuid=True),  
    ForeignKey("customers.id", ondelete="CASCADE"),  
    nullable=False,  
    index=True,  
    comment="Associated customer"  
)  
  
agent_id = Column(  
    UUID(as_uuid=True),  
    ForeignKey("agents.id", ondelete="CASCADE"),  
    nullable=False,  
    index=True,  
    comment="Agent that created recommendation"  
)  
  
type = Column(  
    String(100),  
    nullable=False,  
    index=True,  
    comment="Recommendation type (spot_migration, etc.)"  
)  
  
title = Column(  
    String(255),  
    nullable=False,  
    comment="Recommendation title"  
)  
  
description = Column(  
    Text,  
    nullable=True,  
    comment="Detailed description"  
)  
  
estimated_savings = Column(  
    DECIMAL(10, 2),  
    nullable=True,  
    comment="Estimated monthly savings (USD)"  
)  
  
estimated_improvement = Column(  
    DECIMAL(10, 2),  
    nullable=True,  
    comment="Performance improvement multiplier"  
)  
  
confidence_score = Column(  
    Integer,  
    nullable=True,  
    comment="Confidence score (0-100)"  
)
```

```
        Float,  
        nullable=False,  
        default=0.0,  
        comment="Confidence score (0.0 to 1.0)"  
)  
priority = Column(  
    Enum(RecommendationPriority),  
    nullable=False,  
    default=RecommendationPriority.MEDIUM,  
    index=True,  
    comment="Recommendation priority"  
)  
status = Column(  
    Enum(RecommendationStatus),  
    nullable=False,  
    default=RecommendationStatus.PENDING,  
    index=True,  
    comment="Current status"  
)  
data = Column(  
    JSONB,  
    nullable=False,  
    default={},  
    comment="Recommendation details/metadata"  
)  
created_at = Column(  
    DateTime(timezone=True),  
    nullable=False,  
    server_default=func.now(),  
    index=True,  
    comment="Creation timestamp"  
)  
updated_at = Column(  
    DateTime(timezone=True),  
    nullable=False,  
    server_default=func.now(),  
    onupdate=func.now(),  
    comment="Last update timestamp"  
)  
approved_at = Column(  
    DateTime(timezone=True),
```

```

nullable=True,
comment="Approval timestamp"
)
executed_at = Column(
    DateTime(timezone=True),
    nullable=True,
    comment="Execution start timestamp"
)

# Relationships
customer = relationship("Customer", back_populates="recommendations")
agent = relationship("Agent", back_populates="recommendations")
approvals = relationship("Approval", back_populates="recommendation", cascade="all, delete-orphan")
optimizations = relationship("Optimization", back_populates="recommendation", cascade="all, delete-orphan")

# Indexes and Constraints


```

```

class Approval(Base):
    """
    Customer approvals for recommendations.

```

Tracks which recommendations customers have approved/rejected.

```

    """
tablename_ = "approvals"

id = Column(
    UUID(as_uuid=True),

```

```
primary_key=True,
default=uuid.uuid4,
comment="Unique approval identifier"
)
recommendation_id = Column(
    UUID(as_uuid=True),
    ForeignKey("recommendations.id", ondelete="CASCADE"),
    nullable=False,
    index=True,
    comment="Associated recommendation"
)
approved_by = Column(
    String(255),
    nullable=False,
    comment="User who approved/rejected (email or ID)"
)
status = Column(
    Enum(ApprovalStatus),
    nullable=False,
    default=ApprovalStatus.PENDING,
    index=True,
    comment="Approval status"
)
comment = Column(
    Text,
    nullable=True,
    comment="Optional approval comment"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="Approval request timestamp"
)
updated_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    onupdate=func.now(),
    comment="Last update timestamp"
)
```

```
# Relationships
recommendation = relationship("Recommendation", back_populates="approvals")

# Indexes
__table_args__ = (
    Index("idx_approval_rec", "recommendation_id"),
    Index("idx_approval_status", "status"),
)
def __repr__(self):
    return f"<Approval(id={self.id}, rec_id={self.recommendation_id}, status={self.status.value})>"
```

## class Optimization(Base):

"""

Executed optimizations and their results.

Tracks the execution of approved recommendations,  
including progress, results, and actual impact.

"""

\_\_tablename\_\_ = "optimizations"

id = Column(

UUID(as\_uuid=True),

primary\_key=True,

default=uuid.uuid4,

comment="Unique optimization identifier"

)

recommendation\_id = Column(

UUID(as\_uuid=True),

ForeignKey("recommendations.id", ondelete="CASCADE"),

nullable=False,

index=True,

comment="Associated recommendation"

)

customer\_id = Column(

UUID(as\_uuid=True),

ForeignKey("customers.id", ondelete="CASCADE"),

nullable=False,

index=True,

```
comment="Associated customer"
)
agent_id = Column(
    UUID(as_uuid=True),
    ForeignKey("agents.id", ondelete="CASCADE"),
    nullable=False,
    index=True,
    comment="Agent executing optimization"
)
status = Column(
    Enum(OptimizationStatus),
    nullable=False,
    default=OptimizationStatus.QUEUED,
    index=True,
    comment="Execution status"
)
progress = Column(
    Integer,
    nullable=False,
    default=0,
    comment="Progress percentage (0-100)"
)
result = Column(
    JSONB,
    nullable=True,
    default={},
    comment="Execution results/details"
)
error = Column(
    Text,
    nullable=True,
    comment="Error message if failed"
)
actual_savings = Column(
    DECIMAL(10, 2),
    nullable=True,
    comment="Actual monthly savings achieved (USD)"
)
actual_improvement = Column(
    DECIMAL(10, 2),
    nullable=True,
```

```

comment="Actual performance improvement"
)
started_at = Column(
    DateTime(timezone=True),
    nullable=True,
    comment="Execution start timestamp"
)
completed_at = Column(
    DateTime(timezone=True),
    nullable=True,
    comment="Execution completion timestamp"
)
rolled_back_at = Column(
    DateTime(timezone=True),
    nullable=True,
    comment="Rollback timestamp (if applicable)"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="Creation timestamp"
)

```

```

# Relationships
recommendation = relationship("Recommendation", back_populates="optimizations")
customer = relationship("Customer", back_populates="optimizations")
agent = relationship("Agent", back_populates="optimizations")

```

```

# Indexes and Constraints
__table_args__ =
Index("idx_opt_rec", "recommendation_id"),
Index("idx_opt_customer", "customer_id"),
Index("idx_opt_agent", "agent_id"),
Index("idx_opt_status", "status"),
Index("idx_opt_customer_status", "customer_id", "status"),
CheckConstraint("progress >= 0 AND progress <= 100", name="check_progress_range"),
)

```

```
def __repr__(self):  
    return f"<Optimization(id={self.id}, status={self.status.value}, progress={self.progress}%)>"
```

 **FILE 1 COMPLETE (~400 lines, ~14,000 characters)**

---

## FILE 2: MODIFY shared/database/models/init.py

Update the package exports to include all core models:



python

""""

Database models package.

""""

```
from shared.database.models.core import (
```

Base,

# Enums

CustomerPlan,

CustomerStatus,

AgentType,

AgentStatus,

EventSeverity,

RecommendationPriority,

RecommendationStatus,

ApprovalStatus,

OptimizationStatus,

# Models

Customer,

Agent,

Event,

Recommendation,

Approval,

Optimization,

)

\_\_all\_\_ = [

"Base",

# Enums

"CustomerPlan",

"CustomerStatus",

"AgentType",

"AgentStatus",

"EventSeverity",

"RecommendationPriority",

"RecommendationStatus",

"ApprovalStatus",

"OptimizationStatus",

# Models

"Customer",

"Agent",

"Event",

"Recommendation",

```
"Approval",  
"Optimization",  
]  
]
```

 **FILE 2 COMPLETE**

---

### **FILE 3: CREATE shared/database/migrations/versions/001\_core\_schema.py**

**Alembic migration to create all tables in PostgreSQL:**



:::::

## Create core schema tables

Revision ID: 001\_core\_schema

Revises:

Create Date: 2025-10-18 12:00:00.000000

:::::

```
from alembic import op
import sqlalchemy as sa
from sqlalchemy.dialects import postgresql
```

*# revision identifiers, used by Alembic.*

revision = '001\_core\_schema'

down\_revision = None

branch\_labels = None

depends\_on = None

def upgrade():

"""Create core schema tables"""

# Create ENUM types
op.execute("""
CREATE TYPE customerplan AS ENUM ('free', 'startup', 'enterprise');
CREATE TYPE customerstatus AS ENUM ('active', 'suspended', 'churned');
CREATE TYPE agenttype AS ENUM ('orchestrator', 'cost', 'performance', 'resource', 'application');
CREATE TYPE agentstatus AS ENUM ('starting', 'healthy', 'degraded', 'failed', 'stopped');
CREATE TYPE eventseverity AS ENUM ('info', 'warning', 'error', 'critical');
CREATE TYPE recommendationpriority AS ENUM ('low', 'medium', 'high', 'critical');
CREATE TYPE recommendationstatus AS ENUM ('pending', 'approved', 'rejected', 'executing', 'completed', 'failed', 'none');
CREATE TYPE approvalstatus AS ENUM ('pending', 'approved', 'rejected');
CREATE TYPE optimizationstatus AS ENUM ('queued', 'executing', 'completed', 'failed', 'rolled\_back');
""")

# Create customers table

op.create\_table(

```
'customers',
sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
sa.Column('name', sa.String(255), nullable=False),
sa.Column('email', sa.String(255), nullable=False, unique=True),
sa.Column('api_key', sa.String(64), nullable=False, unique=True),
```

```
sa.Column('plan', postgresql.ENUM(name='customerplan'), nullable=False, server_default='free'),  
sa.Column('status', postgresql.ENUM(name='customerstatus'), nullable=False, server_default='active'),  
sa.Column('metadata', postgresql.JSONB, nullable=True, server_default='{}')
```