

# FOUNDATION-0.2b: Agent State Tables - PART 1 (Code)

## 🎯 CONTEXT

**Phase:** FOUNDATION (Week 1 - Day 1 Afternoon)  
**Component:** Agent State & Configuration Tables  
**Estimated Time:** 20 min AI execution + 10 min verification  
**Complexity:** MEDIUM  
**Risk Level:** LOW  
**Files:** Part 1 of 2 (Code implementation)

---

## 📦 DEPENDENCIES

### Must Complete First:

- **FOUNDATION-0.2a:** Core Database Schema  COMPLETED & VALIDATED
  - All 6 core tables created
  - 13/13 tests passing
  - Seed data loaded

### Required Services Running:



```
# Verify all services are healthy
cd ~/optiinfra
make verify
```

```
# Expected output:
# PostgreSQL...  HEALTHY
# ClickHouse...  HEALTHY
# Qdrant...  HEALTHY
# Redis...  HEALTHY
```

### Database State:



```
# Verify core tables exist
psql postgresql://optiinfra:password@localhost:5432/optiinfra -c "\dt"
```

# Expected: customers, agents, events, recommendations, approvals, optimizations

---

## 🎯 OBJECTIVE

Extend the database schema with **agent-specific state and configuration tables**. These tables enable:

1. **Agent Configuration** - Store agent settings, parameters, thresholds
2. **Agent State** - Track real-time agent state (active workflows, locks, status)
3. **Agent Capabilities** - Detailed capability definitions and versions
4. **Agent Metrics** - Agent-level performance metrics

### What This Enables:

- Store agent configuration (KV cache settings, thresholds, etc.)
- Track agent state (workflows in progress, locks)
- Version capabilities (each agent can have multiple capability versions)
- Monitor agent health metrics
- Enable agent coordination (orchestrator needs this!)

### Success Criteria:

- ✓ 4 new tables created (agent\_configs, agent\_states, agent\_capabilities, agent\_metrics)
- ✓ Alembic migration working (002\_agent\_state\_tables.py)
- ✓ All foreign keys to agents table working
- ✓ Seed data for test agents
- ✓ 10+ tests passing
- ✓ Documentation complete

### Failure Signs:

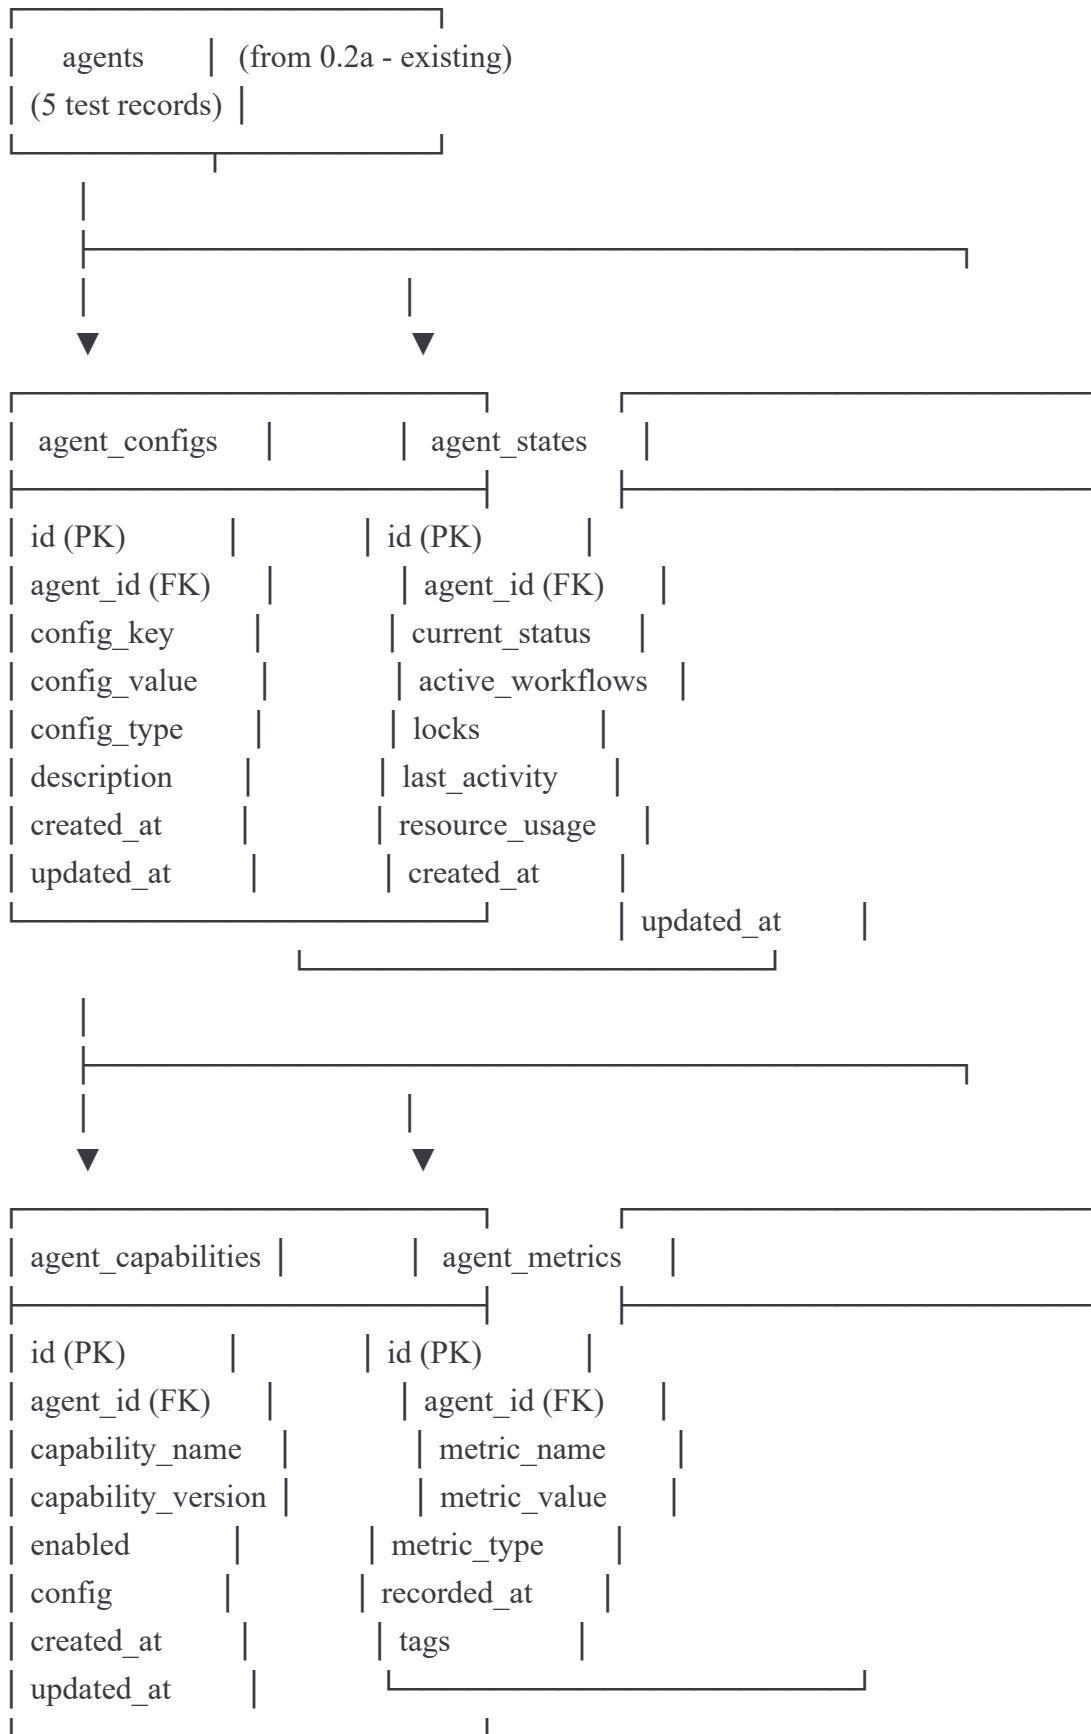
- ✗ Migration fails
- ✗ Foreign key constraints broken
- ✗ Cannot insert test data

---

## 📘 TECHNICAL SPECIFICATION

### Database Schema Design





## File Structure to Create (Part 1)



```
shared/
└── database/
    ├── models/
    │   ├── __init__.py      # MODIFY: Export new models ★
    │   ├── core.py          # KEEP: From 0.2a
    │   └── agent_state.py   # CREATE: Agent state models ★
    └── migrations/
        └── versions/
            ├── 001_core_schema.py    # KEEP: From 0.2a
            └── 002_agent_state_tables.py # CREATE: New migration ★
```



## IMPLEMENTATION - PART 1

### FILE 1: CREATE shared/database/models/agent\_state.py

Complete SQLAlchemy models for agent state management:



python

"""

Agent state and configuration models.

Extends the core schema with agent-specific data.

"""

```
import uuid
from datetime import datetime
from enum import Enum as PyEnum
```

```
from sqlalchemy import (
    Column, String, Integer, Float, Boolean, DateTime, Text,
    ForeignKey, Enum, Index, UniqueConstraint
)
from sqlalchemy.dialects.postgresql import UUID, JSONB
from sqlalchemy.orm import relationship
from sqlalchemy.sql import func

from shared.database.models.core import Base
```

```
# =====
```

```
# ENUMS
```

```
# =====
```

```
class ConfigType(str, PyEnum):
    """Configuration value types"""
    STRING = "string"
    INTEGER = "integer"
    FLOAT = "float"
    BOOLEAN = "boolean"
    JSON = "json"
```

```
class AgentStatusDetail(str, PyEnum):
    """Detailed agent status"""
    IDLE = "idle"
    BUSY = "busy"
    PROCESSING = "processing"
    WAITING = "waiting"
    ERROR = "error"
```

```
class MetricType(str, PyEnum):
    """Metric value types"""
    COUNTER = "counter"
    GAUGE = "gauge"
    HISTOGRAM = "histogram"

# =====
# MODELS
# =====

class AgentConfig(Base):
    """
    Agent configuration storage.

    Stores key-value configuration for each agent.
    Examples: thresholds, timeouts, feature flags, optimization parameters.
    """

    __tablename__ = "agent_configs"

    id = Column(
        UUID(as_uuid=True),
        primary_key=True,
        default=uuid.uuid4,
        comment="Unique config entry identifier"
    )
    agent_id = Column(
        UUID(as_uuid=True),
        ForeignKey("agents.id", ondelete="CASCADE"),
        nullable=False,
        index=True,
        comment="Agent this config belongs to"
    )
    config_key = Column(
        String(255),
        nullable=False,
        index=True,
        comment="Configuration key (e.g., 'kv_cache_size')"
    )
    config_value = Column(
        Text,
```

```

nullable=False,
comment="Configuration value (stored as string)"
)
config_type = Column(
    Enum(ConfigType, create_type=False, values_callable=lambda x: [e.value for e in x]),
    nullable=False,
    default=ConfigType.STRING,
    comment="Type of configuration value"
)
description = Column(
    Text,
    nullable=True,
    comment="Human-readable description of this config"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="Config creation timestamp"
)
updated_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    onupdate=func.now(),
    comment="Last update timestamp"
)
)

# Relationships
agent = relationship("Agent", back_populates="configs")

# Indexes and Constraints
__table_args__ =
    Index("idx_agent_config_agent", "agent_id"),
    Index("idx_agent_config_key", "config_key"),
    Index("idx_agent_config_agent_key", "agent_id", "config_key"),
    UniqueConstraint("agent_id", "config_key", name="uq_agent_config"),
)

def __repr__(self):
    return f"<AgentConfig(agent_id={self.agent_id}, key='{self.config_key}', type={self.config_type.value})>"
```

```
class AgentState(Base):
    """
    Real-time agent state.

    Tracks the current operational state of each agent:
    - Active workflows
    - Resource locks
    - Current status
    - Resource usage
    """

    __tablename__ = "agent_states"

    id = Column(
        UUID(as_uuid=True),
        primary_key=True,
        default=uuid.uuid4,
        comment="Unique state identifier"
    )
    agent_id = Column(
        UUID(as_uuid=True),
        ForeignKey("agents.id", ondelete="CASCADE"),
        nullable=False,
        unique=True,
        index=True,
        comment="Agent this state belongs to (one-to-one)"
    )
    current_status = Column(
        Enum(AgentStatusDetail, create_type=False, values_callable=lambda x: [e.value for e in x]),
        nullable=False,
        default=AgentStatusDetail.IDLE,
        index=True,
        comment="Current detailed status"
    )
    active_workflows = Column(
        JSONB,
        nullable=False,
        default=[],
        comment="List of currently active workflow IDs"
    )
```

```
locks = Column(
    JSONB,
    nullable=False,
    default={},
    comment="Resource locks held by this agent"
)
last_activity = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    index=True,
    comment="Last activity timestamp"
)
resource_usage = Column(
    JSONB,
    nullable=True,
    default={},
    comment="Current resource usage (CPU, memory, etc.)"
)
metadata = Column(
    JSONB,
    nullable=True,
    default={},
    comment="Additional state metadata"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="State creation timestamp"
)
updated_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    onupdate=func.now(),
    comment="Last update timestamp"
)
# Relationships
agent = relationship("Agent", back_populates="state", uselist=False)
```

```

#Indexes
__table_args__ = (
    Index("idx_agent_state_agent", "agent_id"),
    Index("idx_agent_state_status", "current_status"),
    Index("idx_agent_state_activity", "last_activity"),
)
def __repr__(self):
    return f"<AgentState(agent_id={self.agent_id}, status={self.current_status.value})>"
```

## class AgentCapability(Base):

""""

Agent capability definitions.

Tracks what each agent can do, with version information.

Example: Cost Agent can do "spot\_migration" v1.2.0

""""

\_\_tablename\_\_ = "agent\_capabilities"

id = Column(

UUID(as\_uuid=True),

primary\_key=True,

default=uuid.uuid4,

comment="Unique capability identifier"

)

agent\_id = Column(

UUID(as\_uuid=True),

ForeignKey("agents.id", ondelete="CASCADE"),

nullable=False,

index=True,

comment="Agent this capability belongs to"

)

capability\_name = Column(

String(255),

nullable=False,

index=True,

comment="Capability name (e.g., 'spot\_migration')"

)

capability\_version = Column(

```
String(50),
nullable=False,
default="1.0.0",
comment="Capability version (semver)"

)
enabled = Column(
    Boolean,
    nullable=False,
    default=True,
    index=True,
    comment="Whether this capability is currently enabled"
)
config = Column(
    JSONB,
    nullable=True,
    default={},
    comment="Capability-specific configuration"
)
description = Column(
    Text,
    nullable=True,
    comment="Human-readable description"
)
created_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    comment="Capability registration timestamp"
)
updated_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    onupdate=func.now(),
    comment="Last update timestamp"
)

# Relationships
agent = relationship("Agent", back_populates="capabilities")
```

# Indexes and Constraints

```
__table_args__ = (
    Index("idx_agent_cap_agent", "agent_id"),
    Index("idx_agent_cap_name", "capability_name"),
    Index("idx_agent_cap_enabled", "enabled"),
    Index("idx_agent_cap_agent_name", "agent_id", "capability_name"),
    UniqueConstraint("agent_id", "capability_name", name="uq_agent_capability"),
)
def __repr__(self):
    return f"<AgentCapability(agent_id={self.agent_id}, name='{self.capability_name}', version='{self.capability_version}')>"
```

## class AgentMetric(Base):

```
"""
```

Agent-level performance metrics.

Time-series metrics for agent performance:

- Request counts
- Success rates
- Processing times
- Resource usage

```
"""
```

```
tablename = "agent_metrics"
```

```
id = Column(
```

```
    UUID(as_uuid=True),
```

```
    primary_key=True,
```

```
    default=uuid.uuid4,
```

```
    comment="Unique metric identifier"
```

```
)
```

```
agent_id = Column(
```

```
    UUID(as_uuid=True),
```

```
    ForeignKey("agents.id", ondelete="CASCADE"),
```

```
    nullable=False,
```

```
    index=True,
```

```
    comment="Agent this metric belongs to"
```

```
)
```

```
metric_name = Column(
```

```
    String(255),
```

```
    nullable=False,
```

```
    index=True,
```

```

comment="Metric name (e.g., 'requests_total')"
)
metric_value = Column(
    Float,
    nullable=False,
    comment="Metric value"
)
metric_type = Column(
    Enum(MetricType, create_type=False, values_callable=lambda x: [e.value for e in x]),
    nullable=False,
    default=MetricType.GAUGE,
    comment="Type of metric"
)
tags = Column(
    JSONB,
    nullable=True,
    default={},
    comment="Metric tags/labels"
)
recorded_at = Column(
    DateTime(timezone=True),
    nullable=False,
    server_default=func.now(),
    index=True,
    comment="Metric recording timestamp"
)

```

#### *# Relationships*

```
agent = relationship("Agent", back_populates="metrics")
```

#### *# Indexes*

```

__table_args__ = (
    Index("idx_agent_metric_agent", "agent_id"),
    Index("idx_agent_metric_name", "metric_name"),
    Index("idx_agent_metric_recorded", "recorded_at"),
    Index("idx_agent_metric_agent_name", "agent_id", "metric_name"),
    Index("idx_agent_metric_agent_recorded", "agent_id", "recorded_at"),
)
```

```
def __repr__(self):
    return f"<AgentMetric(agent_id={self.agent_id}, name='{self.metric_name}', value={self.metric_value})>"
```

## FILE 1 COMPLETE (~400 lines)

---

## FILE 2: MODIFY shared/database/models/core.py

### Add relationships to the Agent model:

Find the Agent class and add these new relationship lines:



python

```
class Agent(Base):
    # ... existing code ...

# Relationships (ADD THESE NEW LINES)
events = relationship("Event", back_populates="agent", cascade="all, delete-orphan")
recommendations = relationship("Recommendation", back_populates="agent", cascade="all, delete-orphan")
optimizations = relationship("Optimization", back_populates="agent", cascade="all, delete-orphan")

# NEW RELATIONSHIPS - ADD THESE:
configs = relationship("AgentConfig", back_populates="agent", cascade="all, delete-orphan")
state = relationship("AgentState", back_populates="agent", uselist=False, cascade="all, delete-orphan")
capabilities = relationship("AgentCapability", back_populates="agent", cascade="all, delete-orphan")
metrics = relationship("AgentMetric", back_populates="agent", cascade="all, delete-orphan")
```

## FILE 2 MODIFIED

---

## FILE 3: MODIFY shared/database/models/init.py

### Export the new models:



python

||||

Database models package.

||||

```
from shared.database.models.core import (
    Base,
    # Enums
    CustomerPlan,
    CustomerStatus,
    AgentType,
    AgentStatus,
    EventSeverity,
    RecommendationPriority,
    RecommendationStatus,
    ApprovalStatus,
    OptimizationStatus,
    # Models
    Customer,
    Agent,
    Event,
    Recommendation,
    Approval,
    Optimization,
)
```

# NEW IMPORTS - ADD THESE:

```
from shared.database.models.agent_state import (
    # Enums
    ConfigType,
    AgentStatusDetail,
    MetricType,
    # Models
    AgentConfig,
    AgentState,
    AgentCapability,
    AgentMetric,
)
```

\_\_all\_\_ = [

"Base",

# Core Enums

"CustomerPlan",

```
"CustomerStatus",
"AgentType",
"AgentStatus",
"EventSeverity",
"RecommendationPriority",
"RecommendationStatus",
"ApprovalStatus",
"OptimizationStatus",
# Agent State Enums - NEW
"ConfigType",
"AgentStatusDetail",
"MetricType",
# Core Models
"Customer",
"Agent",
"Event",
"Recommendation",
"Approval",
"Optimization",
# Agent State Models - NEW
"AgentConfig",
"AgentState",
"AgentCapability",
"AgentMetric",
]
```

### FILE 3 MODIFIED

---

## FILE 4: CREATE shared/database/migrations/versions/002\_agent\_state\_tables.py

Alembic migration to create agent state tables:



:::::

## Create agent state tables

Revision ID: 002\_agent\_state\_tables

Revises: 001\_core\_schema

Create Date: 2025-10-18 18:00:00.000000

:::::

```
from alembic import op
import sqlalchemy as sa
from sqlalchemy.dialects import postgresql
```

*# revision identifiers, used by Alembic.*

revision = '002\_agent\_state\_tables'

down\_revision = '001\_core\_schema'

branch\_labels = None

depends\_on = None

**def upgrade():**

"""Create agent state tables"""

# Create ENUM types
op.execute("""
CREATE TYPE configtype AS ENUM ('string', 'integer', 'float', 'boolean', 'json');
CREATE TYPE agentstatusdetail AS ENUM ('idle', 'busy', 'processing', 'waiting', 'error');
CREATE TYPE metrictype AS ENUM ('counter', 'gauge', 'histogram');
""")

# Create agent\_configs table

op.create\_table(

```
'agent_configs',
sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
sa.Column('agent_id', postgresql.UUID(as_uuid=True), sa.ForeignKey('agents.id', ondelete='CASCADE'), nullable=False),
sa.Column('config_key', sa.String(255), nullable=False),
sa.Column('config_value', sa.Text, nullable=False),
sa.Column('config_type', postgresql.ENUM(name='configtype'), nullable=False, server_default='string'),
sa.Column('description', sa.Text, nullable=True),
sa.Column('created_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
sa.Column('updated_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
)
```

```

# Create indexes for agent_configs
op.create_index('idx_agent_config_agent', 'agent_configs', ['agent_id'])
op.create_index('idx_agent_config_key', 'agent_configs', ['config_key'])
op.create_index('idx_agent_config_agent_key', 'agent_configs', ['agent_id', 'config_key'])
op.create_unique_constraint('uq_agent_config', 'agent_configs', ['agent_id', 'config_key'])

# Create agent_states table
op.create_table(
    'agent_states',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), sa.ForeignKey('agents.id', ondelete='CASCADE'), nullable=False),
    sa.Column('current_status', postgresql.ENUM(name='agentstatusdetail'), nullable=False, server_default='idle'),
    sa.Column('active_workflows', postgresql.JSONB, nullable=False, server_default='[]'),
    sa.Column('locks', postgresql.JSONB, nullable=False, server_default='{}'),
    sa.Column('last_activity', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
    sa.Column('resource_usage', postgresql.JSONB, nullable=True, server_default='{}'),
    sa.Column('metadata', postgresql.JSONB, nullable=True, server_default='{}'),
    sa.Column('created_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
    sa.Column('updated_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
)

```

```

# Create indexes for agent_states
op.create_index('idx_agent_state_agent', 'agent_states', ['agent_id'])
op.create_index('idx_agent_state_status', 'agent_states', ['current_status'])
op.create_index('idx_agent_state_activity', 'agent_states', ['last_activity'])

```

```

# Create agent_capabilities table
op.create_table(
    'agent_capabilities',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), sa.ForeignKey('agents.id', ondelete='CASCADE'), nullable=False),
    sa.Column('capability_name', sa.String(255), nullable=False),
    sa.Column('capability_version', sa.String(50), nullable=False, server_default='1.0.0'),
    sa.Column('enabled', sa.Boolean, nullable=False, server_default='true'),
    sa.Column('config', postgresql.JSONB, nullable=True, server_default='{}'),
    sa.Column('description', sa.Text, nullable=True),
    sa.Column('created_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
    sa.Column('updated_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
)

```

```
# Create indexes for agent_capabilities
```

```
op.create_index('idx_agent_cap_agent', 'agent_capabilities', ['agent_id'])
op.create_index('idx_agent_cap_name', 'agent_capabilities', ['capability_name'])
op.create_index('idx_agent_cap_enabled', 'agent_capabilities', ['enabled'])
op.create_index('idx_agent_cap_agent_name', 'agent_capabilities', ['agent_id', 'capability_name'])
op.create_unique_constraint('uq_agent_capability', 'agent_capabilities', ['agent_id', 'capability_name'])

# Create agent_metrics table
op.create_table(
    'agent_metrics',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), sa.ForeignKey('agents.id', ondelete='CASCADE'), nullable=False),
    sa.Column('metric_name', sa.String(255), nullable=False),
    sa.Column('metric_value', sa.Float, nullable=False),
    sa.Column('metric_type', postgresql.ENUM(name='metrictype'), nullable=False, server_default='gauge'),
    sa.Column('tags', postgresql.JSONB, nullable=True, server_default='{}'),
    sa.Column('recorded_at', sa.DateTime(timezone=True), nullable=False, server_default=sa.text('now()')),
)
```

```
# Create indexes for agent_metrics
op.create_index('idx_agent_metric_agent', 'agent_metrics', ['agent_id'])
op.create_index('idx_agent_metric_name', 'agent_metrics', ['metric_name'])
op.create_index('idx_agent_metric_recorded', 'agent_metrics', ['recorded_at'])
op.create_index('idx_agent_metric_agent_name', 'agent_metrics', ['agent_id', 'metric_name'])
op.create_index('idx_agent_metric_agent_recorded', 'agent_metrics', ['agent_id', 'recorded_at'])
```

```
def downgrade():
    """Drop agent state tables"""

```

```
# Drop tables in reverse order
op.drop_table('agent_metrics')
op.drop_table('agent_capabilities')
op.drop_table('agent_states')
op.drop_table('agent_configs')
```

```
# Drop ENUM types
op.execute("""
    DROP TYPE IF EXISTS metrictype;
    DROP TYPE IF EXISTS agentstatusdetail;
```

```
DROP TYPE IF EXISTS configtype;
```

```
""")
```

## FILE 4 COMPLETE (~150 lines)

---

## QUICK VERIFICATION (Part 1)

After creating these 4 files, verify they work:

### Step 1: Check Files Exist



bash

```
cd ~/optiinfra
```

```
# Check new model file
```

```
ls -la shared/database/models/agent_state.py
```

```
# Check modified files
```

```
ls -la shared/database/models/core.py
```

```
ls -la shared/database/models/__init__.py
```

```
# Check migration
```

```
ls -la shared/database/migrations/versions/002_agent_state_tables.py
```

### Step 2: Test Python Imports



bash

```
cd ~/optiinfra
```

# Test if new models can be imported

```
python -c "from shared.database.models.agent_state import AgentConfig, AgentState, AgentCapability, AgentMetric; print
```

# Expected output:

```
# ✓ Models import successfully
```

### Step 3: Run Alembic Migration



bash

```
cd ~/optiinfra/shared/database
```

# Check current migration state

```
alembic current
```

# Expected: 001\_core\_schema (current)

# Run the new migration

```
alembic upgrade head
```

# Expected output:

```
# INFO [alembic.runtime.migration] Running upgrade 001_core_schema -> 002_agent_state_tables, Create agent state i
```

### Step 4: Verify Tables Created



bash

```
# Connect to PostgreSQL
psql postgresql://optiinfra:password@localhost:5432/optiinfra
```

```
# List all tables (should now be 10 tables)
```

```
\dt
```

```
# Expected output (6 from 0.2a + 4 new):
```

```
# agent_configs
# agent_states
# agent_capabilities
# agent_metrics
# customers
# agents
# events
# recommendations
# approvals
# optimizations
```

```
# Describe agent_configs table
```

```
\d agent_configs
```

```
# Expected: Should show all columns, indexes, foreign key to agents
```

```
# Exit
```

```
\q
```

## 🎯 PART 1 SUMMARY

### Files Created/Modified:

- shared/database/models/agent\_state.py** - 4 new models (~400 lines)
- shared/database/models/core.py** - Modified (added relationships)
- shared/database/models/init.py** - Modified (exports)
- shared/database/migrations/versions/002\_agent\_state\_tables.py** - Migration (~150 lines)

### What Works Now:

- Can import all new models in Python
- Can run Alembic migration
- 4 new tables created in PostgreSQL
- All foreign keys to agents table working
- All indexes created

## What's in Part 2:

- Seed data (configs, states, capabilities, metrics for test agents)
  - Test fixtures
  - Test cases (10+ comprehensive tests)
  - All validation commands
  - Troubleshooting guide
  - Success criteria checklist
  - Git commit instructions
- 

## NEXT STEP

Download PART 2 to get:

- Complete seed data implementation
- Full test suite
- Step-by-step validation
- Troubleshooting
- Final checklist

**File name:** FOUNDATION-0.2b-Agent-State-Tables-PART2-Testing.md

---

## HOW TO DOWNLOAD THIS FILE

1. Click the "Copy" dropdown button at the top of this artifact
  2. Select "Download as txt"
  3. Save as: FOUNDATION-0.2b-Agent-State-Tables-PART1-Code.md
- 

 PART 1 COMPLETE! Ready for PART 2?