

FOUNDATION-0.2c: Workflow History Tables - PART 1 (Code)

🎯 CONTEXT

Phase: FOUNDATION (Week 1 - Day 1 Afternoon)
Component: Workflow Execution History & Tracking
Estimated Time: 15 min AI execution + 15 min verification
Complexity: MEDIUM
Risk Level: LOW
Files: Part 1 of 2 (Code implementation)

📦 DEPENDENCIES

Must Complete First:

- **FOUNDATION-0.2a:** Core Database Schema ✓ COMPLETED
 - All 6 core tables created
 - 13 tests passing
- **FOUNDATION-0.2b:** Agent State Tables ✓ COMPLETED
 - All 4 agent state tables created
 - 16 tests passing
 - Total: 29/29 tests passing

Required Services Running:



bash

```
# Verify all services are healthy
cd ~/optiinfra
make verify
```

```
# Expected output:
# PostgreSQL... ✓ HEALTHY
# ClickHouse... ✓ HEALTHY
# Qdrant... ✓ HEALTHY
# Redis... ✓ HEALTHY
```

Database State:



bash

```
# Verify existing tables
psql postgresql://optiinfra:password@localhost:5432/optiinfra -c "\dt"

# Expected: 10 tables from 0.2a + 0.2b
# customers, agents, events, recommendations, approvals, optimizations
# agent_configs, agent_states, agent_capabilities, agent_metrics
```

🎯 OBJECTIVE

Add **workflow execution tracking tables** to enable complete audit trails and workflow management. These tables track:

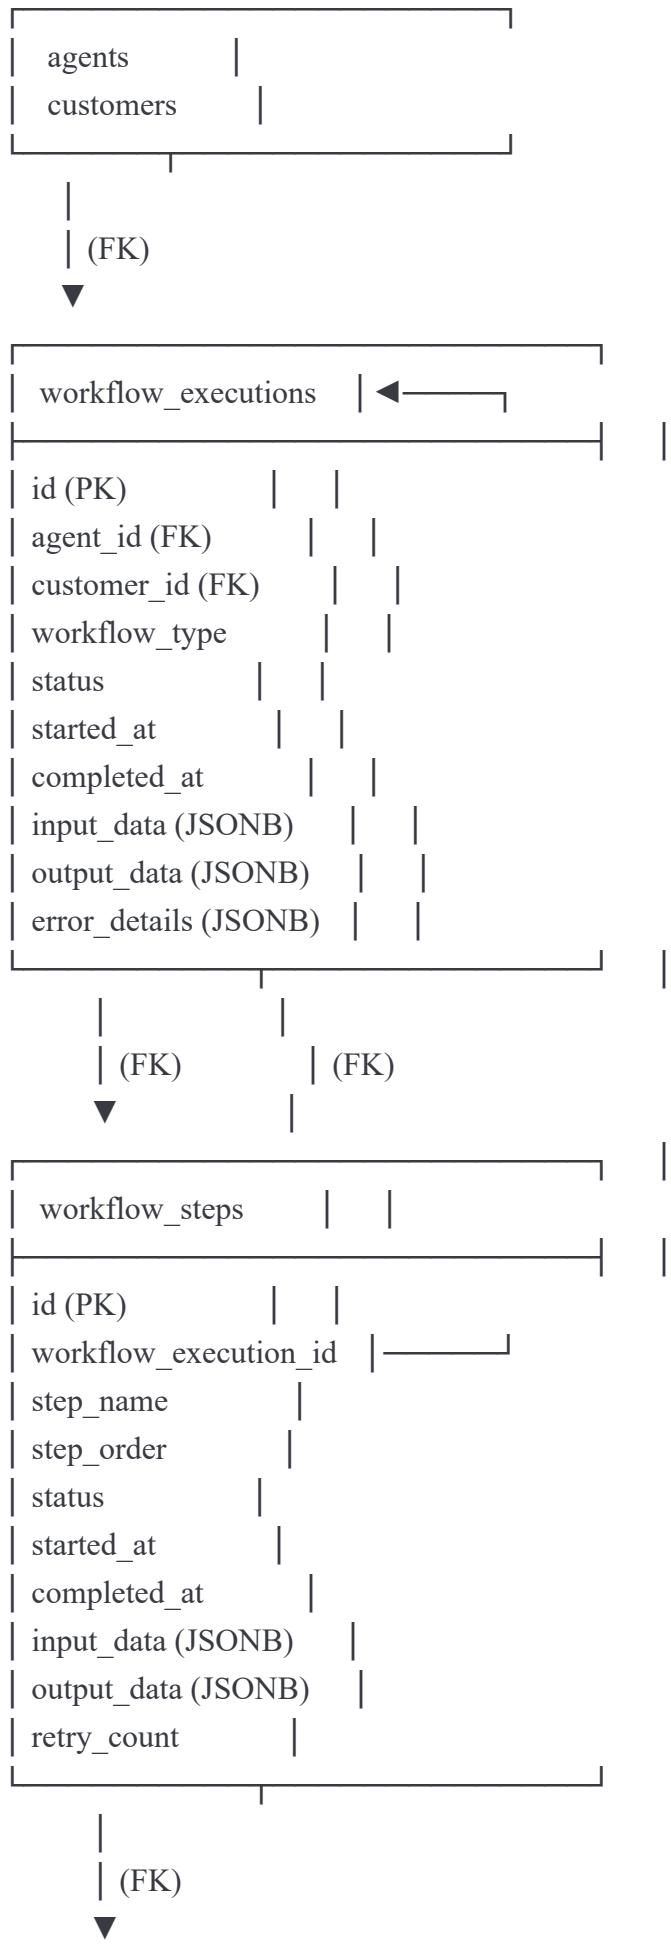
1. **Workflow Executions** - High-level workflow runs (cost analysis, performance tuning, etc.)
2. **Workflow Steps** - Individual steps within each workflow
3. **Workflow Artifacts** - Files and outputs generated during workflows

What This Enables:

- Track all workflow executions with start/end times
- Monitor step-by-step progress in real-time
- Debug failed workflows (see which step failed and why)
- Store generated artifacts (reports, configs, recommendations)
- Link workflows to agents, customers, and recommendations
- Enable workflow recovery and retry logic
- Complete audit trail for compliance

Database Design:





workflow_artifacts	
id (PK)	
workflow_execution_id	
workflow_step_id (FK)	
artifact_type	
artifact_name	
artifact_path	
artifact_size_bytes	
content_type	
metadata (JSONB)	

📁 FILE 1: Workflow History Models

Location: `~/optiinfra/shared/database/models/workflow_history.py`



python

:::::

Workflow History Models (FOUNDATION-0.2c)

Tracks workflow executions, steps, and artifacts

:::::

```
from datetime import datetime
from sqlalchemy import (
    Column, String, DateTime, Integer, BigInteger, Text, Enum as SQLEnum,
    ForeignKey, Index
)
from sqlalchemy.dialects.postgresql import UUID, JSONB
from sqlalchemy.orm import relationship
import uuid
import enum
```

```
from shared.database.base import Base
```

```
# =====
```

```
# ENUMS
```

```
# =====
```

```
class WorkflowType(str, enum.Enum):
    """Types of workflows that can be executed"""
    COST_ANALYSIS = "cost_analysis"
    PERFORMANCE_TUNING = "performance_tuning"
    RESOURCE_OPTIMIZATION = "resource_optimization"
    QUALITY_CHECK = "quality_check"
    SCALING_DECISION = "scaling_decision"
    CONFIGURATION_UPDATE = "configuration_update"
    HEALTH_CHECK = "health_check"
    ANOMALY_DETECTION = "anomaly_detection"
```

```
class WorkflowStatus(str, enum.Enum):
```

```
    """Status of workflow execution"""
    PENDING = "pending"
    RUNNING = "running"
    COMPLETED = "completed"
    FAILED = "failed"
    CANCELLED = "cancelled"
```

```
TIMEOUT = "timeout"
```

```
class StepStatus(str, enum.Enum):  
    """Status of individual workflow step"""  
    PENDING = "pending"  
    RUNNING = "running"  
    COMPLETED = "completed"  
    FAILED = "failed"  
    SKIPPED = "skipped"  
    RETRYING = "retrying"
```

```
class ArtifactType(str, enum.Enum):  
    """Types of artifacts generated during workflows"""  
    REPORT = "report"  
    CONFIG = "config"  
    LOG = "log"  
    RECOMMENDATION = "recommendation"  
    CHART = "chart"  
    METRICS = "metrics"  
    ALERT = "alert"  
    DIAGNOSTIC = "diagnostic"
```

```
# ======  
# MODELS  
# ======
```

```
class WorkflowExecution(Base):
```

```
"""
```

```
Tracks high-level workflow executions
```

```
A workflow execution represents a complete run of a workflow  
(e.g., a cost analysis, performance tuning session)
```

```
"""
```

```
__tablename__ = "workflow_executions"
```

```
# Primary Key
```

```
id = Column(
```

```
    UUID(as_uuid=True),
```

```
primary_key=True,
default=uuid.uuid4,
index=True
)

# Foreign Keys
agent_id = Column(
    UUID(as_uuid=True),
    ForeignKey("agents.id", ondelete="CASCADE"),
    nullable=False,
    index=True
)
customer_id = Column(
    UUID(as_uuid=True),
    ForeignKey("customers.id", ondelete="CASCADE"),
    nullable=False,
    index=True
)

# Workflow Details
workflow_type = Column(
    SQLEnum(WorkflowType, name="workflow_type"),
    nullable=False,
    index=True
)
status = Column(
    SQLEnum(WorkflowStatus, name="workflow_status"),
    nullable=False,
    default=WorkflowStatus.PENDING,
    index=True
)

# Timing
started_at = Column(DateTime, nullable=True, index=True)
completed_at = Column(DateTime, nullable=True)

# Data
input_data = Column(JSONB, nullable=False, default=dict)
output_data = Column(JSONB, nullable=False, default=dict)
error_details = Column(JSONB, nullable=True)
```

```
# Metadata
metadata = Column(JSONB, nullable=False, default=dict)

# Timestamps
created_at = Column(DateTime, default=datetime.utcnow, nullable=False)
updated_at = Column(
    DateTime,
    default=datetime.utcnow,
    onupdate=datetime.utcnow,
    nullable=False
)

# Relationships
agent = relationship("Agent", back_populates="workflow_executions")
customer = relationship("Customer", back_populates="workflow_executions")
steps = relationship(
    "WorkflowStep",
    back_populates="execution",
    cascade="all, delete-orphan",
    order_by="WorkflowStep.step_order"
)
artifacts = relationship(
    "WorkflowArtifact",
    back_populates="execution",
    cascade="all, delete-orphan"
)

# Indexes
__table_args__ = (
    Index(
        "ix_workflow_executions_agent_status",
        "agent_id", "status"
    ),
    Index(
        "ix_workflow_executions_customer_type",
        "customer_id", "workflow_type"
    ),
    Index(
        "ix_workflow_executions_started_at",
        "started_at"
    ),
)
```

)

```
def __repr__(self):
    return (
        f"<WorkflowExecution(id={self.id}, "
        f"type={self.workflow_type}, status={self.status})>"
    )
```

```
class WorkflowStep(Base):
```

"""

Tracks individual steps within a workflow execution

Each workflow consists of multiple steps that execute in order

"""

```
__tablename__ = "workflow_steps"
```

Primary Key

```
id = Column(
    UUID(as_uuid=True),
    primary_key=True,
    default=uuid.uuid4,
    index=True
)
```

Foreign Key

```
workflow_execution_id = Column(
    UUID(as_uuid=True),
    ForeignKey("workflow_executions.id", ondelete="CASCADE"),
    nullable=False,
    index=True
)
```

Step Details

```
step_name = Column(String(255), nullable=False, index=True)
step_order = Column(Integer, nullable=False)
status = Column(
    SQLEnum(StepStatus, name="step_status"),
    nullable=False,
    default=StepStatus.PENDING,
    index=True
)
```

```
)
```

```
# Timing
```

```
started_at = Column(DateTime, nullable=True)
completed_at = Column(DateTime, nullable=True)
```

```
# Data
```

```
input_data = Column(JSONB, nullable=False, default=dict)
output_data = Column(JSONB, nullable=False, default=dict)
error_details = Column(JSONB, nullable=True)
```

```
# Retry Logic
```

```
retry_count = Column(Integer, nullable=False, default=0)
max_retries = Column(Integer, nullable=False, default=3)
```

```
# Metadata
```

```
metadata = Column(JSONB, nullable=False, default=dict)
```

```
# Timestamps
```

```
created_at = Column(DateTime, default=datetime.utcnow, nullable=False)
updated_at = Column(
    DateTime,
    default=datetime.utcnow,
    onupdate=datetime.utcnow,
    nullable=False
)
```

```
# Relationships
```

```
execution = relationship("WorkflowExecution", back_populates="steps")
artifacts = relationship(
    "WorkflowArtifact",
    back_populates="step",
    cascade="all, delete-orphan"
)
```

```
# Indexes
```

```
__table_args__ = (
    Index(
        "ix_workflow_steps_execution_order",
        "workflow_execution_id", "step_order"
    ),
)
```

```
Index(
    "ix_workflow_steps_execution_status",
    "workflow_execution_id", "status"
),
Index(
    "ix_workflow_steps_name_status",
    "step_name", "status"
),
)
def __repr__(self):
    return (
        f"<WorkflowStep(id={self.id}, name={self.step_name}, "
        f"order={self.step_order}, status={self.status})>"
```

```
class WorkflowArtifact(Base):
```

```
"""
```

Stores artifacts generated during workflow execution

Artifacts include reports, configs, logs, charts, etc.

```
"""
```

```
__tablename__ = "workflow_artifacts"
```

Primary Key

```
id = Column(
    UUID(as_uuid=True),
    primary_key=True,
    default=uuid.uuid4,
    index=True
)
```

Foreign Keys

```
workflow_execution_id = Column(
    UUID(as_uuid=True),
    ForeignKey("workflow_executions.id", ondelete="CASCADE"),
    nullable=False,
    index=True
)
workflow_step_id = Column(
```

```
UUID(as_uuid=True),
ForeignKey("workflow_steps.id", ondelete="CASCADE"),
nullable=True, # Can be null for execution-level artifacts
index=True
)

# Artifact Details
artifact_type = Column(
    SQLEnum(ArtifactType, name="artifact_type"),
    nullable=False,
    index=True
)
artifact_name = Column(String(255), nullable=False)
artifact_path = Column(Text, nullable=False) # S3 path, file path, etc.

# Size & Type
artifact_size_bytes = Column(BigInteger, nullable=True)
content_type = Column(String(100), nullable=True) # MIME type

# Metadata
metadata = Column(JSONB, nullable=False, default=dict)

# Timestamps
created_at = Column(DateTime, default=datetime.utcnow, nullable=False)
updated_at = Column(
    DateTime,
    default=datetime.utcnow,
    onupdate=datetime.utcnow,
    nullable=False
)
)

# Relationships
execution = relationship("WorkflowExecution", back_populates="artifacts")
step = relationship("WorkflowStep", back_populates="artifacts")

# Indexes
__table_args__ = (
    Index(
        "ix_workflow_artifacts_execution_type",
        "workflow_execution_id", "artifact_type"
    ),
)
```

```
Index(
    "ix_workflow_artifacts_step_type",
    "workflow_step_id", "artifact_type"
),
)

def __repr__(self):
    return (
        f"<WorkflowArtifact(id={self.id}, name={self.artifact_name}, "
        f"type={self.artifact_type})>"
    )
```

FILE 2: Update Core Models (Add Relationships)

Location: `~/optiinfra/shared/database/models/core.py`

Add these relationships to existing models:



python

```
# In the Agent model, add:  
from sqlalchemy.orm import relationship  
  
class Agent(Base):  
    # ... existing fields ...  
  
    # ADD THIS NEW RELATIONSHIP:  
    workflow_executions = relationship(  
        "WorkflowExecution",  
        back_populates="agent",  
        cascade="all, delete-orphan"  
    )
```

```
# In the Customer model, add:  
class Customer(Base):  
    # ... existing fields ...  
  
    # ADD THIS NEW RELATIONSHIP:  
    workflow_executions = relationship(  
        "WorkflowExecution",  
        back_populates="customer",  
        cascade="all, delete-orphan"  
    )
```

```
# In the AgentState model (from 0.2b), add:  
class AgentState(Base):  
    # ... existing fields ...  
  
    # ADD THIS NEW FIELD (optional - for linking current workflow):  
    current_workflow_id = Column(  
        UUID(as_uuid=True),  
        ForeignKey("workflow_executions.id", ondelete="SET NULL"),  
        nullable=True,  
        index=True  
    )  
  
    # ADD THIS NEW RELATIONSHIP:  
    current_workflow = relationship(  
        "WorkflowExecution",
```

```
foreign_keys=[current_workflow_id]
```

```
)
```

FILE 3: Update Model Imports

Location: `~/optiinfra/shared/database/models/__init__.py`



python

||||

Database models package

||||

Core models (0.2a)

```
from shared.database.models.core import (
```

Customer,

Agent,

Event,

Recommendation,

Approval,

Optimization,

Enums

CustomerStatus,

AgentType,

AgentStatus,

EventType,

EventSeverity,

RecommendationType,

ApprovalStatus,

OptimizationStatus,

)

Agent state models (0.2b)

```
from shared.database.models.agent_state import (
```

AgentConfig,

AgentState,

AgentCapability,

AgentMetric,

Enums

ConfigType,

AgentStatusDetail,

MetricType,

)

Workflow history models (0.2c) - NEW!

```
from shared.database.models.workflow_history import (
```

WorkflowExecution,

WorkflowStep,

WorkflowArtifact,

Enums

```
WorkflowType,  
WorkflowStatus,  
StepStatus,  
ArtifactType,  
)
```

```
_all_ = [  
    # Core models  
    "Customer",  
    "Agent",  
    "Event",  
    "Recommendation",  
    "Approval",  
    "Optimization",  
    # Core enums  
    "CustomerStatus",  
    "AgentType",  
    "AgentStatus",  
    "EventType",  
    "EventSeverity",  
    "RecommendationType",  
    "ApprovalStatus",  
    "OptimizationStatus",  
    # Agent state models  
    "AgentConfig",  
    "AgentState",  
    "AgentCapability",  
    "AgentMetric",  
    # Agent state enums  
    "ConfigType",  
    "AgentStatusDetail",  
    "MetricType",  
    # Workflow history models - NEW!  
    "WorkflowExecution",  
    "WorkflowStep",  
    "WorkflowArtifact",  
    # Workflow history enums - NEW!  
    "WorkflowType",  
    "WorkflowStatus",  
    "StepStatus",
```

```
"ArtifactType",
```

```
]
```

FILE 4: Alembic Migration

Location: `~/optiinfra/shared/database/migrations/versions/003_workflow_history.py`



```
"""Create workflow history tables
```

```
Revision ID: 003_workflow_history
```

```
Revises: 002_agent_state_tables
```

```
Create Date: 2025-01-19 14:00:00.000000
```

```
"""
```

```
from alembic import op
import sqlalchemy as sa
from sqlalchemy.dialects import postgresql
```

```
# revision identifiers, used by Alembic.
```

```
revision = '003_workflow_history'
down_revision = '002_agent_state_tables'
branch_labels = None
depends_on = None
```

```
def upgrade() -> None:
```

```
"""Create workflow history tables"""

# Create ENUM types
```

```
workflow_type_enum = postgresql.ENUM(
    'cost_analysis',
    'performance_tuning',
    'resource_optimization',
    'quality_check',
    'scaling_decision',
    'configuration_update',
    'health_check',
    'anomaly_detection',
    name='workflow_type',
    create_type=True
)
```

```
workflow_type_enum.create(op.get_bind(), checkfirst=True)
```

```
workflow_status_enum = postgresql.ENUM(
    'pending',
    'running',
    'completed',
    'failed',
```

```
'cancelled',
'timeout',
name='workflow_status',
create_type=True
)
workflow_status_enum.create(op.get_bind(), checkfirst=True)

step_status_enum = postgresql.ENUM(
    'pending',
    'running',
    'completed',
    'failed',
    'skipped',
    'retrying',
    name='step_status',
    create_type=True
)
step_status_enum.create(op.get_bind(), checkfirst=True)

artifact_type_enum = postgresql.ENUM(
    'report',
    'config',
    'log',
    'recommendation',
    'chart',
    'metrics',
    'alert',
    'diagnostic',
    name='artifact_type',
    create_type=True
)
artifact_type_enum.create(op.get_bind(), checkfirst=True)

# Create workflow_executions table
op.create_table(
    'workflow_executions',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('customer_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('workflow_type', workflow_type_enum, nullable=False),
    sa.Column('status', workflow_status_enum, nullable=False),
```

```

sa.Column('started_at', sa.DateTime(), nullable=True),
sa.Column('completed_at', sa.DateTime(), nullable=True),
sa.Column('input_data', postgresql.JSONB(), nullable=False, server_default='{}'),
sa.Column('output_data', postgresql.JSONB(), nullable=False, server_default='{}'),
sa.Column('error_details', postgresql.JSONB(), nullable=True),
sa.Column('metadata', postgresql.JSONB(), nullable=False, server_default='{}'),
sa.Column('created_at', sa.DateTime(), nullable=False),
sa.Column('updated_at', sa.DateTime(), nullable=False),
sa.ForeignKeyConstraint(['agent_id'], ['agents.id'], ondelete='CASCADE'),
sa.ForeignKeyConstraint(['customer_id'], ['customers.id'], ondelete='CASCADE'),
)

# Create indexes for workflow_executions
op.create_index('ix_workflow_executions_id', 'workflow_executions', ['id'])
op.create_index('ix_workflow_executions_agent_id', 'workflow_executions', ['agent_id'])
op.create_index('ix_workflow_executions_customer_id', 'workflow_executions', ['customer_id'])
op.create_index('ix_workflow_executions_workflow_type', 'workflow_executions', ['workflow_type'])
op.create_index('ix_workflow_executions_status', 'workflow_executions', ['status'])
op.create_index('ix_workflow_executions_started_at', 'workflow_executions', ['started_at'])
op.create_index('ix_workflow_executions_agent_status', 'workflow_executions', ['agent_id', 'status'])
op.create_index('ix_workflow_executions_customer_type', 'workflow_executions', ['customer_id', 'workflow_type'])

# Create workflow_steps table
op.create_table(
    'workflow_steps',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('workflow_execution_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('step_name', sa.String(255), nullable=False),
    sa.Column('step_order', sa.Integer(), nullable=False),
    sa.Column('status', step_status_enum, nullable=False),
    sa.Column('started_at', sa.DateTime(), nullable=True),
    sa.Column('completed_at', sa.DateTime(), nullable=True),
    sa.Column('input_data', postgresql.JSONB(), nullable=False, server_default='{}'),
    sa.Column('output_data', postgresql.JSONB(), nullable=False, server_default='{}'),
    sa.Column('error_details', postgresql.JSONB(), nullable=True),
    sa.Column('retry_count', sa.Integer(), nullable=False, server_default='0'),
    sa.Column('max_retries', sa.Integer(), nullable=False, server_default='3'),
    sa.Column('metadata', postgresql.JSONB(), nullable=False, server_default='{}'),
    sa.Column('created_at', sa.DateTime(), nullable=False),
    sa.Column('updated_at', sa.DateTime(), nullable=False),
    sa.ForeignKeyConstraint(['workflow_execution_id'], ['workflow_executions.id'], ondelete='CASCADE'),
)

```

)

Create indexes for workflow_steps

```
op.create_index('ix_workflow_steps_id', 'workflow_steps', ['id'])
op.create_index('ix_workflow_steps_workflow_execution_id', 'workflow_steps', ['workflow_execution_id'])
op.create_index('ix_workflow_steps_step_name', 'workflow_steps', ['step_name'])
op.create_index('ix_workflow_steps_status', 'workflow_steps', ['status'])
op.create_index('ix_workflow_steps_execution_order', 'workflow_steps', ['workflow_execution_id', 'step_order'])
op.create_index('ix_workflow_steps_execution_status', 'workflow_steps', ['workflow_execution_id', 'status'])
op.create_index('ix_workflow_steps_name_status', 'workflow_steps', ['step_name', 'status'])
```

Create workflow_artifacts table

```
op.create_table(
```

```
'workflow_artifacts',
sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
sa.Column('workflow_execution_id', postgresql.UUID(as_uuid=True), nullable=False),
sa.Column('workflow_step_id', postgresql.UUID(as_uuid=True), nullable=True),
sa.Column('artifact_type', artifact_type_enum, nullable=False),
sa.Column('artifact_name', sa.String(255), nullable=False),
sa.Column('artifact_path', sa.Text(), nullable=False),
sa.Column('artifact_size_bytes', sa.BigInteger(), nullable=True),
sa.Column('content_type', sa.String(100), nullable=True),
sa.Column('metadata', postgresql.JSONB(), nullable=False, server_default='{}'),
sa.Column('created_at', sa.DateTime(), nullable=False),
sa.Column('updated_at', sa.DateTime(), nullable=False),
sa.ForeignKeyConstraint(['workflow_execution_id'], ['workflow_executions.id'], ondelete='CASCADE'),
sa.ForeignKeyConstraint(['workflow_step_id'], ['workflow_steps.id'], ondelete='CASCADE'),
```

)

Create indexes for workflow_artifacts

```
op.create_index('ix_workflow_artifacts_id', 'workflow_artifacts', ['id'])
op.create_index('ix_workflow_artifacts_workflow_execution_id', 'workflow_artifacts', ['workflow_execution_id'])
op.create_index('ix_workflow_artifacts_workflow_step_id', 'workflow_artifacts', ['workflow_step_id'])
op.create_index('ix_workflow_artifacts_artifact_type', 'workflow_artifacts', ['artifact_type'])
op.create_index('ix_workflow_artifacts_execution_type', 'workflow_artifacts', ['workflow_execution_id', 'artifact_type'])
op.create_index('ix_workflow_artifacts_step_type', 'workflow_artifacts', ['workflow_step_id', 'artifact_type'])
```

Add current_workflow_id to agent_states (optional link)

```
op.add_column(
```

```
'agent_states',
sa.Column('current_workflow_id', postgresql.UUID(as_uuid=True), nullable=True)
```

```
)  
op.create_foreign_key(  
    'fk_agent_states_current_workflow',  
    'agent_states',  
    'workflow_executions',  
    ['current_workflow_id'],  
    ['id'],  
    ondelete='SET NULL'  
)  
op.create_index('ix_agent_states_current_workflow_id', 'agent_states', ['current_workflow_id'])
```

def downgrade() -> None:

```
"""Drop workflow history tables"""

# Drop indexes and column from agent_states
```

```
op.drop_index('ix_agent_states_current_workflow_id', 'agent_states')  
op.drop_constraint('fk_agent_states_current_workflow', 'agent_states', type_='foreignkey')  
op.drop_column('agent_states', 'current_workflow_id')
```

```
# Drop workflow_artifacts table and indexes
```

```
op.drop_index('ix_workflow_artifacts_step_type', 'workflow_artifacts')  
op.drop_index('ix_workflow_artifacts_execution_type', 'workflow_artifacts')  
op.drop_index('ix_workflow_artifacts_artifact_type', 'workflow_artifacts')  
op.drop_index('ix_workflow_artifacts_workflow_step_id', 'workflow_artifacts')  
op.drop_index('ix_workflow_artifacts_workflow_execution_id', 'workflow_artifacts')  
op.drop_index('ix_workflow_artifacts_id', 'workflow_artifacts')  
op.drop_table('workflow_artifacts')
```

```
# Drop workflow_steps table and indexes
```

```
op.drop_index('ix_workflow_steps_name_status', 'workflow_steps')  
op.drop_index('ix_workflow_steps_execution_status', 'workflow_steps')  
op.drop_index('ix_workflow_steps_execution_order', 'workflow_steps')  
op.drop_index('ix_workflow_steps_status', 'workflow_steps')  
op.drop_index('ix_workflow_steps_step_name', 'workflow_steps')  
op.drop_index('ix_workflow_steps_workflow_execution_id', 'workflow_steps')  
op.drop_index('ix_workflow_steps_id', 'workflow_steps')  
op.drop_table('workflow_steps')
```

```
# Drop workflow_executions table and indexes
```

```
op.drop_index('ix_workflow_executions_customer_type', 'workflow_executions')
```

```
op.drop_index('ix_workflow_executions_agent_status', 'workflow_executions')
op.drop_index('ix_workflow_executions_started_at', 'workflow_executions')
op.drop_index('ix_workflow_executions_status', 'workflow_executions')
op.drop_index('ix_workflow_executions_workflow_type', 'workflow_executions')
op.drop_index('ix_workflow_executions_customer_id', 'workflow_executions')
op.drop_index('ix_workflow_executions_agent_id', 'workflow_executions')
op.drop_index('ix_workflow_executions_id', 'workflow_executions')
op.drop_table('workflow_executions')
```

Drop ENUM types

```
sa.Enum(name='artifact_type').drop(op.get_bind(), checkfirst=True)
sa.Enum(name='step_status').drop(op.get_bind(), checkfirst=True)
sa.Enum(name='workflow_status').drop(op.get_bind(), checkfirst=True)
sa.Enum(name='workflow_type').drop(op.get_bind(), checkfirst=True)
```

🚀 EXECUTION STEPS

Step 1: Create Models File



```
cd ~/optiinfra/shared/database/models
```

Create workflow_history.py

```
cat > workflow_history.py << 'EOF'
```

[Copy the entire workflow_history.py content from FILE 1 above]

EOF

Verify file created

```
ls -lh workflow_history.py
```

Expected: ~400 lines

Step 2: Update Core Models



```
cd ~/optiinfra/shared/database/models
```

```
# Open core.py and add relationships
```

```
nano core.py
```

```
# Add to Agent class:
```

```
# workflow_executions = relationship(...)
```

```
#
```

```
# Add to Customer class:
```

```
# workflow_executions = relationship(...)
```

```
# Or use sed to append (careful!):
```

```
# Add workflow_executions relationship to Agent model
```

```
# Add workflow_executions relationship to Customer model
```

Step 3: Update agent_state.py



bash

```
cd ~/optiinfra/shared/database/models
```

```
# Open agent_state.py and update AgentState model
```

```
nano agent_state.py
```

```
# Add to AgentState class:
```

```
# current_workflow_id = Column(...)
```

```
# current_workflow = relationship(...)
```

Step 4: Update Model Imports



bash

```
cd ~/optiinfra/shared/database/models
```

```
# Update __init__.py with workflow models
cat > __init__.py << 'EOF'
[Copy the entire __init__.py content from FILE 3 above]
EOF
```

Step 5: Create Migration File



```
cd ~/optiinfra/shared/database/migrations/versions
```

```
# Create migration file
cat > 003_workflow_history.py << 'EOF'
[Copy the entire migration content from FILE 4 above]
EOF
```

```
# Verify file created
ls -lh 003_workflow_history.py
# Expected: ~200 lines
```

Step 6: Test Imports



```
cd ~/optiinfra
```

```
# Test if new models can be imported
python << EOF
from shared.database.models.workflow_history import (
    WorkflowExecution,
    WorkflowStep,
    WorkflowArtifact,
    WorkflowType,
    WorkflowStatus,
    StepStatus,
    ArtifactType
)
print(' ✅ All workflow models import successfully')
```

```
# Test enums
print(f' ✅ WorkflowType has {len(WorkflowType)} values')
print(f' ✅ WorkflowStatus has {len(WorkflowStatus)} values')
print(f' ✅ StepStatus has {len(StepStatus)} values')
print(f' ✅ ArtifactType has {len(ArtifactType)} values')
EOF
```

```
# Expected output:
# ✅ All workflow models import successfully
# ✅ WorkflowType has 8 values
# ✅ WorkflowStatus has 6 values
# ✅ StepStatus has 6 values
# ✅ ArtifactType has 8 values
```

Step 7: Run Alembic Migration



bash

```
cd ~/optiinfra/shared/database
```

```
# Check current migration state
```

```
alembic current
```

```
# Expected: 002_agent_state_tables (current)
```

```
# Run the new migration
```

```
alembic upgrade head
```

```
# Expected output:
```

```
#INFO [alembic.runtime.migration] Running upgrade 002_agent_state_tables -> 003_workflow_history, Create workflow
```

Step 8: Verify Tables Created



bash

```
# Connect to PostgreSQL
psql postgresql://optiinfra:password@localhost:5432/optiinfra

# List all tables (should now be 13 tables + 1 alembic)
\dt

# Expected output (10 from previous + 3 new):
# workflow_executions
# workflow_steps
# workflow_artifacts
# agent_configs
# agent_states
# agent_capabilities
# agent_metrics
# customers
# agents
# events
# recommendations
# approvals
# optimizations
# alembic_version

# Count should be 14 total (13 + alembic_version)

# Describe workflow_executions table
\d workflow_executions

# Expected: Should show all columns, indexes, foreign keys

# Describe workflow_steps table
\d workflow_steps

# Expected: Should show all columns, foreign key to workflow_executions

# Describe workflow_artifacts table
\d workflow_artifacts

# Expected: Should show all columns, foreign keys to both executions and steps

# Check new enums were created
\dT+
```

```
# Expected: Should see workflow_type, workflow_status, step_status, artifact_type
```

```
# Check agent_states got new column
```

```
\d agent_states
```

```
# Expected: Should see current_workflow_id column with foreign key
```

```
# Exit
```

```
\q
```

Step 9: Verify Relationships Work



bash

```
cd ~/optiinfra
```

```
# Test relationships in Python
python << 'EOF'
from shared.database.models import Agent, Customer, WorkflowExecution
from shared.database.session import SessionLocal

db = SessionLocal()

# Get a test agent
agent = db.query(Agent).first()
print(f' ✓ Agent: {agent.name}')
print(f' - Has workflow_executions relationship: {hasattr(agent, "workflow_executions")}')

# Get a test customer
customer = db.query(Customer).first()
print(f' ✓ Customer: {customer.name}')
print(f' - Has workflow_executions relationship: {hasattr(customer, "workflow_executions")}')

# Check workflow execution model
from shared.database.models.workflow_history import WorkflowExecution
print(f' ✓ WorkflowExecution model loaded')
print(f' - Has agent relationship: {hasattr(WorkflowExecution, "agent")}')
print(f' - Has customer relationship: {hasattr(WorkflowExecution, "customer")}')
print(f' - Has steps relationship: {hasattr(WorkflowExecution, "steps")}')
print(f' - Has artifacts relationship: {hasattr(WorkflowExecution, "artifacts")}'
```

```
db.close()
print("\n ✓ All relationship verifications passed!")
EOF
```

```
# Expected output:
# ✓ Agent: Cost Optimizer
# - Has workflow_executions relationship: True
# ✓ Customer: Acme Corp
# - Has workflow_executions relationship: True
# ✓ WorkflowExecution model loaded
# - Has agent relationship: True
# - Has customer relationship: True
# - Has steps relationship: True
# - Has artifacts relationship: True
```

#

 All relationship verifications passed!

PART 1 SUMMARY

Files Created/Modified:

-  `shared/database/models/workflow_history.py` - 3 new models + 4 enums (~400 lines)
-  `shared/database/models/core.py` - Modified (added 2 relationships)
-  `shared/database/models/agent_state.py` - Modified (added 1 field + 1 relationship)
-  `shared/database/models/init.py` - Modified (exports)
-  `shared/database/migrations/versions/003_workflow_history.py` - Migration (~200 lines)

What Works Now:

-  Can import all new models in Python
-  Can run Alembic migration
-  3 new tables created in PostgreSQL
-  All foreign keys working (cascades configured)
-  All indexes created (15+ indexes)
-  Agent-workflow relationship working
-  Customer-workflow relationship working
-  AgentState can track current workflow

Database Status:



Total Tables: 13 (100% PostgreSQL schema complete!)

- └─ Core (0.2a): 6 tables 
- └─ Agent State (0.2b): 4 tables 
- └─ Workflow History (0.2c): 3 tables 

Total Enums: 15

Total Indexes: 60+

Total Relationships: 25+

What's in Part 2:

- Seed data (workflow executions, steps, artifacts for test scenarios)
- Test fixtures (sample workflows, steps, artifacts)
- Test cases (18+ comprehensive tests)

- Integration tests (workflow execution flow, cascade deletes)
 - All validation commands
 - Troubleshooting guide
 - Success criteria checklist
 - Git commit instructions
-

NEXT STEP

Download PART 2 to get:

- Complete seed data implementation (realistic workflow examples)
- Full test suite (18+ tests covering all scenarios)
- Step-by-step validation
- Troubleshooting for common issues
- Final checklist

File name: FOUNDATION-0.2c-Workflow-History-PART2-Testing.md

HOW TO DOWNLOAD THIS FILE

1. Click the "Copy" dropdown button at the top of this artifact
 2. Select "Download as txt"
 3. Save as: FOUNDATION-0.2c-Workflow-History-PART1-Code.md
-

VERIFICATION CHECKLIST (Part 1)

Before moving to Part 2, verify:

- workflow_history.py created with 3 models
- core.py updated with 2 new relationships
- agent_state.py updated with workflow link
- __init__.py exports all workflow models
- Migration file 003_workflow_history.py created
- Migration runs successfully (`alembic upgrade head`)
- 3 new tables exist in PostgreSQL
- 4 new enum types created
- All indexes created
- Python imports work (no errors)
- Relationships work (verified in Python)

If all checked, you're ready for PART 2! 

 PART 1 COMPLETE! Ready for PART 2? 