

FOUNDATION-0.2e: Application Schema - PART 1 (Code)

🎯 CONTEXT

Phase: FOUNDATION (Week 1 - Day 2 Afternoon)

Component: Application Agent Database Schema

Estimated Time: 15 min AI execution + 15 min verification

Complexity: MEDIUM

Risk Level: LOW

Files: Part 1 of 2 (Code implementation)

MILESTONE: This completes the PostgreSQL schema (100%)! 🎉

📦 DEPENDENCIES

Must Complete First:

- **FOUNDATION-0.2a:** Core Database Schema ✓ COMPLETED
- **FOUNDATION-0.2b:** Agent State Tables ✓ COMPLETED
- **FOUNDATION-0.2c:** Workflow History Tables ✓ COMPLETED
- **FOUNDATION-0.2d:** Resource Schema Tables ✓ COMPLETED

Required Services Running:

```
bash

# Verify all services are healthy
cd ~/optiinfra
make verify

# Expected output:
# PostgreSQL... ✓ HEALTHY
# ClickHouse... ✓ HEALTHY
# Qdrant... ✓ HEALTHY
# Redis... ✓ HEALTHY
```

Database State:

```
bash
```

```
# Verify existing tables
psql postgresql://optiinfra:password@localhost:5432/optiinfra -c "\dt"
```

Expected: 15 tables from 0.2a + 0.2b + 0.2c + 0.2d

🎯 OBJECTIVE

Add **Application Agent-specific database tables** to track LLM output quality, baselines, and regressions.

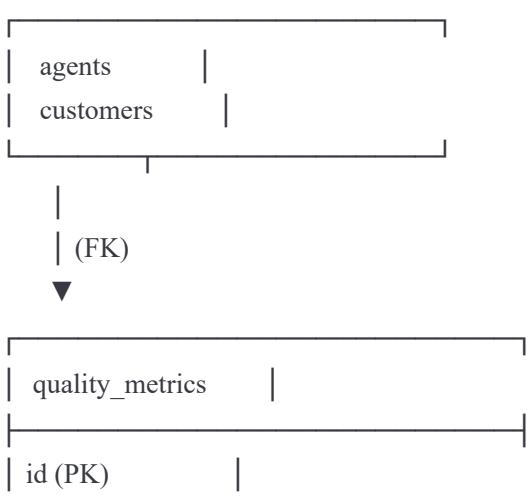
These tables enable:

1. **Quality Metrics** - Track per-request quality scores (relevance, coherence, factuality, hallucinations)
2. **Quality Baselines** - Establish and maintain quality baselines per model
3. **Quality Regressions** - Detect and track quality degradations

What This Enables:

- Monitor LLM output quality in real-time
- Detect hallucinations and toxicity
- Establish quality baselines per model/version
- Automatically detect quality regressions
- Trigger rollbacks when quality drops
- Prevent bad optimizations from reaching production
- A/B testing with quality validation

Database Design:



agent_id (FK)	
customer_id (FK)	
request_id	
model_name	
model_version	
relevance_score	
coherence_score	
factuality_score	
hallucination_detected	
overall_quality_score	
timestamp	
metadata (JSONB)	

quality_baselines	
id (PK)	
agent_id (FK)	
customer_id (FK)	
model_name	
model_version	
baseline_type	
sample_size	
avg_relevance_score	
avg_coherence_score	
avg_overall_score	
established_at	
valid_until	

quality_regressions	
id (PK)	
agent_id (FK)	
customer_id (FK)	
baseline_id (FK)	
workflow_execution_id	
regression_type	
severity	
metric_name	
baseline_value	
current_value	

```
| delta_percent      |
| action_taken       |
| resolved_at        |
```

📁 FILE 1: Application Schema Models

Location: `(~/optiinfra/shared/database/models/application_schema.py)`

```
python
```

....

Application Schema Models (FOUNDATION-0.2e)
Tracks LLM quality metrics, baselines, and regressions

....

```
from datetime import datetime
from sqlalchemy import (
    Column, String, DateTime, Float, Integer, Boolean, Text,
    Enum as SQLEnum, ForeignKey, Index
)
from sqlalchemy.dialects.postgresql import UUID, JSONB
from sqlalchemy.orm import relationship
import uuid
import enum

from shared.database.base import Base

# =====
# ENUMS
# =====

class BaselineType(str, enum.Enum):
    """Types of quality baselines"""
    INITIAL = "initial"
    UPDATED = "updated"
    ROLLBACK = "rollback"

class RegressionType(str, enum.Enum):
    """Types of quality regressions"""
    QUALITY_DROP = "quality_drop"
    LATENCY_INCREASE = "latency_increase"
    HALLUCINATION_SPIKE = "hallucination_spike"
    TOXICITY_INCREASE = "toxicity_increase"

class RegressionSeverity(str, enum.Enum):
    """Severity levels for regressions"""
    LOW = "low"
    MEDIUM = "medium"
    HIGH = "high"
    CRITICAL = "critical"
```

```
class RegressionAction(str, enum.Enum):
    """Actions taken for regressions"""
    ALERT_ONLY = "alert_only"
    ROLLBACK_TRIGGERED = "rollback_triggered"
    MANUAL REVIEW = "manual_review"
    AUTO_FIXED = "auto_fixed"

# =====
# MODELS
# =====

class QualityMetric(Base):
    """
    Tracks per-request quality metrics for LLM outputs

    Stores detailed quality scores for each LLM request to enable
    quality monitoring, regression detection, and A/B testing.
    """

    __tablename__ = "quality_metrics"

    # Primary Key
    id = Column(
        UUID(as_uuid=True),
        primary_key=True,
        default=uuid.uuid4,
        index=True
    )

    # Foreign Keys
    agent_id = Column(
        UUID(as_uuid=True),
        ForeignKey("agents.id", ondelete="CASCADE"),
        nullable=False,
        index=True
    )
    customer_id = Column(
        UUID(as_uuid=True),
        ForeignKey("customers.id", ondelete="CASCADE"),
        nullable=False,
        index=True
    )
```

```
# Request Details
request_id = Column(
    String(255),
    nullable=False,
    unique=True,
    index=True,
    comment="Unique identifier for the LLM request"
)
model_name = Column(
    String(100),
    nullable=False,
    index=True,
    comment="Model name: gpt-4, claude-3-opus, llama-3-70b, etc."
)
model_version = Column(
    String(50),
    nullable=False,
    comment="Model version: 0314, 20240229, etc."
)
```

```
# Token Usage
prompt_tokens = Column(
    Integer,
    nullable=True,
    comment="Number of tokens in prompt"
)
completion_tokens = Column(
    Integer,
    nullable=True,
    comment="Number of tokens in completion"
)
```

```
# Performance
latency_ms = Column(
    Float,
    nullable=True,
    comment="Request latency in milliseconds"
)
```

```
# Quality Scores (0.0 to 1.0)
relevance_score = Column(
    Float,
    nullable=True,
```

```
    comment="How relevant is the response to the prompt (0-1)"  
)  
coherence_score = Column(  
    Float,  
    nullable=True,  
    comment="How coherent and logical is the response (0-1)"  
)  
factuality_score = Column(  
    Float,  
    nullable=True,  
    comment="How factually accurate is the response (0-1)"  
)  
hallucination_detected = Column(  
    Boolean,  
    nullable=False,  
    default=False,  
    index=True,  
    comment="Whether hallucination was detected"  
)  
toxicity_score = Column(  
    Float,  
    nullable=True,  
    comment="Toxicity level of the response (0-1, lower is better)"  
)  
overall_quality_score = Column(  
    Float,  
    nullable=False,  
    index=True,  
    comment="Aggregated quality score (0-1)"  
)
```

```
# Timing  
timestamp = Column(  
    DateTime,  
    nullable=False,  
    index=True,  
    comment="When the request was made"  
)
```

```
# Additional Data  
metadata = Column(  
    JSONB,  
    nullable=False,  
    default=dict,
```

```

comment="Additional context: user_id, session_id, etc."
)

# Audit
created_at = Column(DateTime, default=datetime.utcnow, nullable=False)

# Relationships
agent = relationship("Agent", back_populates="quality_metrics")
customer = relationship("Customer", back_populates="quality_metrics")

# Indexes
__table_args__ = (
    Index(
        "ix_quality_metrics_agent_model",
        "agent_id", "model_name"
    ),
    Index(
        "ix_quality_metrics_customer_model",
        "customer_id", "model_name"
    ),
    Index(
        "ix_quality_metrics_timestamp",
        "timestamp"
    ),
    Index(
        "ix_quality_metrics_overall_score",
        "overall_quality_score"
    ),
    Index(
        "ix_quality_metrics_model_timestamp",
        "model_name", "timestamp"
    ),
)
)

def __repr__(self):
    return (
        f"<QualityMetric(id={self.id}, "
        f'model={self.model_name}, score={self.overall_quality_score:.2f})>"
    )
}

class QualityBaseline(Base):
    """
    Stores quality baselines for comparison and regression detection

```

Baselines are established from sample sets and used to detect when quality degrades after changes.

"""

```
__tablename__ = "quality_baselines"
```

Primary Key

```
id = Column(  
    UUID(as_uuid=True),  
    primary_key=True,  
    default=uuid.uuid4,  
    index=True  
)
```

Foreign Keys

```
agent_id = Column(  
    UUID(as_uuid=True),  
    ForeignKey("agents.id", ondelete="CASCADE"),  
    nullable=False,  
    index=True  
)  
  
customer_id = Column(  
    UUID(as_uuid=True),  
    ForeignKey("customers.id", ondelete="CASCADE"),  
    nullable=False,  
    index=True  
)
```

Model Details

```
model_name = Column(  
    String(100),  
    nullable=False,  
    index=True  
)  
  
model_version = Column(  
    String(50),  
    nullable=False  
)
```

Baseline Type

```
baseline_type = Column(  
    SQLEnum(BaselineType, name="baseline_type"),  
    nullable=False,  
    default=BaselineType.INITIAL,
```

```
index=True
)

# Sample Info
sample_size = Column(
    Integer,
    nullable=False,
    comment="Number of requests used to establish baseline"
)

# Average Scores
avg_relevance_score = Column(
    Float,
    nullable=True,
    comment="Average relevance score from samples"
)
avg_coherence_score = Column(
    Float,
    nullable=True,
    comment="Average coherence score from samples"
)
avg_factuality_score = Column(
    Float,
    nullable=True,
    comment="Average factuality score from samples"
)
avg_overall_score = Column(
    Float,
    nullable=False,
    comment="Average overall quality score"
)

# Performance Baseline
p95_latency_ms = Column(
    Float,
    nullable=True,
    comment="P95 latency from samples"
)

# Validity
established_at = Column(
    DateTime,
    nullable=False,
    index=True,
```

```

    comment="When the baseline was established"
)
valid_until = Column(
    DateTime,
    nullable=True,
    comment="When the baseline expires (null = indefinite)"
)

# Additional Data
metadata = Column(
    JSONB,
    nullable=False,
    default=dict,
    comment="Additional context: confidence_interval, std_dev, etc."
)

# Audit
created_at = Column(DateTime, default=datetime.utcnow, nullable=False)

# Relationships
agent = relationship("Agent", back_populates="quality_baselines")
customer = relationship("Customer", back_populates="quality_baselines")
regressions = relationship(
    "QualityRegression",
    back_populates="baseline",
    cascade="all, delete-orphan"
)

# Indexes
__table_args__ = (
    Index(
        "ix_quality_baselines_agent_model",
        "agent_id", "model_name"
    ),
    Index(
        "ix_quality_baselines_customer_model",
        "customer_id", "model_name"
    ),
    Index(
        "ix_quality_baselines_established_at",
        "established_at"
    ),
    Index(
        "ix_quality_baselines_type",
        "type"
    )
)
```

```
        "baseline_type"
    ),
)

def __repr__(self):
    return (
        f"<QualityBaseline(id={self.id}, "
        f'model={self.model_name}, type={self.baseline_type}, '
        f'score={self.avg_overall_score:.2f})>"
    )

class QualityRegression(Base):
    """
    Tracks detected quality regressions

    Records when quality drops below acceptable thresholds,
    enabling automated rollbacks and alerts.
    """

    __tablename__ = "quality_regressions"

    # Primary Key
    id = Column(
        UUID(as_uuid=True),
        primary_key=True,
        default=uuid.uuid4,
        index=True
    )

    # Foreign Keys
    agent_id = Column(
        UUID(as_uuid=True),
        ForeignKey("agents.id", ondelete="CASCADE"),
        nullable=False,
        index=True
    )
    customer_id = Column(
        UUID(as_uuid=True),
        ForeignKey("customers.id", ondelete="CASCADE"),
        nullable=False,
        index=True
    )
    baseline_id = Column(
        UUID(as_uuid=True),
```

```
    ForeignKey("quality_baseline.id", onDelete="CASCADE"),
    nullable=False,
    index=True,
    comment="Baseline used for comparison"
)
workflow_execution_id = Column(
    UUID(as_uuid=True),
    ForeignKey("workflow_executions.id", onDelete="SET NULL"),
    nullable=True,
    index=True,
    comment="Workflow that caused the regression (if known)"
)
```

Regression Details

```
regression_type = Column(
    SQLEnum(RegressionType, name="regression_type"),
    nullable=False,
    index=True
)
severity = Column(
    SQLEnum(RegressionSeverity, name="regression_severity"),
    nullable=False,
    index=True
)
```

Detection Time

```
detected_at = Column(
    DateTime,
    nullable=False,
    index=True
)
```

Metric Details

```
metric_name = Column(
    String(100),
    nullable=False,
    comment="Which metric regressed: overall_score, relevance, latency, etc."
)
baseline_value = Column(
    Float,
    nullable=False,
    comment="Expected value from baseline"
)
current_value = Column(
```

```
        Float,  
        nullable=False,  
        comment="Actual value that triggered regression"  
)  
delta_percent = Column(  
    Float,  
    nullable=False,  
    comment="Percentage change from baseline (negative = worse)"  
)  
sample_size = Column(  
    Integer,  
    nullable=False,  
    comment="Number of samples used to detect regression"  
)
```

Action & Resolution

```
action_taken = Column(  
    SQLEnum(RegressionAction, name="regression_action"),  
    nullable=False,  
    default=RegressionAction.ALERT_ONLY,  
    index=True  
)  
resolved_at = Column(  
    DateTime,  
    nullable=True,  
    comment="When the regression was resolved"  
)  
resolution_notes = Column(  
    Text,  
    nullable=True,  
    comment="How the regression was resolved"  
)
```

Additional Data

```
metadata = Column(  
    JSONB,  
    nullable=False,  
    default=dict,  
    comment="Additional context: affected_requests, p_value, etc."  
)
```

Audit

```
created_at = Column(DateTime, default=datetime.utcnow, nullable=False)  
updated_at = Column(
```

```
DateTime,  
    default=datetime.utcnow,  
    onupdate=datetime.utcnow,  
    nullable=False  
)  
  
# Relationships  
agent = relationship("Agent", back_populates="quality_regressions")  
customer = relationship("Customer", back_populates="quality_regressions")  
baseline = relationship("QualityBaseline", back_populates="regressions")  
workflow = relationship(  
    "WorkflowExecution",  
    foreign_keys=[workflow_execution_id],  
    backref="quality_regressions"  
)  
  
# Indexes  
__table_args__ = (  
    Index(  
        "ix_quality_regressions_agent_severity",  
        "agent_id", "severity"  
    ),  
    Index(  
        "ix_quality_regressions_customer_type",  
        "customer_id", "regression_type"  
    ),  
    Index(  
        "ix_quality_regressions_detected_at",  
        "detected_at"  
    ),  
    Index(  
        "ix_quality_regressions_baseline",  
        "baseline_id"  
    ),  
    Index(  
        "ix_quality_regressions_workflow",  
        "workflow_execution_id"  
    ),  
    Index(  
        "ix_quality_regressions_unresolved",  
        "resolved_at"  
    ),  
)
```

```
def __repr__(self):
    return (
        f"<QualityRegression(id={self.id}, "
        f"type={self.regression_type}, severity={self.severity}, "
        f'resolved={self.resolved_at is not None})>"
    )

@property
def is_resolved(self):
    """Check if regression has been resolved"""
    return self.resolved_at is not None

@property
def time_to_resolve_minutes(self):
    """Calculate time to resolution in minutes"""
    if self.resolved_at and self.detected_at:
        return (self.resolved_at - self.detected_at).total_seconds() / 60
    return None
```

📁 FILE 2: Update Core Models (Add Relationships)

Location: `~/optiinfra/shared/database/models/core.py`

Add these relationships to existing models:

```
python
```

```
# In the Agent model, add:  
from sqlalchemy.orm import relationship  
  
class Agent(Base):  
    # ... existing fields ...  
  
    # ADD THESE NEW RELATIONSHIPS:  
    quality_metrics = relationship(  
        "QualityMetric",  
        back_populates="agent",  
        cascade="all, delete-orphan"  
    )  
    quality_baselines = relationship(  
        "QualityBaseline",  
        back_populates="agent",  
        cascade="all, delete-orphan"  
    )  
    quality_regressions = relationship(  
        "QualityRegression",  
        back_populates="agent",  
        cascade="all, delete-orphan"  
    )
```

```
# In the Customer model, add:  
class Customer(Base):  
    # ... existing fields ...  
  
    # ADD THESE NEW RELATIONSHIPS:  
    quality_metrics = relationship(  
        "QualityMetric",  
        back_populates="customer",  
        cascade="all, delete-orphan"  
    )  
    quality_baselines = relationship(  
        "QualityBaseline",  
        back_populates="customer",  
        cascade="all, delete-orphan"  
    )  
    quality_regressions = relationship(  
        "QualityRegression",  
        back_populates="customer",  
        cascade="all, delete-orphan"  
    )
```

```
    cascade="all, delete-orphan"
```

```
)
```

FILE 3: Update Model Imports

Location: `~/optiinfra/shared/database/models/_init__.py`

```
python
```

....

Database models package

....

Core models (0.2a)

```
from shared.database.models.core import (
```

Customer,

Agent,

Event,

Recommendation,

Approval,

Optimization,

Enums

CustomerStatus,

AgentType,

AgentStatus,

EventType,

EventSeverity,

RecommendationType,

ApprovalStatus,

OptimizationStatus,

)

Agent state models (0.2b)

```
from shared.database.models.agent_state import (
```

AgentConfig,

AgentState,

AgentCapability,

AgentMetric,

Enums

ConfigType,

AgentStatusDetail,

MetricType,

)

Workflow history models (0.2c)

```
from shared.database.models.workflow_history import (
```

WorkflowExecution,

WorkflowStep,

WorkflowArtifact,

Enums

WorkflowType,

WorkflowStatus,

```
StepStatus,  
ArtifactType,  
)  
  
# Resource schema models (0.2d)  
from shared.database.models.resource_schema import (  
    ResourceMetric,  
    ScalingEvent,  
    # Enums  
    ResourceType,  
    ScalingEventType,  
)
```

```
# Application schema models (0.2e) - NEW!  
from shared.database.models.application_schema import (  
    QualityMetric,  
    QualityBaseline,  
    QualityRegression,  
    # Enums  
    BaselineType,  
    RegressionType,  
    RegressionSeverity,  
    RegressionAction,  
)
```

```
__all__ = [  
    # Core models  
    "Customer",  
    "Agent",  
    "Event",  
    "Recommendation",  
    "Approval",  
    "Optimization",  
    # Core enums  
    "CustomerStatus",  
    "AgentType",  
    "AgentStatus",  
    "EventType",  
    "EventSeverity",  
    "RecommendationType",  
    "ApprovalStatus",  
    "OptimizationStatus",  
    # Agent state models  
    "AgentConfig",
```

```
"AgentState",
"AgentCapability",
"AgentMetric",
# Agent state enums
"ConfigType",
"AgentStatusDetail",
"MetricType",
# Workflow history models
"WorkflowExecution",
"WorkflowStep",
"WorkflowArtifact",
# Workflow history enums
"WorkflowType",
"WorkflowStatus",
"StepStatus",
"ArtifactType",
# Resource schema models
"ResourceMetric",
"ScalingEvent",
# Resource schema enums
"ResourceType",
"ScalingEventType",
# Application schema models - NEW!
"QualityMetric",
"QualityBaseline",
"QualityRegression",
# Application schema enums - NEW!
"BaselineType",
"RegressionType",
"RegressionSeverity",
"RegressionAction",
]
```

📁 FILE 4: Alembic Migration

Location: `(~/optiinfra/shared/database/migrations/versions/005_application_schema.py)`

```
python
```

```
"""Create application schema tables
```

Revision ID: 005_application_schema

Revises: 004_resource_schema

Create Date: 2025-10-20 18:00:00.000000

```
"""
```

```
from alembic import op
import sqlalchemy as sa
from sqlalchemy.dialects import postgresql
```

revision identifiers, used by Alembic.

```
revision = '005_application_schema'
down_revision = '004_resource_schema'
branch_labels = None
depends_on = None
```

```
def upgrade() -> None:
```

```
"""Create application schema tables"""

# Create ENUM types
```

```
baseline_type_enum = postgresql.ENUM(
    'initial',
    'updated',
    'rollback',
    name='baseline_type',
    create_type=True
)
baseline_type_enum.create(op.get_bind(), checkfirst=True)
```

```
regression_type_enum = postgresql.ENUM(
    'quality_drop',
    'latency_increase',
    'hallucination_spike',
    'toxicity_increase',
    name='regression_type',
    create_type=True
)
regression_type_enum.create(op.get_bind(), checkfirst=True)
```

```
regression_severity_enum = postgresql.ENUM(
    'low',
```

```
'medium',
'high',
'critical',
name='regression_severity',
create_type=True
)
regression_severity_enum.create(op.get_bind(), checkfirst=True)

regression_action_enum = postgresql.ENUM(
    'alert_only',
    'rollback_triggered',
    'manual_review',
    'auto_fixed',
    name='regression_action',
    create_type=True
)
regression_action_enum.create(op.get_bind(), checkfirst=True)

# Create quality_metrics table
op.create_table(
    'quality_metrics',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('customer_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('request_id', sa.String(255), nullable=False, unique=True),
    sa.Column('model_name', sa.String(100), nullable=False),
    sa.Column('model_version', sa.String(50), nullable=False),
    sa.Column('prompt_tokens', sa.Integer(), nullable=True),
    sa.Column('completion_tokens', sa.Integer(), nullable=True),
    sa.Column('latency_ms', sa.Float(), nullable=True),
    sa.Column('relevance_score', sa.Float(), nullable=True),
    sa.Column('coherence_score', sa.Float(), nullable=True),
    sa.Column('factuality_score', sa.Float(), nullable=True),
    sa.Column('hallucination_detected', sa.Boolean(), nullable=False, server_default='false'),
    sa.Column('toxicity_score', sa.Float(), nullable=True),
    sa.Column('overall_quality_score', sa.Float(), nullable=False),
    sa.Column('timestamp', sa.DateTime(), nullable=False),
    sa.Column('metadata', postgresql.JSONB(), nullable=False, server_default='{}'),
    sa.Column('created_at', sa.DateTime(), nullable=False),
    sa.ForeignKeyConstraint(['agent_id'], ['agents.id'], onDelete='CASCADE'),
    sa.ForeignKeyConstraint(['customer_id'], ['customers.id'], onDelete='CASCADE')
)

# Create indexes for quality_baselines
```

```

op.create_index('ix_quality_baselines_id', 'quality_baselines', ['id'])
op.create_index('ix_quality_baselines_agent_id', 'quality_baselines', ['agent_id'])
op.create_index('ix_quality_baselines_customer_id', 'quality_baselines', ['customer_id'])
op.create_index('ix_quality_baselines_model_name', 'quality_baselines', ['model_name'])
op.create_index('ix_quality_baselines_baseline_type', 'quality_baselines', ['baseline_type'])
op.create_index('ix_quality_baselines_established_at', 'quality_baselines', ['established_at'])
op.create_index('ix_quality_baselines_agent_model', 'quality_baselines', ['agent_id', 'model_name'])
op.create_index('ix_quality_baselines_customer_model', 'quality_baselines', ['customer_id', 'model_name'])
op.create_index('ix_quality_baselines_type', 'quality_baselines', ['baseline_type'])

# Create quality_regressions table
op.create_table(
    'quality_regressions',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('customer_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('baseline_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('workflow_execution_id', postgresql.UUID(as_uuid=True), nullable=True),
    sa.Column('regression_type', regression_type_enum, nullable=False),
    sa.Column('severity', regression_severity_enum, nullable=False),
    sa.Column('detected_at', sa.DateTime(), nullable=False),
    sa.Column('metric_name', sa.String(100), nullable=False),
    sa.Column('baseline_value', sa.Float(), nullable=False),
    sa.Column('current_value', sa.Float(), nullable=False),
    sa.Column('delta_percent', sa.Float(), nullable=False),
    sa.Column('sample_size', sa.Integer(), nullable=False),
    sa.Column('action_taken', regression_action_enum, nullable=False),
    sa.Column('resolved_at', sa.DateTime(), nullable=True),
    sa.Column('resolution_notes', sa.Text(), nullable=True),
    sa.Column('metadata', postgresql.JSONB(), nullable=False, server_default='{}'),
    sa.Column('created_at', sa.DateTime(), nullable=False),
    sa.Column('updated_at', sa.DateTime(), nullable=False),
    sa.ForeignKeyConstraint(['agent_id'], ['agents.id'], ondelete='CASCADE'),
    sa.ForeignKeyConstraint(['customer_id'], ['customers.id'], ondelete='CASCADE'),
    sa.ForeignKeyConstraint(['baseline_id'], ['quality_baselines.id'], ondelete='CASCADE'),
    sa.ForeignKeyConstraint(['workflow_execution_id'], ['workflow_executions.id'], ondelete='SET NULL'),
)
# Create indexes for quality_regressions
op.create_index('ix_quality_regressions_id', 'quality_regressions', ['id'])
op.create_index('ix_quality_regressions_agent_id', 'quality_regressions', ['agent_id'])
op.create_index('ix_quality_regressions_customer_id', 'quality_regressions', ['customer_id'])
op.create_index('ix_quality_regressions_baseline_id', 'quality_regressions', ['baseline_id'])
op.create_index('ix_quality_regressions_workflow_execution_id', 'quality_regressions', ['workflow_execution_id'])

```

```
op.create_index('ix_quality_regressions_regression_type', 'quality_regressions', ['regression_type'])
op.create_index('ix_quality_regressions_severity', 'quality_regressions', ['severity'])
op.create_index('ix_quality_regressions_action_taken', 'quality_regressions', ['action_taken'])
op.create_index('ix_quality_regressions_detected_at', 'quality_regressions', ['detected_at'])
op.create_index('ix_quality_regressions_agent_severity', 'quality_regressions', ['agent_id', 'severity'])
op.create_index('ix_quality_regressions_customer_type', 'quality_regressions', ['customer_id', 'regression_type'])
op.create_index('ix_quality_regressions_baseline', 'quality_regressions', ['baseline_id'])
op.create_index('ix_quality_regressions_workflow', 'quality_regressions', ['workflow_execution_id'])
op.create_index('ix_quality_regressions_unresolved', 'quality_regressions', ['resolved_at'])
```

```
def downgrade() -> None:
    """Drop application schema tables"""

    # Drop quality_regressions table and indexes
```

```
op.drop_index('ix_quality_regressions_unresolved', 'quality_regressions')
op.drop_index('ix_quality_regressions_workflow', 'quality_regressions')
op.drop_index('ix_quality_regressions_baseline', 'quality_regressions')
op.drop_index('ix_quality_regressions_customer_type', 'quality_regressions')
op.drop_index('ix_quality_regressions_agent_severity', 'quality_regressions')
op.drop_index('ix_quality_regressions_detected_at', 'quality_regressions')
op.drop_index('ix_quality_regressions_action_taken', 'quality_regressions')
op.drop_index('ix_quality_regressions_severity', 'quality_regressions')
op.drop_index('ix_quality_regressions_regression_type', 'quality_regressions')
op.drop_index('ix_quality_regressions_workflow_execution_id', 'quality_regressions')
op.drop_index('ix_quality_regressions_baseline_id', 'quality_regressions')
op.drop_index('ix_quality_regressions_customer_id', 'quality_regressions')
op.drop_index('ix_quality_regressions_agent_id', 'quality_regressions')
op.drop_index('ix_quality_regressions_id', 'quality_regressions')
op.drop_table('quality_regressions')
```

```
# Drop quality_baselines table and indexes
op.drop_index('ix_quality_baselines_type', 'quality_baselines')
op.drop_index('ix_quality_baselines_customer_model', 'quality_baselines')
op.drop_index('ix_quality_baselines_agent_model', 'quality_baselines')
op.drop_index('ix_quality_baselines_established_at', 'quality_baselines')
op.drop_index('ix_quality_baselines_baseline_type', 'quality_baselines')
op.drop_index('ix_quality_baselines_model_name', 'quality_baselines')
op.drop_index('ix_quality_baselines_customer_id', 'quality_baselines')
op.drop_index('ix_quality_baselines_agent_id', 'quality_baselines')
op.drop_index('ix_quality_baselines_id', 'quality_baselines')
op.drop_table('quality_baselines')
```

```
# Drop quality_metrics table and indexes
```

```
op.drop_index('ix_quality_metrics_model_timestamp', 'quality_metrics')
op.drop_index('ix_quality_metrics_customer_model', 'quality_metrics')
op.drop_index('ix_quality_metrics_agent_model', 'quality_metrics')
op.drop_index('ix_quality_metrics_timestamp', 'quality_metrics')
op.drop_index('ix_quality_metrics_overall_quality_score', 'quality_metrics')
op.drop_index('ix_quality_metrics_hallucination_detected', 'quality_metrics')
op.drop_index('ix_quality_metrics_model_name', 'quality_metrics')
op.drop_index('ix_quality_metrics_request_id', 'quality_metrics')
op.drop_index('ix_quality_metrics_customer_id', 'quality_metrics')
op.drop_index('ix_quality_metrics_agent_id', 'quality_metrics')
op.drop_index('ix_quality_metrics_id', 'quality_metrics')
op.drop_table('quality_metrics')
```

Drop ENUM types

```
sa.Enum(name='regression_action').drop(op.get_bind(), checkfirst=True)
sa.Enum(name='regression_severity').drop(op.get_bind(), checkfirst=True)
sa.Enum(name='regression_type').drop(op.get_bind(), checkfirst=True)
sa.Enum(name='baseline_type').drop(op.get_bind(), checkfirst=True)
```

EXECUTION STEPS

Step 1: Create Models File

```
bash
cd ~/optiinfra/shared/database/models

# Create application_schema.py
cat > application_schema.py << EOF
[Copy the entire application_schema.py content from FILE 1 above]
EOF

# Verify file created
ls -lh application_schema.py
# Expected: ~550 lines
```

Step 2: Update Core Models

```
bash
```

```
cd ~/optiinfra/shared/database/models
```

```
# Open core.py and add relationships
```

```
nano core.py
```

```
# Add to Agent class:
```

```
# quality_metrics = relationship(...)
```

```
# quality_baselines = relationship(...)
```

```
# quality_regressions = relationship(...)
```

```
#
```

```
# Add to Customer class:
```

```
# quality_metrics = relationship(...)
```

```
# quality_baselines = relationship(...)
```

```
# quality_regressions = relationship(...)
```

Step 3: Update Model Imports

```
bash
```

```
cd ~/optiinfra/shared/database/models
```

```
# Update __init__.py with application models
```

```
cat > __init__.py << 'EOF'
```

```
[Copy the entire __init__.py content from FILE 3 above]
```

```
EOF
```

Step 4: Create Migration File

```
bash
```

```
cd ~/optiinfra/shared/database/migrations/versions
```

```
# Create migration file
```

```
cat > 005_application_schema.py << 'EOF'
```

```
[Copy the entire migration content from FILE 4 above]
```

```
EOF
```

```
# Verify file created
```

```
ls -lh 005_application_schema.py
```

```
# Expected: ~220 lines
```

Step 5: Test Imports

```
bash

cd ~/optiinfra

# Test if new models can be imported
python << EOF
from shared.database.models.application_schema import (
    QualityMetric,
    QualityBaseline,
    QualityRegression,
    BaselineType,
    RegressionType,
    RegressionSeverity,
    RegressionAction
)
print(' ✅ All application models import successfully')

# Test enums
print(f ✅ BaselineType has {len(BaselineType)} values')
print(f ✅ RegressionType has {len(RegressionType)} values')
print(f ✅ RegressionSeverity has {len(RegressionSeverity)} values')
print(f ✅ RegressionAction has {len(RegressionAction)} values')
EOF

# Expected output:
# ✅ All application models import successfully
# ✅ BaselineType has 3 values
# ✅ RegressionType has 4 values
# ✅ RegressionSeverity has 4 values
# ✅ RegressionAction has 4 values
```

Step 6: Run Alembic Migration

```
bash
```

```
cd ~/optiinfra/shared/database
```

```
# Check current migration state
```

```
alembic current
```

```
# Expected: 004_resource_schema (current)
```

```
# Run the new migration
```

```
alembic upgrade head
```

```
# Expected output:
```

```
# INFO [alembic.runtime.migration] Running upgrade 004_resource_schema -> 005_application_schema, Create applicatio
```

Step 7: Verify Tables Created

```
bash
```

```
# Connect to PostgreSQL
psql postgresql://optiinfra:password@localhost:5432/optiinfra

# List all tables (should now be 18 tables)
\dt

# Expected output (15 from previous + 3 new):
# quality_metrics
# quality_baselines
# quality_regressions
# resource_metrics
# scaling_events
# workflow_executions
# workflow_steps
# workflow_artifacts
# agent_configs
# agent_states
# agent_capabilities
# agent_metrics
# customers
# agents
# events
# recommendations
# approvals
# optimizations

# Describe quality_metrics table
\d quality_metrics

# Expected: Should show all columns, indexes, foreign keys

# Describe quality_baselines table
\d quality_baselines

# Expected: Should show all columns, relationships

# Describe quality_regressions table
\d quality_regressions

# Expected: Should show all columns, foreign keys to baselines and workflows

# Check new enums were created
\dT+
```

```
# Expected: Should see baseline_type, regression_type, regression_severity, regression_action
```

```
# Exit
```

```
\q
```

Step 8: Verify Relationships Work

```
bash
```

cd ~/optiinfra

```
# Test relationships in Python
python << 'EOF'
from shared.database.models import Agent, Customer, QualityMetric, QualityBaseline, QualityRegression
from shared.database.session import SessionLocal

db = SessionLocal()

# Get a test agent
agent = db.query(Agent).first()
print(f" ✅ Agent: {agent.name}")
print(f" - Has quality_metrics relationship: {hasattr(agent, 'quality_metrics')}")
print(f" - Has quality_baselines relationship: {hasattr(agent, 'quality_baselines')}")
print(f" - Has quality_regressions relationship: {hasattr(agent, 'quality_regressions')")

# Get a test customer
customer = db.query(Customer).first()
print(f" ✅ Customer: {customer.name}")
print(f" - Has quality_metrics relationship: {hasattr(customer, 'quality_metrics')}")
print(f" - Has quality_baselines relationship: {hasattr(customer, 'quality_baselines')}")
print(f" - Has quality_regressions relationship: {hasattr(customer, 'quality_regressions')")

# Check models
print(f" ✅ QualityMetric model loaded")
print(f" - Has agent relationship: {hasattr(QualityMetric, 'agent')}")
print(f" - Has customer relationship: {hasattr(QualityMetric, 'customer')}")

print(f" ✅ QualityBaseline model loaded")
print(f" - Has agent relationship: {hasattr(QualityBaseline, 'agent')}")
print(f" - Has customer relationship: {hasattr(QualityBaseline, 'customer')}")
print(f" - Has regressions relationship: {hasattr(QualityBaseline, 'regressions')}")

print(f" ✅ QualityRegression model loaded")
print(f" - Has agent relationship: {hasattr(QualityRegression, 'agent')}")
print(f" - Has baseline relationship: {hasattr(QualityRegression, 'baseline')}")
print(f" - Has workflow relationship: {hasattr(QualityRegression, 'workflow')}")

db.close()
print("\n ✅ All relationship verifications passed!")
EOF
```

Expected output similar to previous verifications

🎯 PART 1 SUMMARY

Files Created/Modified:

- ✓ **shared/database/models/application_schema.py** - 3 new models + 4 enums (~550 lines)
- ✓ **shared/database/models/core.py** - Modified (added 6 relationships)
- ✓ **shared/database/models/init.py** - Modified (exports)
- ✓ **shared/database/migrations/versions/005_application_schema.py** - Migration (~220 lines)

What Works Now:

- ✓ Can import all new models in Python
- ✓ Can run Alembic migration
- ✓ 3 new tables created in PostgreSQL
- ✓ All foreign keys working (cascades configured)
- ✓ All indexes created (35+ indexes)
- ✓ Agent-quality relationship working
- ✓ Customer-quality relationship working
- ✓ Regression-baseline-workflow relationships working

Database Status:

Total Tables: 18 (100% PostgreSQL schema COMPLETE!) 🎉

- Core (0.2a): 6 tables ✓
- Agent State (0.2b): 4 tables ✓
- Workflow History (0.2c): 3 tables ✓
- Resource Schema (0.2d): 2 tables ✓
- Application Schema (0.2e): 3 tables ✓

Total Enums: 21

Total Indexes: 115+

Total Relationships: 35+

🎉 MILESTONE ACHIEVED:

PostgreSQL schema is 100% COMPLETE!

All 4 agents now have complete database support:

- Cost Agent (core tables + metrics)
- Performance Agent (core tables + metrics)
- Resource Agent (resource_metrics, scaling_events)
- Application Agent (quality_metrics, baselines, regressions)

What's in Part 2:

- Seed data (quality metrics, baselines, regressions)
- Test fixtures (sample quality data, regression scenarios)
- Test cases (14-16 comprehensive tests)
- Integration tests (complete regression detection flow)
- All validation commands
- Troubleshooting guide
- Success criteria checklist
- Git commit instructions
- **Celebration of PostgreSQL completion!** 🎉

▶ NEXT STEP

Download PART 2 to get:

- Complete seed data implementation (realistic quality metrics)
- Full test suite (14-16 tests covering all scenarios)
- Step-by-step validation
- Troubleshooting for common issues
- Final checklist
- **PostgreSQL 100% completion celebration!**

File name: FOUNDATION-0.2e-Application-Schema-PART2-Testing.md

⬇ HOW TO DOWNLOAD THIS FILE

1. Click the "Copy" dropdown button at the top of this artifact

2. Select "Download as txt"

3. Save as: FOUNDATION-0.2e-Application-Schema-PART1-Code.md

VERIFICATION CHECKLIST (Part 1)

Before moving to Part 2, verify:

- application_schema.py created with 3 models
- core.py updated with 6 new relationships
- init_.py exports all application models
- Migration file 005_application_schema.py created
- Migration runs successfully (`alembic upgrade head`)
- 3 new tables exist in PostgreSQL
- 4 new enum types created
- All indexes created
- Python imports work (no errors)
- Relationships work (verified in Python)

If all checked, you're ready for PART 2! 

 PART 1 COMPLETE! Ready for PART 2 and PostgreSQL completion?  lete='CASCADE'),)

```

# Create indexes for quality_metrics
op.create_index('ix_quality_metrics_id', 'quality_metrics', ['id'])
op.create_index('ix_quality_metrics_agent_id', 'quality_metrics', ['agent_id'])
op.create_index('ix_quality_metrics_customer_id', 'quality_metrics', ['customer_id'])
op.create_index('ix_quality_metrics_request_id', 'quality_metrics', ['request_id'], unique=True)
op.create_index('ix_quality_metrics_model_name', 'quality_metrics', ['model_name'])
op.create_index('ix_quality_metrics_hallucination_detected', 'quality_metrics', ['hallucination_detected'])
op.create_index('ix_quality_metrics_overall_quality_score', 'quality_metrics', ['overall_quality_score'])
op.create_index('ix_quality_metrics_timestamp', 'quality_metrics', ['timestamp'])
op.create_index('ix_quality_metrics_agent_model', 'quality_metrics', ['agent_id', 'model_name'])
op.create_index('ix_quality_metrics_customer_model', 'quality_metrics', ['customer_id', 'model_name'])
op.create_index('ix_quality_metrics_model_timestamp', 'quality_metrics', ['model_name', 'timestamp'])

# Create quality_baselines table
op.create_table(
    'quality_baselines',
    sa.Column('id', postgresql.UUID(as_uuid=True), primary_key=True),
    sa.Column('agent_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('customer_id', postgresql.UUID(as_uuid=True), nullable=False),
    sa.Column('model_name', sa.String(100), nullable=False),
    sa.Column('model_version', sa.String(50), nullable=False),
    sa.Column('baseline_type', baseline_type_enum, nullable=False),
    sa.Column('sample_size', sa.Integer(), nullable=False),
    sa.Column('avg_relevance_score', sa.Float(), nullable=True),
    sa.Column('avg_coherence_score', sa.Float(), nullable=True),
    sa.Column('avg_factuality_score', sa.Float(), nullable=True),
    sa.Column('avg_overall_score', sa.Float(), nullable=False),
    sa.Column('p95_latency_ms', sa.Float(), nullable=True),
    sa.Column('established_at', sa.DateTime(), nullable=False),
    sa.Column('valid_until', sa.DateTime(), nullable=True),
    sa.Column('metadata', postgresql.JSONB(), nullable=False, server_default='{}'),
    sa.Column('created_at', sa.DateTime(), nullable=False),
    sa.ForeignKeyConstraint(['agent_id'], ['agents.id'], ondelete='CASCADE'),
    sa.ForeignKeyConstraint(['customer_id'], ['customers.id'], onde

```