

FOUNDATION-0.3: ClickHouse Time-Series Schema - PART 1 (Code)

🎯 CONTEXT

Phase: FOUNDATION (Week 1 - Day 2 Evening)

Component: ClickHouse Time-Series Database Setup

Estimated Time: 15 min AI execution + 10 min verification

Complexity: MEDIUM

Risk Level: LOW

Files: Part 1 of 2 (Code implementation)

MILESTONE: High-frequency metrics storage for all agents! 

📦 DEPENDENCIES

Must Complete First:

- **P-01:** Bootstrap Project Structure  COMPLETED
- **FOUNDATION-0.2a:** Core Database Schema  COMPLETED
- **FOUNDATION-0.2b:** Agent State Tables  COMPLETED
- **FOUNDATION-0.2c:** Workflow History Tables  COMPLETED
- **FOUNDATION-0.2d:** Resource Schema Tables  COMPLETED
- **FOUNDATION-0.2e:** Application Schema Tables  COMPLETED

Required Services Running:

```
bash

# Verify all services are healthy
cd ~/optiinfra
make verify

# Expected output:
# PostgreSQL...  HEALTHY
# ClickHouse...  HEALTHY
# Qdrant...  HEALTHY
# Redis...  HEALTHY
```

Verify ClickHouse is Accessible:

```
bash  
  
# Test ClickHouse connection  
docker exec optiinfra-clickhouse clickhouse-client --query="SELECT 1"  
  
# Expected output: 1
```

⌚ OBJECTIVE

Set up **ClickHouse time-series database** to store high-frequency metrics that would overwhelm PostgreSQL.

What We're Building:

4 Time-Series Tables (raw metrics):

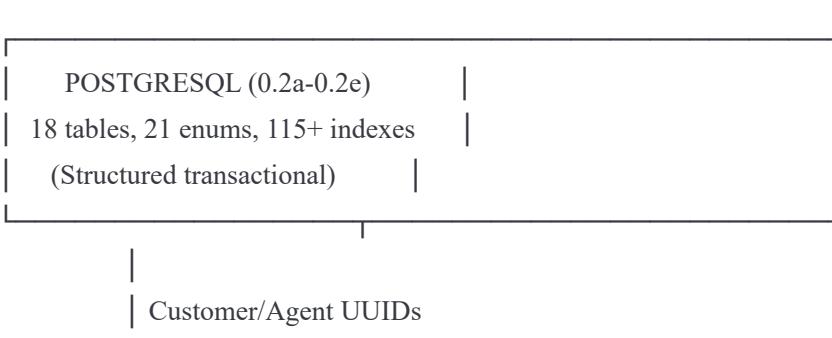
1. **cost_metrics_ts** - Cloud costs per instance per hour (90-day TTL)
2. **performance_metrics_ts** - LLM latency/throughput per request (30-day TTL)
3. **resource_metrics_ts** - GPU/CPU utilization per minute (90-day TTL)
4. **quality_metrics_ts** - Quality scores per request (30-day TTL)

4 Materialized Views (hourly aggregations):

1. **cost_metrics_hourly_mv** - Pre-aggregated hourly costs
2. **performance_metrics_hourly_mv** - P95/P99 latency by hour
3. **resource_metrics_hourly_mv** - GPU utilization by hour
4. **quality_metrics_hourly_mv** - Quality trends by hour

Python Client - Easy-to-use interface for inserts and queries

Database Architecture:



CLICKHOUSE (0.3) <input checked="" type="checkbox"/> NEW	
8 tables (4 base + 4 views)	
(High-frequency time-series)	
Performance:	
- 1M+ inserts/second	
- Millisecond query latency	
- 10-20x compression	
- Automatic TTL cleanup	

Why ClickHouse?

Feature	PostgreSQL	ClickHouse
Time-series inserts	~10K/sec	~1M/sec (100x faster)
Query latency	Seconds	Milliseconds
Storage compression	2-3x	10-20x
Best for	Structured data	Time-series metrics

📁 FILE 1: ClickHouse Initialization Script

Location: `~/optiinfra/shared/clickhouse/migrations/init.sql`

sql

```

-- =====
-- OptiInfra ClickHouse Time-Series Database
-- Foundation Phase 0.3
-- =====

-- Create database
CREATE DATABASE IF NOT EXISTS optiinfra;
USE optiinfra;

-- =====

-- TABLE 1: cost_metrics_ts
-- Purpose: Track cloud costs per instance per hour
-- Retention: 90 days
-- Granularity: 1 hour
-- =====

CREATE TABLE IF NOT EXISTS cost_metrics_ts (
    timestamp DateTime,
    customer_id UUID,
    cloud_provider String,      -- 'aws', 'gcp', 'azure'
    service_name String,        -- 'ec2', 'compute-engine', 'vm'
    instance_id String,
    instance_type String,       -- 'm5.xlarge', 'n1-standard-4'
    region String,              -- 'us-east-1', 'us-central1'
    cost_per_hour Float64,
    utilization_percent Float32,
    is_spot UInt8,              -- 0 or 1 (boolean)
    is_reserved UInt8           -- 0 or 1 (boolean)
)
ENGINE = MergeTree()
PARTITION BY toYYYYMM(timestamp)
ORDER BY (customer_id, cloud_provider, timestamp)
TTL timestamp + INTERVAL 90 DAY
SETTINGS index_granularity = 8192;

-- =====

-- TABLE 2: performance_metrics_ts
-- Purpose: Track LLM inference performance per request
-- Retention: 30 days
-- Granularity: 1 second (per request)
-- =====

CREATE TABLE IF NOT EXISTS performance_metrics_ts (

```

```
timestamp DateTime,  
customer_id UUID,  
service_id UUID,  
service_type String,          -- 'vllm', 'tgi', 'sglang'  
model_name String,           -- 'gpt-4', 'llama-2-70b'  
request_id UUID,  
latency_ms Float32,  
throughput_tokens_per_sec Float32,  
gpu_utilization Float32,  
kv_cache_utilization Float32,  
batch_size UInt32,  
prompt_tokens UInt32,  
completion_tokens UInt32,  
total_tokens UInt32
```

```
)
```

```
ENGINE = MergeTree()
```

```
PARTITION BY toYYYYMM(timestamp)
```

```
ORDER BY (customer_id, service_id, timestamp)
```

```
TTL timestamp + INTERVAL 30 DAY
```

```
SETTINGS index_granularity = 8192;
```

```
-- =====
```

```
-- TABLE 3: resource_metrics_ts
```

```
-- Purpose: Track GPU/CPU utilization per minute
```

```
-- Retention: 90 days
```

```
-- Granularity: 1 minute
```

```
-- =====
```

```
CREATE TABLE IF NOT EXISTS resource_metrics_ts (
```

```
timestamp DateTime,  
customer_id UUID,  
instance_id String,  
instance_type String,  
gpu_index UInt8,           -- 0-7 (which GPU on instance)  
gpu_utilization Float32,   -- 0-100%  
gpu_memory_used_mb Float32,  
gpu_memory_total_mb Float32,  
gpu_temperature Float32,  
cpu_utilization Float32,    -- 0-100%  
memory_used_gb Float32,  
memory_total_gb Float32,  
network_rx_mbps Float32,  
network_tx_mbps Float32
```

```
)
```

```
ENGINE = MergeTree()
PARTITION BY toYYYYMM(timestamp)
ORDER BY (customer_id, instance_id, timestamp)
TTL timestamp + INTERVAL 90 DAY
SETTINGS index_granularity = 8192;
```

```
-- =====
-- TABLE 4: quality_metrics_ts
-- Purpose: Track LLM output quality per request
-- Retention: 30 days
-- Granularity: Per request
-- =====
```

```
CREATE TABLE IF NOT EXISTS quality_metrics_ts (
    timestamp DateTime,
    customer_id UUID,
    service_id UUID,
    request_id UUID,
    model_name String,
    relevance_score Float32,      -- 0-1
    coherence_score Float32,      -- 0-1
    factuality_score Float32,      -- 0-1
    hallucination_detected UInt8,  -- 0 or 1
    toxicity_score Float32,        -- 0-1
    overall_quality_score Float32, -- weighted average
    prompt_hash String,           -- for grouping similar prompts
    latency_ms UInt32
)
```

```
ENGINE = MergeTree()
PARTITION BY toYYYYMM(timestamp)
ORDER BY (customer_id, service_id, timestamp)
TTL timestamp + INTERVAL 30 DAY
SETTINGS index_granularity = 8192;
```

```
-- =====
-- MATERIALIZED VIEW 1: cost_metrics_hourly_mv
-- Purpose: Pre-aggregate cost data by hour for fast dashboard queries
-- =====
```

```
CREATE MATERIALIZED VIEW IF NOT EXISTS cost_metrics_hourly_mv
ENGINE = SummingMergeTree()
PARTITION BY toYYYYMM(hour)
ORDER BY (customer_id, cloud_provider, hour)
POPULATE
```

AS SELECT

```
toStartOfHour(timestamp) as hour,
customer_id,
cloud_provider,
service_name,
sum(cost_per_hour) as total_cost,
avg(utilization_percent) as avg_utilization,
count() as sample_count
FROM cost_metrics_ts
GROUP BY hour, customer_id, cloud_provider, service_name;
```

```
-- =====
-- MATERIALIZED VIEW 2: performance_metrics_hourly_mv
-- Purpose: Pre-aggregate latency P95/P99 by hour
-- =====
```

CREATE MATERIALIZED VIEW IF NOT EXISTS performance_metrics_hourly_mv

```
ENGINE = AggregatingMergeTree()
PARTITION BY toYYYYMM(hour)
ORDER BY (customer_id, service_id, hour)
POPULATE
```

AS SELECT

```
toStartOfHour(timestamp) as hour,
customer_id,
service_id,
service_type,
avgState(latency_ms) as avg_latency,
quantileState(0.95)(latency_ms) as p95_latency,
quantileState(0.99)(latency_ms) as p99_latency,
avgState(throughput_tokens_per_sec) as avg_throughput,
avgState(gpu_utilization) as avg_gpu_util,
count() as request_count
FROM performance_metrics_ts
GROUP BY hour, customer_id, service_id, service_type;
```

```
-- =====
-- MATERIALIZED VIEW 3: resource_metrics_hourly_mv
-- Purpose: Pre-aggregate GPU utilization by hour
-- =====
```

CREATE MATERIALIZED VIEW IF NOT EXISTS resource_metrics_hourly_mv

```
ENGINE = AggregatingMergeTree()
PARTITION BY toYYYYMM(hour)
ORDER BY (customer_id, instance_id, hour)
```

```
POPULATE
AS SELECT
    toStartOfHour(timestamp) as hour,
    customer_id,
    instance_id,
    instance_type,
    avgState(gpu_utilization) as avg_gpu_util,
    maxState(gpu_utilization) as max_gpu_util,
    avgState(gpu_memory_used_mb) as avg_gpu_memory,
    avgState(cpu_utilization) as avg_cpu_util,
    count() as sample_count
FROM resource_metrics_ts
GROUP BY hour, customer_id, instance_id, instance_type;

-- =====
-- MATERIALIZED VIEW 4: quality_metrics_hourly_mv
-- Purpose: Pre-aggregate quality scores by hour
-- =====

CREATE MATERIALIZED VIEW IF NOT EXISTS quality_metrics_hourly_mv
ENGINE = AggregatingMergeTree()
PARTITION BY toYYYYMM(hour)
ORDER BY (customer_id, service_id, hour)
POPULATE
AS SELECT
    toStartOfHour(timestamp) as hour,
    customer_id,
    service_id,
    model_name,
    avgState(overall_quality_score) as avg_quality,
    avgState(relevance_score) as avg_relevance,
    avgState(coherence_score) as avg_coherence,
    avgState(factuality_score) as avg_factuality,
    sumState(hallucination_detected) as hallucination_count,
    count() as request_count
FROM quality_metrics_ts
GROUP BY hour, customer_id, service_id, model_name;

-- =====
-- VERIFICATION QUERIES
-- =====
```

```
-- Show all tables
SHOW TABLES;
```

```
-- Check table row counts
SELECT
    database,
    table,
    formatReadableSize(total_bytes) as size,
    total_rows as rows
FROM system.tables
WHERE database = 'optiinfra'
ORDER BY table;
```

📁 FILE 2: Python ClickHouse Client

Location: `~/optiinfra/shared/clickhouse/client.py`

```
python
```

....

ClickHouse client for high-frequency time-series metrics.

Provides easy-to-use interface for inserting and querying metrics.

Usage:

```
from shared.clickhouse.client import get_clickhouse_client

client = get_clickhouse_client()

# Insert cost metrics
client.insert_cost_metrics([
    {
        'timestamp': datetime.now(),
        'customer_id': '123e4567-e89b-12d3-a456-426614174000',
        'cloud_provider': 'aws',
        'service_name': 'ec2',
        'instance_id': 'i-1234567',
        'instance_type': 'm5.xlarge',
        'region': 'us-east-1',
        'cost_per_hour': 0.192,
        'utilization_percent': 45.5,
        'is_spot': 0,
        'is_reserved': 0
    }
])
```

```
# Query hourly costs
```

```
results = client.query_cost_hourly(
    customer_id='123e4567-e89b-12d3-a456-426614174000',
    start_date=datetime.now() - timedelta(days=7),
    end_date=datetime.now()
)
```

```
from clickhouse_driver import Client
from typing import Dict, List, Any, Optional
import os
from datetime import datetime, timedelta
import logging
```

```
logger = logging.getLogger(__name__)
```

```
class ClickHouseClient:  
    """Client for inserting and querying time-series metrics in ClickHouse."""  
  
    def __init__(self):  
        """Initialize ClickHouse client with connection parameters."""  
        self.client = Client(  
            host=os.getenv('CLICKHOUSE_HOST', 'localhost'),  
            port=int(os.getenv('CLICKHOUSE_PORT', 9000)),  
            database=os.getenv('CLICKHOUSE_DB', 'optiinfra'),  
            user=os.getenv('CLICKHOUSE_USER', 'default'),  
            password=os.getenv('CLICKHOUSE_PASSWORD', '')  
        )  
        logger.info(f"ClickHouse client initialized: {self.client.connection.host}")
```

def ping(self) -> bool:

"""

Check if ClickHouse is accessible.

Returns:

bool: True if ClickHouse responds, False otherwise

"""

try:

result = self.client.execute('SELECT 1')

return result[0][0] == 1

except Exception as e:

logger.error(f"ClickHouse ping failed: {e}")

return False

```
# ======  
# COST METRICS  
# ======
```

```
def insert_cost_metrics(self, metrics: List[Dict[str, Any]]) -> int:  
    """
```

Insert cost metrics in batch.

Args:

metrics: List of cost metric dictionaries with keys:

- timestamp (datetime)
- customer_id (str UUID)
- cloud_provider (str)
- service_name (str)
- instance_id (str)
- instance_type (str)

- region (str)
- cost_per_hour (float)
- utilization_percent (float)
- is_spot (int: 0 or 1)
- is_reserved (int: 0 or 1)

Returns:

int: Number of rows inserted

Example:

```
client.insert_cost_metrics([
    {
        'timestamp': datetime.now(),
        'customer_id': '123e4567-e89b-12d3-a456-426614174000',
        'cloud_provider': 'aws',
        'service_name': 'ec2',
        'instance_id': 'i-1234567',
        'instance_type': 'm5.xlarge',
        'region': 'us-east-1',
        'cost_per_hour': 0.192,
        'utilization_percent': 45.5,
        'is_spot': 0,
        'is_reserved': 0
    }
])
"""

if not metrics:
    return 0

query = """
INSERT INTO cost_metrics_ts
(timestamp, customer_id, cloud_provider, service_name, instance_id,
instance_type, region, cost_per_hour, utilization_percent, is_spot, is_reserved)
VALUES
"""

self.client.execute(query, metrics)
logger.info(f"Inserted {len(metrics)} cost metrics")
return len(metrics)

def query_cost_hourly(
    self,
    customer_id: str,
    start_date: datetime,
```

```
    end_date: datetime,  
    cloud_provider: Optional[str] = None  
) -> List[Dict[str, Any]]:  
    """
```

Query hourly cost aggregations.

Args:

```
    customer_id: Customer UUID  
    start_date: Start datetime  
    end_date: End datetime  
    cloud_provider: Optional filter by cloud provider
```

Returns:

List of dictionaries with hourly cost data

Example:

```
results = client.query_cost_hourly(  
    customer_id='123e4567-e89b-12d3-a456-426614174000',  
    start_date=datetime.now() - timedelta(days=7),  
    end_date=datetime.now(),  
    cloud_provider='aws'  
)  
"""  
query = """  
SELECT  
    hour,  
    cloud_provider,  
    service_name,  
    total_cost,  
    avg_utilization,  
    sample_count  
FROM cost_metrics_hourly_mv  
WHERE customer_id = %(customer_id)s  
    AND hour >= %(start_date)s  
    AND hour < %(end_date)s  
"""  
  
params = {  
    'customer_id': customer_id,  
    'start_date': start_date,  
    'end_date': end_date  
}  
  
if cloud_provider:
```

```
query += " AND cloud_provider = %(cloud_provider)s"
params['cloud_provider'] = cloud_provider

query += " ORDER BY hour"

result = self.client.execute(query, params)

return [
{
    'hour': row[0],
    'cloud_provider': row[1],
    'service_name': row[2],
    'total_cost': row[3],
    'avg_utilization': row[4],
    'sample_count': row[5]
}
for row in result
]
```

```
# =====
```

```
# PERFORMANCE METRICS
```

```
# =====
```

```
def insert_performance_metrics(self, metrics: List[Dict[str, Any]]) -> int:
```

```
"""
```

Insert performance metrics in batch.

Args:

metrics: List of performance metric dictionaries

Returns:

int: Number of rows inserted

```
"""
```

```
if not metrics:
```

```
    return 0
```

```
query = """
```

```
INSERT INTO performance_metrics_ts
(timestamp, customer_id, service_id, service_type, model_name, request_id,
latency_ms, throughput_tokens_per_sec, gpu_utilization, kv_cache_utilization,
batch_size, prompt_tokens, completion_tokens, total_tokens)
```

```
VALUES
```

```
"""
```

```
self.client.execute(query, metrics)
logger.info(f"Inserted {len(metrics)} performance metrics")
return len(metrics)

def query_performance_p95(
    self,
    customer_id: str,
    service_id: str,
    hours: int = 24
) -> Dict[str, float]:
    """
    Query P95 latency over last N hours.

```

Args:

```
customer_id: Customer UUID
service_id: Service UUID
hours: Number of hours to query (default 24)
```

Returns:

Dictionary with performance metrics

Example:

```
results = client.query_performance_p95(
    customer_id='123e4567-e89b-12d3-a456-426614174000',
    service_id='456e7890-e89b-12d3-a456-426614174000',
    hours=24
)
# Returns: {'avg_latency_ms': 245.3, 'p95_latency_ms': 450.2, ...}
"""

end_time = datetime.now()
start_time = end_time - timedelta(hours=hours)

query = """
SELECT
    avgMerge(avg_latency) as avg_latency,
    quantileMerge(0.95)(p95_latency) as p95_latency,
    quantileMerge(0.99)(p99_latency) as p99_latency,
    avgMerge(avg_throughput) as avg_throughput,
    sum(request_count) as total_requests
FROM performance_metrics_hourly_mv
WHERE customer_id = %(customer_id)s
    AND service_id = %(service_id)s
    AND hour >= %(start_time)s
    AND hour < %(end_time)s

```

```
"""
result = self.client.execute(query, {
    'customer_id': customer_id,
    'service_id': service_id,
    'start_time': start_time,
    'end_time': end_time
})

if not result:
    return {}

row = result[0]
return {
    'avg_latency_ms': row[0],
    'p95_latency_ms': row[1],
    'p99_latency_ms': row[2],
    'avg_throughput': row[3],
    'total_requests': row[4]
}

# =====
# RESOURCE METRICS
# =====

def insert_resource_metrics(self, metrics: List[Dict[str, Any]]) -> int:
    """Insert resource metrics in batch."""
    if not metrics:
        return 0

    query = """
INSERT INTO resource_metrics_ts
(timestamp, customer_id, instance_id, instance_type, gpu_index,
gpu_utilization, gpu_memory_used_mb, gpu_memory_total_mb, gpu_temperature,
cpu_utilization, memory_used_gb, memory_total_gb, network_rx_mbps, network_tx_mbps)
VALUES
"""

    self.client.execute(query, metrics)
    logger.info(f'Inserted {len(metrics)} resource metrics')
    return len(metrics)

def query_resource_utilization(
    self,
```

```

customer_id: str,
instance_id: str,
hours: int = 24
) -> Dict[str, float]:
    """Query resource utilization over last N hours."""
    end_time = datetime.now()
    start_time = end_time - timedelta(hours=hours)

query = """
SELECT
    avgMerge(avg_gpu_util) as avg_gpu_util,
    maxMerge(max_gpu_util) as max_gpu_util,
    avgMerge(avg_gpu_memory) as avg_gpu_memory,
    avgMerge(avg_cpu_util) as avg_cpu_util
FROM resource_metrics_hourly_mv
WHERE customer_id = %(customer_id)s
    AND instance_id = %(instance_id)s
    AND hour >= %(start_time)s
    AND hour < %(end_time)s
"""

```

```
result = self.client.execute(query, {
```

```

    'customer_id': customer_id,
    'instance_id': instance_id,
    'start_time': start_time,
    'end_time': end_time
})
```

```
if not result:
```

```
    return {}
```

```
row = result[0]
return {
    'avg_gpu_utilization': row[0],
    'max_gpu_utilization': row[1],
    'avg_gpu_memory_mb': row[2],
    'avg_cpu_utilization': row[3]
}
```

```
# =====
```

```
# QUALITY METRICS
```

```
# =====
```

```
def insert_quality_metrics(self, metrics: List[Dict[str, Any]]) -> int:
```

```
"""Insert quality metrics in batch."""
if not metrics:
    return 0

query = """
INSERT INTO quality_metrics_ts
(timestamp, customer_id, service_id, request_id, model_name,
relevance_score, coherence_score, factuality_score, hallucination_detected,
toxicity_score, overall_quality_score, prompt_hash, latency_ms)
VALUES
"""

self.client.execute(query, metrics)
logger.info(f"Inserted {len(metrics)} quality metrics")
return len(metrics)

def query_quality_trends(
    self,
    customer_id: str,
    service_id: str,
    hours: int = 24
) -> List[Dict[str, Any]]:
    """Query quality trends over last N hours."""
    end_time = datetime.now()
    start_time = end_time - timedelta(hours=hours)

    query = """
SELECT
    hour,
    avgMerge(avg_quality) as avg_quality,
    avgMerge(avg_relevance) as avg_relevance,
    avgMerge(avg_coherence) as avg_coherence,
    avgMerge(avg_factuality) as avg_factuality,
    sum(hallucination_count) as hallucinations,
    sum(request_count) as total_requests
FROM quality_metrics_hourly_mv
WHERE customer_id = %(customer_id)s
    AND service_id = %(service_id)s
    AND hour >= %(start_time)s
    AND hour < %(end_time)s
GROUP BY hour
ORDER BY hour
"""

```

```
        result = self.client.execute(query, {
            'customer_id': customer_id,
            'service_id': service_id,
            'start_time': start_time,
            'end_time': end_time
        })
```

```
    return [
        {
            'hour': row[0],
            'avg_quality': row[1],
            'avg_relevance': row[2],
            'avg_coherence': row[3],
            'avg_factuality': row[4],
            'hallucinations': row[5],
            'total_requests': row[6]
        }
    ]
]
```

```
# =====  
# SINGLETON PATTERN  
# =====
```

```
_clickhouse_client = None
```

```
def get_clickhouse_client() -> ClickHouseClient:
```

```
    """
```

```
    Get singleton ClickHouse client instance.
```

Returns:

ClickHouseClient: Singleton client instance

Example:

```
client = get_clickhouse_client()  
if client.ping():  
    print("ClickHouse is ready!")
```

```
    """
```

```
global _clickhouse_client  
if _clickhouse_client is None:  
    _clickhouse_client = ClickHouseClient()  
return _clickhouse_client
```

FILE 3: Package Initialization

Location: `~/optiinfra/shared/clickhouse/_init_.py`

```
python
"""

ClickHouse time-series database package.

Provides high-performance storage and querying for:
- Cost metrics (hourly cloud spending)
- Performance metrics (per-request LLM latency)
- Resource metrics (GPU/CPU utilization)
- Quality metrics (LLM output quality scores)

Usage:
from shared.clickhouse import get_clickhouse_client

client = get_clickhouse_client()

# Insert metrics
client.insert_cost_metrics(...)

# Query aggregations
results = client.query_cost_hourly(...)

"""

from shared.clickhouse.client import (
    ClickHouseClient,
    get_clickhouse_client
)

__all__ = [
    'ClickHouseClient',
    'get_clickhouse_client'
]
```

FILE 4: Schema Initialization Helper

Location: `~/optiinfra/shared/clickhouse/schemas/_init_.py`

```
python
```

```
"""
```

```
ClickHouse schema initialization helpers.
```

```
"""
```

```
# Empty file for now, but can add schema utilities later
```

📁 FILE 5: Update Requirements

Location: `~/optiinfra/shared/requirements.txt`

```
txt
```

```
# Existing dependencies...
```

```
sqlalchemy==2.0.23
```

```
alembic==1.12.1
```

```
psycopg2-binary==2.9.9
```

```
# ClickHouse driver (ADD THIS)
```

```
clickhouse-driver==0.2.6
```

```
# Other dependencies...
```

📁 FILE 6: README Documentation

Location: `~/optiinfra/shared/clickhouse/README.md`

```
markdown
```

ClickHouse Time-Series Database

High-performance time-series storage for OptiInfra metrics.

Overview

ClickHouse stores high-frequency metrics that would overwhelm PostgreSQL:

- **1M+ inserts/second** vs PostgreSQL's ~10K/second
- **10-20x compression** vs PostgreSQL's 2-3x
- **Millisecond queries** vs PostgreSQL's seconds

Architecture

Tables

1. **cost_metrics_ts** - Cloud costs (1-hour granularity, 90-day retention)
2. **performance_metrics_ts** - LLM latency (per-request, 30-day retention)
3. **resource_metrics_ts** - GPU/CPU utilization (1-minute, 90-day retention)
4. **quality_metrics_ts** - Quality scores (per-request, 30-day retention)

Materialized Views

Pre-aggregated hourly rollups for fast dashboard queries:

- cost_metrics_hourly_mv
- performance_metrics_hourly_mv
- resource_metrics_hourly_mv
- quality_metrics_hourly_mv

Usage

Initialize Database

```
```bash
Run initialization script
docker exec -i optiinfra-clickhouse clickhouse-client < shared/clickhouse/migrations/init.sql
````
```

Python Client

```
```python
from shared.clickhouse import get_clickhouse_client
from datetime import datetime, timedelta

Get client
client = get_clickhouse_client()

Check connection
```
```

```
if client.ping():
    print("✅ ClickHouse connected!")
```

```
# Insert
```