

# FOUNDATION-0.3: ClickHouse Time-Series Schema - PART 2 (Execution & Testing)

## CONTEXT

**Phase:** FOUNDATION (Week 1 - Day 2 Evening)

**Component:** ClickHouse Time-Series Database - Execution & Validation

**Estimated Time:** 10 min execution + 15 min testing

**Complexity:** MEDIUM

**Risk Level:** LOW

**Files:** Part 2 of 2 (Execution, testing, validation)

---

## PREREQUISITES

### Must Have Completed:

-  **PART 1** - All code files created
-  **0.2a-0.2e** - PostgreSQL schemas complete

### Verify PART 1 Files Exist:

```
bash

cd ~/optiinfra

# Check all files created
ls -lh shared/clickhouse/migrations/init.sql
ls -lh shared/clickhouse/client.py
ls -lh shared/clickhouse/__init__.py
ls -lh shared/clickhouse/schemas/__init__.py
ls -lh shared/clickhouse/README.md

# Expected: All files present
```

---

## STEP-BY-STEP EXECUTION

### Step 1: Create Directory Structure

```
bash
```

```
cd ~/optiinfra/shared
```

```
# Create clickhouse directory structure
```

```
mkdir -p clickhouse/schemas
```

```
mkdir -p clickhouse/migrations
```

```
# Verify structure
```

```
tree clickhouse/ -L 2
```

```
# Expected:
```

```
# clickhouse/
```

```
#   └── migrations/
```

```
#   └── schemas/
```

## Validation:

```
bash
```

```
# Check directories exist
```

```
[ -d "clickhouse/migrations" ] && echo "✓ migrations/ created" || echo "✗ migrations/ missing"
```

```
[ -d "clickhouse/schemas" ] && echo "✓ schemas/ created" || echo "✗ schemas/ missing"
```

## Step 2: Create Initialization Script

```
bash
```

```
cd ~/optiinfra/shared/clickhouse/migrations
```

```
# Create init.sql (copy from PART 1, FILE 1)
```

```
cat > init.sql << 'EOF'
```

```
-- =====  
-- OptiInfra ClickHouse Time-Series Database
```

```
-- Foundation Phase 0.3
```

```
-- =====
```

```
-- Create database
```

```
CREATE DATABASE IF NOT EXISTS optiinfra;
```

```
USE optiinfra;
```

```
[... COPY ENTIRE INIT.SQL FROM PART 1, FILE 1 ...]
```

```
EOF
```

```
# Verify file created
```

```
ls -lh init.sql
```

```
wc -l init.sql
```

```
# Expected: ~350 lines
```

## Validation:

```
bash
```

```
# Check file size
```

```
FILE_SIZE=$(wc -l < init.sql)
```

```
if [ "$FILE_SIZE" -gt 300 ]; then
```

```
    echo "✓ init.sql created ($FILE_SIZE lines)"
```

```
else
```

```
    echo "✗ init.sql too small ($FILE_SIZE lines)"
```

```
fi
```

```
# Check for key tables
```

```
grep -q "cost_metrics_ts" init.sql && echo "✓ cost_metrics_ts found" || echo "✗ Missing"
```

```
grep -q "performance_metrics_ts" init.sql && echo "✓ performance_metrics_ts found" || echo "✗ Missing"
```

```
grep -q "resource_metrics_ts" init.sql && echo "✓ resource_metrics_ts found" || echo "✗ Missing"
```

```
grep -q "quality_metrics_ts" init.sql && echo "✓ quality_metrics_ts found" || echo "✗ Missing"
```

## Step 3: Create Python Client

```
bash

cd ~/optiinfra/shared/clickhouse

# Create client.py (copy from PART 1, FILE 2)
cat > client.py << EOF
=====
ClickHouse client for high-frequency time-series metrics.
[... COPY ENTIRE CLIENT.PY FROM PART 1, FILE 2 ...]
=====
EOF

# Verify file
ls -lh client.py
wc -l client.py
# Expected: ~450 lines
```

## Validation:

```
bash

# Check file size
FILE_SIZE=$(wc -l < client.py)
if [ "$FILE_SIZE" -gt 400 ]; then
    echo "✓ client.py created ($FILE_SIZE lines)"
else
    echo "✗ client.py too small ($FILE_SIZE lines)"
fi

# Check for key classes
grep -q "class ClickHouseClient" client.py && echo "✓ ClickHouseClient class found" || echo "✗ Missing"
grep -q "def insert_cost_metrics" client.py && echo "✓ insert_cost_metrics found" || echo "✗ Missing"
grep -q "def query_cost_hourly" client.py && echo "✓ query_cost_hourly found" || echo "✗ Missing"
grep -q "get_clickhouse_client" client.py && echo "✓ get_clickhouse_client found" || echo "✗ Missing"
```

## Step 4: Create Package Files

```
bash
```

```
cd ~/optiinfra/shared/clickhouse
```

```
# Create __init__.py (copy from PART 1, FILE 3)
```

```
cat > __init__.py << 'EOF'
```

```
'''
```

```
ClickHouse time-series database package.
```

```
[... COPY FROM PART 1, FILE 3 ...]
```

```
'''
```

```
EOF
```

```
# Create schemas/__init__.py (copy from PART 1, FILE 4)
```

```
cat > schemas/__init__.py << 'EOF'
```

```
'''
```

```
ClickHouse schema initialization helpers.
```

```
'''
```

```
# Empty file for now, but can add schema utilities later
```

```
EOF
```

```
# Verify files
```

```
ls -lh __init__.py
```

```
ls -lh schemas/__init__.py
```

## Validation:

```
bash
```

```
# Check files exist
```

```
[ -f "__init__.py" ] && echo "✓ __init__.py created" || echo "✗ Missing"
```

```
[ -f "schemas/__init__.py" ] && echo "✓ schemas/__init__.py created" || echo "✗ Missing"
```

## Step 5: Update Requirements

```
bash
```

```
cd ~/optiinfra/shared
```

```
# Add clickhouse-driver to requirements.txt
echo "" >> requirements.txt
echo "# ClickHouse driver (FOUNDATION-0.3)" >> requirements.txt
echo "clickhouse-driver==0.2.6" >> requirements.txt
```

```
# Verify addition
tail -5 requirements.txt
```

## Validation:

```
bash
```

```
# Check if clickhouse-driver is in requirements
grep -q "clickhouse-driver" requirements.txt && echo "✅ clickhouse-driver added" || echo "✗ Missing"
```

## Step 6: Install Dependencies

```
bash
```

```
cd ~/optiinfra/shared
```

```
# Install clickhouse-driver
pip install clickhouse-driver==0.2.6
```

```
# Verify installation
python -c "import clickhouse_driver; print('✅ clickhouse-driver installed successfully')"
# Expected: ✅ clickhouse-driver installed successfully
```

## Validation:

```
bash
```

```
# Test import
python << 'EOF'
try:
    import clickhouse_driver
    print("✓ clickhouse_driver imports successfully")
    print(f" Version: {clickhouse_driver.__version__}")
except ImportError as e:
    print(f"✗ Import failed: {e}")
EOF
```

## Step 7: Create README

```
bash

cd ~/optiinfra/shared/clickhouse

# Create README.md (copy from PART 1, FILE 6)
cat > README.md << 'EOF'
# ClickHouse Time-Series Database
[... COPY ENTIRE README FROM PART 1, FILE 6 ...]
EOF

# Verify file
ls -lh README.md
```

## Validation:

```
bash

# Check README exists and has content
FILE_SIZE=$(wc -l < README.md)
if [ "$FILE_SIZE" -gt 50 ]; then
    echo "✓ README.md created ($FILE_SIZE lines)"
else
    echo "✗ README.md too small"
fi
```

## Step 8: Verify ClickHouse is Running

```
bash
```

```
cd ~/optiinfra
```

```
# Check if ClickHouse container is running
```

```
docker ps | grep clickhouse
```

```
# If not running, start all services
```

```
docker-compose up -d
```

```
# Wait for ClickHouse to be ready (takes ~10 seconds)
```

```
echo "Waiting for ClickHouse to start..."
```

```
sleep 10
```

```
# Test connection
```

```
docker exec optiinfra-clickhouse clickhouse-client --query="SELECT 1"
```

```
# Expected output: 1
```

## Validation:

```
bash
```

```

# Comprehensive ClickHouse check
echo "==== ClickHouse Health Check ==="

# 1. Container running?
if docker ps | grep -q clickhouse; then
    echo "✓ ClickHouse container is running"
else
    echo "✗ ClickHouse container is NOT running"
    echo " Run: docker-compose up -d clickhouse"
    exit 1
fi

# 2. Can connect?
if docker exec optiinfra-clickhouse clickhouse-client --query="SELECT 1" > /dev/null 2>&1; then
    echo "✓ ClickHouse is responding"
else
    echo "✗ ClickHouse is not responding"
    echo " Check logs: docker logs optiinfra-clickhouse"
    exit 1
fi

# 3. Check version
VERSION=$(docker exec optiinfra-clickhouse clickhouse-client --query="SELECT version()")
echo "✓ ClickHouse version: $VERSION"

echo "==== All checks passed! ==="

```

## Step 9: Initialize ClickHouse Database

bash

```
cd ~/optiinfra
```

```
# Run initialization script
echo "Initializing ClickHouse database..."
docker exec -i optiinfra-clickhouse clickhouse-client < shared/clickhouse/migrations/init.sql
```

```
# Expected output:
# CREATE DATABASE IF NOT EXISTS optiinfra
# USE optiinfra
# CREATE TABLE IF NOT EXISTS cost_metrics_ts ...
# [etc for all 8 tables/views]
```

```
echo "✅ ClickHouse initialization complete"
```

## Validation:

```
bash
```

```

# Check all tables were created
echo "==== Verifying ClickHouse Tables ==="

TABLES=$(docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="SHOW TABLES" | wc -l)

if [ "$TABLES" -eq 8 ]; then
    echo " ✅ All 8 tables/views created"
else
    echo " ❌ Expected 8 tables, found $TABLES"
    echo " Tables created:"
    docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="SHOW TABLES"
    exit 1
fi

# List all tables
echo """
echo "ClickHouse tables:"
docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="SHOW TABLES"

# Expected output:
# cost_metrics_hourly_mv
# cost_metrics_ts
# performance_metrics_hourly_mv
# performance_metrics_ts
# quality_metrics_hourly_mv
# quality_metrics_ts
# resource_metrics_hourly_mv
# resource_metrics_ts

```

## Step 10: Verify Table Structures

bash

```
cd ~/optiinfra
```

```
# Check cost_metrics_ts structure
echo "==== cost_metrics_ts structure ===="
docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="DESCRIBE TABLE cost_metrics_ts"

# Check materialized view
echo """
echo "==== cost_metrics_hourly_mv structure ===="
docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="DESCRIBE TABLE cost_metrics_hourly_mv"

# Check partitioning
echo """
echo "==== Partition configuration ===="
docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="
SELECT
    table,
    partition_key,
    sorting_key
FROM system.tables
WHERE database = 'optiinfra' AND table LIKE '%_ts'
FORMAT Vertical
""
```

## Validation:

```
bash
```

```
# Verify all base tables have TTL
echo "==== Checking TTL Configuration ==="

for table in cost_metrics_ts performance_metrics_ts resource_metrics_ts quality_metrics_ts; do
    TTL=$(docker exec optiinfra-clickhouse clickhouse-client --database=optiinfra --query="
        SELECT ttl_expression
        FROM system.tables
        WHERE database='optiinfra' AND table='$table'
    ")

    if [ -n "$TTL" ]; then
        echo "✓ $table has TTL: $TTL"
    else
        echo "✗ $table missing TTL"
    fi
done
```

## Step 11: Test Python Client Connection

```
bash
```

```
cd ~/optiinfra
```

```
# Test client import and connection
python << EOF
from shared.clickhouse.client import get_clickhouse_client

print("== Testing ClickHouse Client ==")

# Get client
client = get_clickhouse_client()
print(" ✅ Client initialized")

# Test ping
if client.ping():
    print(" ✅ ClickHouse connection successful!")
else:
    print(" ❌ ClickHouse connection failed")
    exit(1)

# Show tables
result = client.client.execute("SHOW TABLES")
print(f"\n ✅ Found {len(result)} tables:")
for table in result:
    print(f" - {table[0]}")

print("\n== All tests passed! ==")
EOF
```

## Expected Output:

==== Testing ClickHouse Client ====

- Client initialized
- ClickHouse connection successful!

- Found 8 tables:

- cost\_metrics\_hourly\_mv
- cost\_metrics\_ts
- performance\_metrics\_hourly\_mv
- performance\_metrics\_ts
- quality\_metrics\_hourly\_mv
- quality\_metrics\_ts
- resource\_metrics\_hourly\_mv
- resource\_metrics\_ts

==== All tests passed! ====

## Step 12: Test Data Insertion

bash

```
cd ~/optiinfra
```

```
# Insert test cost metrics
python << 'EOF'
from shared.clickhouse.client import get_clickhouse_client
from datetime import datetime
import uuid

print("== Testing Data Insertion ==")

client = get_clickhouse_client()
test_customer_id = '123e4567-e89b-12d3-a456-426614174000'

# Insert 2 test cost metrics
cost_metrics = [
    {
        'timestamp': datetime.now(),
        'customer_id': test_customer_id,
        'cloud_provider': 'aws',
        'service_name': 'ec2',
        'instance_id': 'i-test123',
        'instance_type': 'm5.xlarge',
        'region': 'us-east-1',
        'cost_per_hour': 0.192,
        'utilization_percent': 45.5,
        'is_spot': 0,
        'is_reserved': 0
    },
    {
        'timestamp': datetime.now(),
        'customer_id': test_
```