

FOUNDATION-0.4: Qdrant Vector Database - PART 1 (Code)

🎯 CONTEXT

Phase: FOUNDATION (Week 1 - Day 2 Evening)

Component: Qdrant Vector Database Setup

Estimated Time: 15 min AI execution + 10 min verification

Complexity: MEDIUM

Risk Level: LOW

Files: Part 1 of 2 (Code implementation)

MILESTONE: Vector storage for agent memory and learning! 💬

📦 DEPENDENCIES

Must Complete First:

- **P-01:** Bootstrap Project Structure ✅ COMPLETED
- **FOUNDATION-0.2a-0.2e:** PostgreSQL schemas ✅ COMPLETED
- **FOUNDATION-0.3:** ClickHouse ✅ COMPLETED

Required Services Running:

```
bash

# Verify all services are healthy
cd ~/optinfra
make verify

# Expected output:
# PostgreSQL... ✅ HEALTHY
# ClickHouse... ✅ HEALTHY
# Qdrant... ✅ HEALTHY
# Redis... ✅ HEALTHY
```

🎯 OBJECTIVE

Set up **Qdrant vector database** to enable agents to learn from past decisions and retrieve relevant context.

What We're Building:

3 Vector Collections:

1. **cost_optimization_knowledge** - Past cost optimization decisions and outcomes
2. **performance_patterns** - Successful performance optimization patterns
3. **customer_context** - Customer-specific learning and preferences

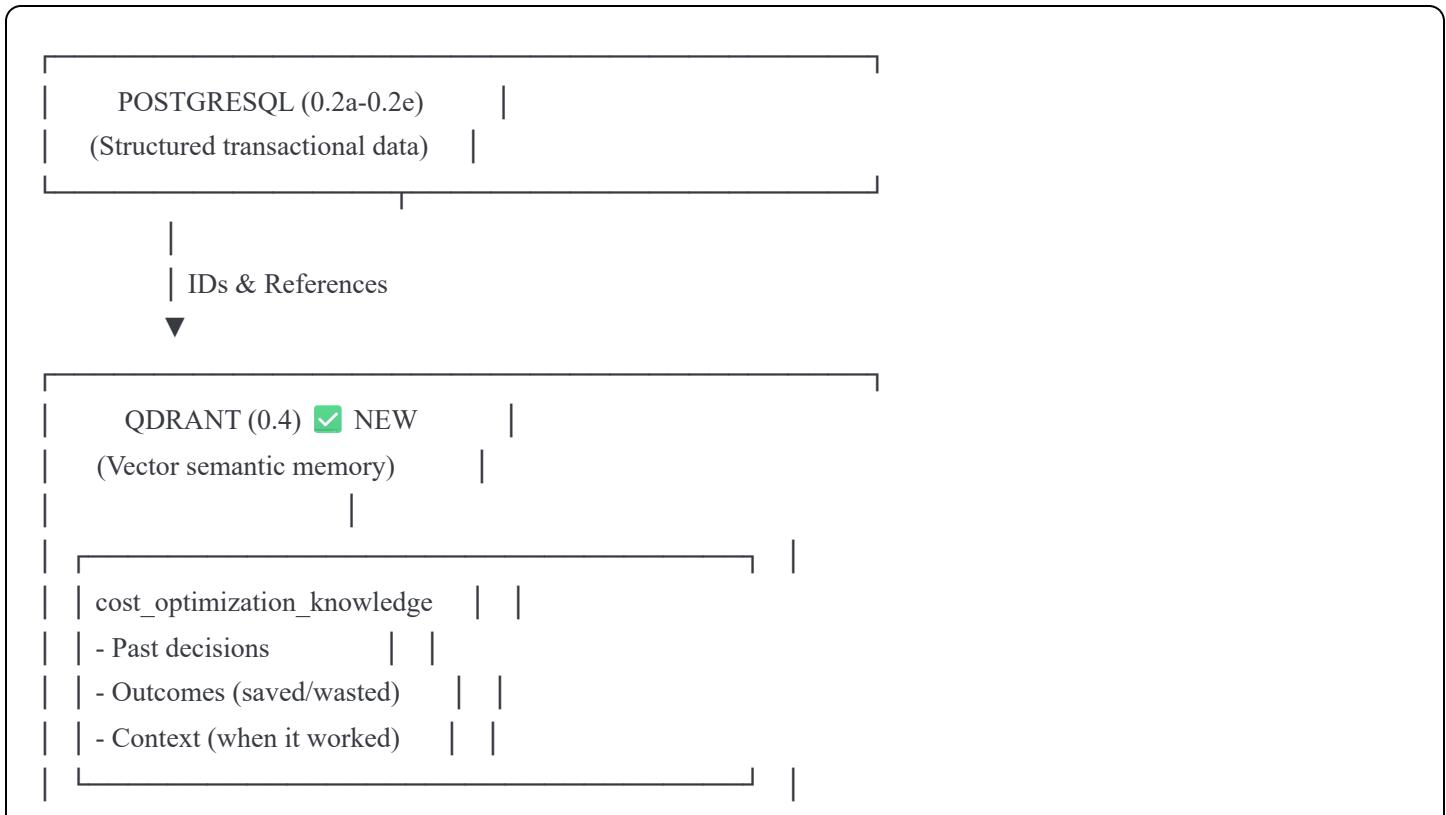
Python Client:

- Easy-to-use interface for embeddings
- Similarity search for context retrieval
- Outcome tracking for learning

Why Qdrant?

Feature	Traditional DB	Qdrant
Semantic search	✗	✓
Similarity matching	✗	✓
Vector storage	✗	✓ Optimized
Speed	N/A	Sub-millisecond
Best for	Exact match	Semantic similarity

Architecture:



```
    |  
    | performance_patterns | |  
    | - Successful optimizations | |  
    | - Configuration that worked | |  
    | - Similar past scenarios | |
```

```
    |  
    | customer_context | |  
    | - Customer preferences | |  
    | - Past interactions | |  
    | - Specific constraints | |
```

Use Cases:

Cost Agent:

"I want to migrate to spot instances. What worked for similar customers?" → Searches `cost_optimization_knowledge` for similar migrations → Returns: "3 similar cases, 2 succeeded (35% savings), 1 failed (workload too variable)"

Performance Agent:

"Customer has high P95 latency. What optimizations worked before?" → Searches `performance_patterns` for similar scenarios → Returns: "KV cache tuning worked in 5 similar cases (2.3x improvement)"

Application Agent:

"Customer prefers slow rollouts. What's their risk tolerance?" → Searches `customer_context` for preferences → Returns: "Customer X prefers 10%→25%→50%→100% rollout with 24hr validation"

FILE 1: Qdrant Collection Schemas

Location: `~/optiinfra/shared/qdrant/schemas/collections.py`

```
python
```

"""

Qdrant collection schemas and configurations.

Defines the structure and settings for all vector collections
used by OptiInfra agents for memory and learning.

"""

```
from dataclasses import dataclass
from typing import Dict, Any
from qdrant_client.models import Distance, VectorParams
```

```
@dataclass
class CollectionConfig:
    """Configuration for a Qdrant collection."""

    name: str
    vector_size: int
    distance: Distance
    description: str
    payload_schema: Dict[str, str]
```

```
def to_vector_params(self) -> VectorParams:
    """Convert to Qdrant VectorParams."""
    return VectorParams(
        size=self.vector_size,
        distance=self.distance
    )
```

```
# =====
# COLLECTION CONFIGURATIONS
# =====
```

```
COST_OPTIMIZATION KNOWLEDGE = CollectionConfig(
    name="cost_optimization_knowledge",
    vector_size=1536, # OpenAI ada-002 embedding size
    distance=Distance.COSINE,
    description="Knowledge base of past cost optimization decisions and outcomes",
    payload_schema={
        "optimization_id": "UUID - Link to optimizations table",
        "customer_id": "UUID - Which customer",
        "optimization_type": "String - spot_migration, reserved_instance, right_sizing",
```

```
"decision_context": "String - Why this decision was made",
"outcome": "String - success, failed, rolled_back",
"savings_percent": "Float - Actual savings achieved (if success)",
"cost_impact": "Float - Dollar savings per month",
"execution_date": "DateTime - When this was executed",
"cloud_provider": "String - aws, gcp, azure",
"instance_type": "String - m5.xlarge, etc",
"workload_characteristics": "String - Description of workload",
"lessons_learned": "String - What we learned from this",
"confidence_score": "Float - How confident we were (0-1)",
"customer_feedback": "String - Customer's feedback (if any)"
}
)
```

```
PERFORMANCE_PATTERNS = CollectionConfig(
    name="performance_patterns",
    vector_size=1536,
    distance=Distance.COSINE,
    description="Successful performance optimization patterns",
    payload_schema={
        "optimization_id": "UUID - Link to optimizations table",
        "customer_id": "UUID - Which customer",
        "service_type": "String - vllm, tgi, sclang",
        "model_name": "String - Which LLM model",
        "optimization_applied": "String - What was changed",
        "problem_description": "String - Original performance issue",
        "solution_description": "String - How it was solved",
        "before_latency_p95": "Float - P95 latency before (ms)",
        "after_latency_p95": "Float - P95 latency after (ms)",
        "improvement_factor": "Float - How much faster (2.3x, etc)",
        "config_changes": "JSONB - Specific configuration changes",
        "side_effects": "String - Any negative impacts observed",
        "execution_date": "DateTime - When applied",
        "stability_score": "Float - How stable post-change (0-1)",
        "replicable": "Boolean - Can this be repeated for similar cases"
    }
)
```

```
CUSTOMER_CONTEXT = CollectionConfig(
    name="customer_context",
    vector_size=1536,
    distance=Distance.COSINE,
    description="Customer-specific context, preferences, and constraints",
    payload_schema={
```

```
"customer_id": "UUID - Which customer",
"context_type": "String - preference, constraint, historical_note",
"topic": "String - What this context is about",
"content": "String - The actual context/note",
"source": "String - How we learned this (conversation, observation, explicit)",
"confidence": "Float - How confident we are (0-1)",
"created_at": "DateTime - When this context was added",
"updated_at": "DateTime - Last update",
"priority": "String - low, medium, high, critical",
"applies_to_agents": "List[String] - Which agents should use this",
"examples": "String - Example scenarios where this applies",
"exceptions": "String - When this doesn't apply"
}

)
```

```
# =====
# COLLECTION REGISTRY
# =====
```

```
ALL_COLLECTIONS = [
    COST_OPTIMIZATION KNOWLEDGE,
    PERFORMANCE PATTERNS,
    CUSTOMER CONTEXT
]
```

```
def get_collection_config(collection_name: str) -> CollectionConfig:
```

```
    """
```

Get configuration for a specific collection.

Args:

collection_name: Name of the collection

Returns:

CollectionConfig for the requested collection

Raises:

ValueError: If collection not found

```
    """
```

```
for config in ALL_COLLECTIONS:
    if config.name == collection_name:
        return config
```

```
raise ValueError(f"Unknown collection: {collection_name}")
```

FILE 2: Qdrant Client

Location: `~/optiinfra/shared/qdrant/client.py`

```
python
```

....

Qdrant client for vector storage and similarity search.

Provides easy-to-use interface for:

- Storing optimization decisions with embeddings
- Searching for similar past decisions
- Learning from outcomes

Usage:

```
from shared.qdrant.client import get_qdrant_client
```

```
client = get_qdrant_client()
```

```
# Store a decision
```

```
client.store_cost_decision(  
    optimization_id="...",  
    decision_context="Migrating to spot for stable batch workload",  
    outcome="success",  
    savings_percent=38.5,  
    ...  
)
```

```
# Search for similar decisions
```

```
results = client.search_similar_cost_decisions(  
    query="migrate to spot instances for batch processing",  
    limit=5  
)
```

....

```
from qdrant_client import QdrantClient  
from qdrant_client.models import (  
    Distance, VectorParams, PointStruct, Filter,  
    FieldCondition, MatchValue, SearchRequest  
)  
from typing import List, Dict, Any, Optional  
import os  
import uuid  
from datetime import datetime  
import logging  
  
from shared.qdrant.schemas.collections import (  
    ALL_COLLECTIONS,  
    COST_OPTIMIZATION KNOWLEDGE,
```

```
PERFORMANCE_PATTERNS,
CUSTOMER_CONTEXT,
get_collection_config
)

logger = logging.getLogger(__name__)

class QdrantVectorClient:
    """Client for vector storage and similarity search in Qdrant."""

    def __init__(self):
        """Initialize Qdrant client."""
        self.client = QdrantClient(
            host=os.getenv('QDRANT_HOST', 'localhost'),
            port=int(os.getenv('QDRANT_PORT', 6333)),
            api_key=os.getenv('QDRANT_API_KEY', None)
        )

        # Embedding function (you'll integrate OpenAI/other provider)
        self.embedding_function = self._default_embedding_function

        logger.info(f"Qdrant client initialized: {self.client._client.rest_uri}")
```

```
def _default_embedding_function(self, text: str) -> List[float]:
```

```
    """
```

```
    Default embedding function (placeholder).
```

In production, replace with actual embedding model:

- OpenAI ada-002
- Sentence-Transformers
- Custom model

Args:

text: Text to embed

Returns:

List of floats (embedding vector)

```
    """
```

TODO: Replace with actual embedding model

For now, return random vector for testing

```
import random
```

```
return [random.random() for _ in range(1536)]
```

```
def ping(self) -> bool:  
    """  
    Check if Qdrant is accessible.  
  
    Returns:  
        bool: True if Qdrant responds, False otherwise  
    """
```

```
try:  
    collections = self.client.get_collections()  
    return True  
except Exception as e:  
    logger.error(f"Qdrant ping failed: {e}")  
    return False
```

```
def initialize_collections(self):  
    """  
    Create all collections if they don't exist.  
  
    This is idempotent - safe to call multiple times.  
    """  
    existing = {c.name for c in self.client.get_collections().collections}
```

```
for config in ALL_COLLECTIONS:  
    if config.name not in existing:  
        logger.info(f"Creating collection: {config.name}")  
        self.client.create_collection(  
            collection_name=config.name,  
            vectors_config=config.to_vector_params()  
        )  
        logger.info(f"✓ Created collection: {config.name}")  
    else:  
        logger.info(f"Collection already exists: {config.name}")
```

```
# ======  
# COST OPTIMIZATION KNOWLEDGE  
# ======
```

```
def store_cost_decision(  
    self,  
    optimization_id: str,  
    customer_id: str,  
    optimization_type: str,  
    decision_context: str,  
    outcome: str,
```

```
savings_percent: Optional[float] = None,  
cost_impact: Optional[float] = None,  
**kwargs  
) -> str:  
    """
```

Store a cost optimization decision with embedding.

Args:

```
optimization_id: UUID from optimizations table  
customer_id: UUID from customers table  
optimization_type: spot_migration, reserved_instance, right_sizing  
decision_context: Why this decision was made (text for embedding)  
outcome: success, failed, rolled_back  
savings_percent: Actual savings (if success)  
cost_impact: Dollar savings per month  
**kwargs: Additional payload fields
```

Returns:

```
str: Point ID in Qdrant
```

Example:

```
point_id = client.store_cost_decision(  
    optimization_id="123e4567-e89b-12d3-a456-426614174000",  
    customer_id="789e0123-e89b-12d3-a456-426614174000",  
    optimization_type="spot_migration",  
    decision_context="Migrating batch processing workload to spot instances. Workload is tolerant to interruptions and  
    outcome='success',  
    savings_percent=38.5,  
    cost_impact=18000,  
    cloud_provider='aws',  
    instance_type='m5.xlarge',  
    workload_characteristics="Batch ETL jobs, 4-6 hour runtime, can checkpoint",  
    lessons_learned="Spot worked well for this workload. No interruptions in first month."  
)
```

```
    """
```

Generate embedding from decision context

```
embedding = self.embedding_function(decision_context)
```

Create payload

```
payload = {  
    "optimization_id": optimization_id,  
    "customer_id": customer_id,  
    "optimization_type": optimization_type,  
    "decision_context": decision_context,
```

```

    "outcome": outcome,
    "savings_percent": savings_percent,
    "cost_impact": cost_impact,
    "execution_date": datetime.now().isoformat(),
    **kwargs
}

# Generate point ID
point_id = str(uuid.uuid4())

# Store in Qdrant
self.client.upsert(
    collection_name=COST_OPTIMIZATION KNOWLEDGE.name,
    points=[
        PointStruct(
            id=point_id,
            vector=embedding,
            payload=payload
        )
    ]
)
)

logger.info(f"Stored cost decision: {point_id}")
return point_id

```

```

def search_similar_cost_decisions(
    self,
    query: str,
    limit: int = 5,
    filter_outcome: Optional[str] = None,
    filter_type: Optional[str] = None
) -> List[Dict[str, Any]]:
    """
    Search for similar cost optimization decisions.

```

Args:

- query: Natural language query describing the scenario
- limit: Number of results to return
- filter_outcome: Optional filter (success, failed, rolled_back)
- filter_type: Optional filter by optimization type

Returns:

- List of similar decisions with scores

Example:

```
results = client.search_similar_cost_decisions(  
    query="migrate batch processing to spot instances",  
    limit=5,  
    filter_outcome="success"  
)  
  
for result in results:  
    print(f"Score: {result['score']:.3f}")  
    print(f"Context: {result['payload']['decision_context']}")  
    print(f"Outcome: {result['payload']['outcome']}")  
    print(f"Savings: {result['payload']['savings_percent']}%")  
      
    """  
    # Generate query embedding  
    query_embedding = self.embedding_function(query)  
  
    # Build filter  
    must_conditions = []  
    if filter_outcome:  
        must_conditions.append(  
            FieldCondition(  
                key="outcome",  
                match=MatchValue(value=filter_outcome)  
            )  
        )  
    if filter_type:  
        must_conditions.append(  
            FieldCondition(  
                key="optimization_type",  
                match=MatchValue(value=filter_type)  
            )  
        )  
  
    search_filter = Filter(must=must_conditions) if must_conditions else None  
  
    # Search  
    results = self.client.search(  
        collection_name=COST_OPTIMIZATION KNOWLEDGE.name,  
        query_vector=query_embedding,  
        limit=limit,  
        query_filter=search_filter  
)  
  
    return [
```

```

{
    "id": result.id,
    "score": result.score,
    "payload": result.payload
}
for result in results
]

# =====
# PERFORMANCE PATTERNS
# =====

def store_performance_pattern(
    self,
    optimization_id: str,
    customer_id: str,
    service_type: str,
    model_name: str,
    problem_description: str,
    solution_description: str,
    before_latency_p95: float,
    after_latency_p95: float,
    **kwargs
) -> str:
    """
    Store a successful performance optimization pattern.
    """

    Store a successful performance optimization pattern.

```

Args:

- optimization_id: UUID from optimizations table
- customer_id: UUID from customers table
- service_type: vllm, tgi, sglang
- model_name: Which LLM model
- problem_description: Original issue (for embedding)
- solution_description: How it was solved (for embedding)
- before_latency_p95: P95 latency before (ms)
- after_latency_p95: P95 latency after (ms)
- **kwargs: Additional payload fields

Returns:

- str: Point ID in Qdrant

"""

```

# Combine problem + solution for embedding
combined_text = f'{problem_description}\n\n{solution_description}'
embedding = self.embedding_function(combined_text)

```

```

# Calculate improvement
improvement_factor = before_latency_p95 / after_latency_p95 if after_latency_p95 > 0 else 0

# Create payload
payload = {
    "optimization_id": optimization_id,
    "customer_id": customer_id,
    "service_type": service_type,
    "model_name": model_name,
    "problem_description": problem_description,
    "solution_description": solution_description,
    "before_latency_p95": before_latency_p95,
    "after_latency_p95": after_latency_p95,
    "improvement_factor": improvement_factor,
    "execution_date": datetime.now().isoformat(),
    **kwargs
}

point_id = str(uuid.uuid4())

self.client.upsert(
    collection_name=PERFORMANCE_PATTERNS.name,
    points=[
        PointStruct(
            id=point_id,
            vector=embedding,
            payload=payload
        )
    ]
)

logger.info(f"Stored performance pattern: {point_id}")

return point_id

def search_similar_performance_patterns(
    self,
    query: str,
    limit: int = 5,
    filter_service_type: Optional[str] = None
) -> List[Dict[str, Any]]:
    """
    Search for similar performance optimization patterns.

```

Args:

```
query: Description of the performance issue
limit: Number of results
filter_service_type: Optional filter (vllm, tgi, sclang)
```

Returns:

```
List of similar patterns with scores
```

=====

```
query_embedding = self.embedding_function(query)
```

```
search_filter = None
if filter_service_type:
    search_filter = Filter(
        must=[
            FieldCondition(
                key="service_type",
                match=MatchValue(value=filter_service_type)
            )
        ]
    )
```

```
results = self.client.search(
    collection_name=PERFORMANCE_PATTERNS.name,
    query_vector=query_embedding,
    limit=limit,
    query_filter=search_filter
)
```

```
return [
{
    "id": result.id,
    "score": result.score,
    "payload": result.payload
}
for result in results
]
```

```
# =====
```

```
# CUSTOMER CONTEXT
```

```
# =====
```

```
def store_customer_context(
    self,
    customer_id: str,
```

```
    context_type: str,  
    topic: str,  
    content: str,  
    confidence: float = 0.8,  
    **kwargs  
) -> str:  
    """
```

Store customer-specific context or preference.

Args:

```
    customer_id: UUID from customers table  
    context_type: preference, constraint, historical_note  
    topic: What this context is about  
    content: The actual context (for embedding)  
    confidence: How confident we are (0-1)  
    **kwargs: Additional payload fields
```

Returns:

```
    str: Point ID in Qdrant
```

Example:

```
client.store_customer_context(  
    customer_id="123e4567-e89b-12d3-a456-426614174000",  
    context_type="preference",  
    topic="rollout_strategy",  
    content="Customer prefers slow, cautious rollouts with 24-hour validation periods between stages. They value stability and consistency.",  
    confidence=0.9,  
    source="conversation with CTO on 2025-01-15",  
    priority="high",  
    applies_to_agents=["performance_agent", "cost_agent"]  
)  
"""  
embedding = self.embedding_function(content)  
  
payload = {  
    "customer_id": customer_id,  
    "context_type": context_type,  
    "topic": topic,  
    "content": content,  
    "confidence": confidence,  
    "created_at": datetime.now().isoformat(),  
    "updated_at": datetime.now().isoformat(),  
    **kwargs  
}
```

```
point_id = str(uuid.uuid4())

self.client.upsert(
    collection_name=CUSTOMER_CONTEXT.name,
    points=[
        PointStruct(
            id=point_id,
            vector=embedding,
            payload=payload
        )
    ]
)

logger.info(f"Stored customer context: {point_id}")
return point_id
```

```
def search_customer_context(
    self,
    customer_id: str,
    query: str,
    limit: int = 3
) -> List[Dict[str, Any]]:
    """
```

Search for relevant customer context.

Args:

```
customer_id: Which customer
query: What to search for
limit: Number of results
```

Returns:

List of relevant context with scores

"""

```
query_embedding = self.embedding_function(query)
```

```
search_filter = Filter(
    must=[
        FieldCondition(
            key="customer_id",
            match=MatchValue(value=customer_id)
        )
    ]
)
```

```
results = self.client.search(
    collection_name=CUSTOMER_CONTEXT.name,
    query_vector=query_embedding,
    limit=limit,
    query_filter=search_filter
)

return [
{
    "id": result.id,
    "score": result.score,
    "payload": result.payload
}
for result in results
]
```

```
# =====
```

```
# SINGLETON PATTERN
```

```
# =====
```

```
_qdrant_client = None
```

```
def get_qdrant_client() -> QdrantVectorClient:
```

```
"""
```

```
Get singleton Qdrant client instance.
```

```
Returns:
```

```
QdrantVectorClient: Singleton client instance
```

```
Example:
```

```
client = get_qdrant_client()
```

```
if client.ping():
```

```
    print("Qdrant is ready!")
```

```
"""
```

```
global _qdrant_client
```

```
if _qdrant_client is None:
```

```
    _qdrant_client = QdrantVectorClient()
```

```
return _qdrant_client
```

FILE 3: Package Initialization

Location: `(~/optiinfra/shared/qdrant/_init_.py)`

```
python
```

....

Qdrant vector database package.

Provides vector storage and semantic search for:

- Cost optimization knowledge (past decisions → outcomes)
- Performance patterns (successful optimizations)
- Customer context (preferences, constraints)

Usage:

```
from shared.qdrant import get_qdrant_client

client = get_qdrant_client()
client.initialize_collections()

# Store decision
client.store_cost_decision(...)

# Search for similar
results = client.search_similar_cost_decisions(...)
```

....

```
from shared.qdrant.client import (
    QdrantVectorClient,
    get_qdrant_client
)
```

```
from shared.qdrant.schemas.collections import (
    COST_OPTIMIZATION KNOWLEDGE,
    PERFORMANCE PATTERNS,
    CUSTOMER CONTEXT,
    ALL COLLECTIONS,
    get_collection_config
)
```

```
_all_ = [
    'QdrantVectorClient',
    'get_qdrant_client',
    'COST_OPTIMIZATION KNOWLEDGE',
    'PERFORMANCE PATTERNS',
    'CUSTOMER CONTEXT',
    'ALL COLLECTIONS',
```

```
'get_collection_config'  
]
```

📁 FILE 4: Schema Package Init

Location: `~/optiinfra/shared/qdrant/schemas/_init_.py`

```
python  
=====  
Qdrant schema definitions.  
=====  
  
from shared.qdrant.schemas.collections import (  
    CollectionConfig,  
    COST_OPTIMIZATION KNOWLEDGE,  
    PERFORMANCE PATTERNS,  
    CUSTOMER CONTEXT,  
    ALL COLLECTIONS,  
    get_collection_config  
)  
  
__all__ = [  
    'CollectionConfig',  
    'COST_OPTIMIZATION KNOWLEDGE',  
    'PERFORMANCE PATTERNS',  
    'CUSTOMER CONTEXT',  
    'ALL COLLECTIONS',  
    'get_collection_config'  
]
```

📁 FILE 5: Update Requirements

Location: `~/optiinfra/shared/requirements.txt`

```
txt
```

```
# Existing dependencies...
```

```
sqlalchemy==2.0.23
```

```
alembic==1.12.1
```

```
psycopg2-binary==2.9.9
```

```
clickhouse-driver==0.2.6
```

```
# Qdrant client (ADD THIS)
```

```
qdrant-client==1.7.0
```

```
# Other dependencies...
```

FILE 6: README Documentation

Location: `~/optiinfra/shared/qdrant/README.md`

markdown

Qdrant Vector Database

Vector storage and semantic search for OptiInfra agent memory and learning.

Overview

Qdrant enables agents to:

- **Learn from past decisions** - "What worked for similar scenarios?"
- **Retrieve relevant context** - "What do we know about this customer?"
- **Semantic search** - Find similar situations, not just exact matches

Architecture

Collections

1. **cost_optimization_knowledge** - Past cost optimization decisions
2. **performance_patterns** - Successful performance optimizations
3. **customer_context** - Customer-specific preferences and constraints

Each collection stores:

- **Vector embeddings** (1536-dim from OpenAI ada-002)
- **Metadata payload** (structured data about the decision/pattern)
- **Cosine similarity** for semantic matching

Usage

Initialize Collections

```
```bash
Run once to create all collections
python << 'EOF'
from shared.qdrant import get_qdrant_client
```

```
client = get_qdrant_client()
client.initialize_collections()
```

```
print("✅ Collections initialized")
```

```
EOF
```

```
```
```

Python Client

```
```python
from shared.qdrant import get_qdrant_client
```

##### # Get client

```
client = get_qdrant_client()
```

```
Check connection
if client.ping():
 print("✅ Qdrant connected!")

Store a cost optimization decision
point_id = client.store_cost_decision(
 optimization_id="123e4567-e89b-12d3-a456-426614174000",
 customer_id="789e0123-e89b-12d3-a456-426614174000",
 optimization_type="spot_migration",
 decision_context="Migrating batch ETL workload to spot instances. Workload can handle interruptions with checkpointing",
 outcome="success",
 savings_percent=38.5,
 cost_impact=18000,
 cloud_provider="aws",
 instance_type="m5.xlarge"
)

Search for similar decisions
results = client.search_similar_cost_decisions(
 query="migrate batch processing to spot",
 limit=5,
 filter_outcome="success"
)

for result in results:
 print(
```