

# FOUNDATION-0.4: Qdrant Vector Database - PART 2 (Execution & Testing)

## 🎯 CONTEXT

**Phase:** FOUNDATION (Week 1 - Day 2 Evening)

**Component:** Qdrant Vector Database - Execution & Validation

**Estimated Time:** 10 min execution + 15 min testing

**Complexity:** MEDIUM

**Risk Level:** LOW

**Files:** Part 2 of 2 (Execution, testing, validation)

---

## 📦 PREREQUISITES

**Must Have Completed:**

- **PART 1** - All code files created
- **0.2a-0.2e** - PostgreSQL schemas complete
- **0.3** - ClickHouse complete

**Verify PART 1 Files Exist:**

```
bash

cd ~/optiinfra

# Check all files created
ls -lh shared/qdrant/schemas/collections.py
ls -lh shared/qdrant/client.py
ls -lh shared/qdrant/__init__.py
ls -lh shared/qdrant/schemas/__init__.py
ls -lh shared/qdrant/README.md
```

*# Expected: All files present*

# STEP-BY-STEP EXECUTION

## Step 1: Create Directory Structure

```
bash

cd ~/optiinfra/shared

# Create qdrant directory structure
mkdir -p qdrant/schemas

# Verify structure
tree qdrant/ -L 2
# Expected:
# qdrant/
#   └── schemas/
```

## Validation:

```
bash

# Check directories exist
[ -d "qdrant" ] && echo "✓ qdrant/ created" || echo "✗ qdrant/ missing"
[ -d "qdrant/schemas" ] && echo "✓ schemas/ created" || echo "✗ schemas/ missing"
```

---

## Step 2: Create Collection Schemas

```
bash
```

```
cd ~/optiinfra/shared/qdrant/schemas
```

```
# Create collections.py (copy from PART 1, FILE 1)
cat > collections.py << 'EOF'
=====
Qdrant collection schemas and configurations.
[... COPY ENTIRE COLLECTIONS.PY FROM PART 1, FILE 1 ...]
=====
EOF

# Verify file created
ls -lh collections.py
wc -l collections.py
# Expected: ~200 lines
```

## Validation:

```
bash
```

```
# Check file size
FILE_SIZE=$(wc -l < collections.py)
if [ "$FILE_SIZE" -gt 150 ]; then
    echo "✓ collections.py created ($FILE_SIZE lines)"
else
    echo "✗ collections.py too small ($FILE_SIZE lines)"
fi

# Check for key classes
grep -q "class CollectionConfig" collections.py && echo "✓ CollectionConfig found" || echo "✗ Missing"
grep -q "COST_OPTIMIZATION KNOWLEDGE" collections.py && echo "✓ COST_OPTIMIZATION KNOWLEDGE" || echo "✗ Missing"
grep -q "PERFORMANCE PATTERNS" collections.py && echo "✓ PERFORMANCE PATTERNS found" || echo "✗ Missing"
grep -q "CUSTOMER CONTEXT" collections.py && echo "✓ CUSTOMER CONTEXT found" || echo "✗ Missing"
```

## Step 3: Create Python Client

```
bash
```

```
cd ~/optiinfra/shared/qdrant
```

```
# Create client.py (copy from PART 1, FILE 2)
cat > client.py << 'EOF'
=====
Qdrant client for vector storage and similarity search.
[... COPY ENTIRE CLIENT.PY FROM PART 1, FILE 2 ...]
=====
EOF

# Verify file
ls -lh client.py
wc -l client.py
# Expected: ~550 lines
```

## Validation:

```
bash

# Check file size
FILE_SIZE=$(wc -l < client.py)
if [ "$FILE_SIZE" -gt 500 ]; then
    echo "✓ client.py created ($FILE_SIZE lines)"
else
    echo "✗ client.py too small ($FILE_SIZE lines)"
fi

# Check for key classes and methods
grep -q "class QdrantVectorClient" client.py && echo "✓ QdrantVectorClient class found" || echo "✗ Missing"
grep -q "def store_cost_decision" client.py && echo "✓ store_cost_decision found" || echo "✗ Missing"
grep -q "def search_similar_cost_decisions" client.py && echo "✓ search_similar_cost_decisions found" || echo "✗ Missing"
grep -q "get_qdrant_client" client.py && echo "✓ get_qdrant_client found" || echo "✗ Missing"
```

## Step 4: Create Package Files

```
bash
```

```
cd ~/optiinfra/shared/qdrant
```

```
# Create __init__.py (copy from PART 1, FILE 3)
```

```
cat > __init__.py << 'EOF'
```

```
'''
```

```
Qdrant vector database package.
```

```
[... COPY FROM PART 1, FILE 3 ...]
```

```
'''
```

```
EOF
```

```
# Create schemas/__init__.py (copy from PART 1, FILE 4)
```

```
cat > schemas/__init__.py << 'EOF'
```

```
'''
```

```
Qdrant schema definitions.
```

```
[... COPY FROM PART 1, FILE 4 ...]
```

```
'''
```

```
EOF
```

```
# Verify files
```

```
ls -lh __init__.py
```

```
ls -lh schemas/__init__.py
```

## Validation:

```
bash
```

```
# Check files exist
```

```
[ -f "__init__.py" ] && echo "✓ __init__.py created" || echo "✗ Missing"
```

```
[ -f "schemas/__init__.py" ] && echo "✓ schemas/__init__.py created" || echo "✗ Missing"
```

## Step 5: Update Requirements

```
bash
```

```
cd ~/optiinfra/shared
```

```
# Add qdrant-client to requirements.txt
echo "" >> requirements.txt
echo "# Qdrant vector database (FOUNDATION-0.4)" >> requirements.txt
echo "qdrant-client==1.7.0" >> requirements.txt
```

```
# Verify addition
```

```
tail -5 requirements.txt
```

## Validation:

```
bash
```

```
# Check if qdrant-client is in requirements
grep -q "qdrant-client" requirements.txt && echo "✓ qdrant-client added" || echo "✗ Missing"
```

## Step 6: Install Dependencies

```
bash
```

```
cd ~/optiinfra/shared
```

```
# Install qdrant-client
pip install qdrant-client==1.7.0

# Verify installation
python -c "import qdrant_client; print('✓ qdrant-client installed successfully')"
# Expected: ✓ qdrant-client installed successfully
```

## Validation:

```
bash
```

```
# Test import
python << 'EOF'
try:
    import qdrant_client
    from qdrant_client.models import Distance, VectorParams
    print("✓ qdrant_client imports successfully")
    print(f" Version: {qdrant_client.__version__}")
except ImportError as e:
    print(f"✗ Import failed: {e}")
EOF
```

## Step 7: Create README

```
bash

cd ~/optiinfra/shared/qdrant

# Create README.md (copy from PART 1, FILE 6)
cat > README.md << 'EOF'
# Qdrant Vector Database
[... COPY ENTIRE README FROM PART 1, FILE 6 ...]
EOF

# Verify file
ls -lh README.md
```

## Validation:

```
bash

# Check README exists and has content
FILE_SIZE=$(wc -l< README.md)
if [ "$FILE_SIZE" -gt 50 ]; then
    echo "✓ README.md created ($FILE_SIZE lines)"
else
    echo "✗ README.md too small"
fi
```

## Step 8: Verify Qdrant is Running

```
bash

cd ~/optiinfra

# Check if Qdrant container is running
docker ps | grep qdrant

# If not running, start all services
docker-compose up -d

# Wait for Qdrant to be ready (takes ~5 seconds)
echo "Waiting for Qdrant to start..."
sleep 5

# Test connection via REST API
curl -s http://localhost:6333/collections | head -20
# Expected: JSON response with collections (may be empty initially)
```

### Validation:

```
bash
```

```

# Comprehensive Qdrant check
echo "==== Qdrant Health Check ===="

# 1. Container running?
if docker ps | grep -q qdrant; then
    echo "✓ Qdrant container is running"
else
    echo "✗ Qdrant container is NOT running"
    echo "Run: docker-compose up -d qdrant"
    exit 1
fi

# 2. Can connect via REST?
if curl -s http://localhost:6333/collections > /dev/null 2>&1; then
    echo "✓ Qdrant REST API is responding"
else
    echo "✗ Qdrant is not responding"
    echo "Check logs: docker logs optiinfra-qdrant"
    exit 1
fi

# 3. Check version (if available)
VERSION=$(curl -s http://localhost:6333 | grep -o '"version": "[^"]*" | cut -d '"' -f4 || echo "unknown")
echo "✓ Qdrant version: $VERSION"

echo "==== All checks passed! ===="

```

## Step 9: Test Python Client Connection

bash

```
cd ~/optinfra
```

```
# Test client import and connection
python << 'EOF'
from shared.qdrant.client import get_qdrant_client

print("== Testing Qdrant Client ==")

# Get client
client = get_qdrant_client()
print("✓ Client initialized")

# Test ping
if client.ping():
    print("✓ Qdrant connection successful!")
else:
    print("✗ Qdrant connection failed")
    exit(1)

# Check existing collections
collections = client.client.get_collections()
print(f"\n✓ Current collections: {len(collections.collections)}")
for collection in collections.collections:
    print(f" - {collection.name}")

print("\n== All tests passed! ==")
EOF
```

## Expected Output:

```
== Testing Qdrant Client ==
✓ Client initialized
✓ Qdrant connection successful!

✓ Current collections: 0

== All tests passed! ==
```

## Step 10: Initialize Collections

```
bash
```

```
cd ~/optiinfra
```

```
# Create all collections
python << EOF
from shared.qdrant.client import get_qdrant_client

print("== Initializing Qdrant Collections ==")

client = get_qdrant_client()

# Initialize all collections
client.initialize_collections()

# Verify collections created
collections = client.client.get_collections()
print(f"\n✓ Collections created: {len(collections.collections)}")
for collection in collections.collections:
    info = client.client.get_collection(collection.name)
    print(f" - {collection.name}")
    print(f"   Vectors: {info.vectors_count if hasattr(info, 'vectors_count') else 0}")
    print(f"   Points: {info.points_count}")

print("\n== Initialization complete! ==")
EOF
```

### Expected Output:

```
==== Initializing Qdrant Collections ====
```

- ✓ Created collection: cost\_optimization\_knowledge
- ✓ Created collection: performance\_patterns
- ✓ Created collection: customer\_context

```
✓ Collections created: 3
```

```
- cost_optimization_knowledge
```

```
    Vectors: 0
```

```
    Points: 0
```

```
- performance_patterns
```

```
    Vectors: 0
```

```
    Points: 0
```

```
- customer_context
```

```
    Vectors: 0
```

```
    Points: 0
```

```
==== Initialization complete! ====
```

## Step 11: Test Storing Data

```
bash
```

```
cd ~/optiinfra
```

```
# Test storing cost optimization decision
python << EOF
from shared.qdrant.client import get_qdrant_client

print("== Testing Data Storage ==")

client = get_qdrant_client()

# Store a cost optimization decision
print("\n1. Storing cost optimization decision...")
point_id = client.store_cost_decision(
    optimization_id="123e4567-e89b-12d3-a456-426614174000",
    customer_id="789e0123-e89b-12d3-a456-426614174000",
    optimization_type="spot_migration",
    decision_context="Migrating batch ETL workload to spot instances. Workload can handle interruptions with checkpointing",
    outcome="success",
    savings_percent=38.5,
    cost_impact=18000,
    cloud_provider="aws",
    instance_type="m5.xlarge",
    workload_characteristics="Batch ETL, 4-6 hour runtime, checkpoint-able",
    lessons_learned="Spot worked well. No interruptions in first month. Checkpointing every 30min was sufficient."
)
print(f" ✅ Stored with point_id: {point_id}")

# Store a performance pattern
print("\n2. Storing performance pattern...")
point_id = client.store_performance_pattern(
    optimization_id="456e7890-e89b-12d3-a456-426614174000",
    customer_id="789e0123-e89b-12d3-a456-426614174000",
    service_type="vllm",
    model_name="llama-2-70b",
    problem_description="High P95 latency (800ms) due to inefficient KV cache settings",
    solution_description="Tuned KV cache block size from 16 to 32, enabled prefix caching",
    before_latency_p95=800.0,
    after_latency_p95=280.0,
    config_changes={"kv_cache_block_size": 32, "enable_prefix_caching": True},
    replicable=True
)
print(f" ✅ Stored with point_id: {point_id}")
```

```
# Store customer context
print("\n3. Storing customer context...")
point_id = client.store_customer_context(
    customer_id="789e0123-e89b-12d3-a456-426614174000",
    context_type="preference",
    topic="rollout_strategy",
    content="Customer prefers slow, cautious rollouts with 24-hour validation periods between stages. They value stability over speed and reliability over rapid deployment.",
    confidence=0.9,
    source="conversation with CTO on 2025-01-15",
    priority="high",
    applies_to_agents=["performance_agent", "cost_agent"]
)
print(f" ✅ Stored with point_id: {point_id}")

# Verify points stored
collections = client.client.get_collections()
print("\n==== Verification ===")
for collection in collections.collections:
    info = client.client.get_collection(collection.name)
    print(f"{collection.name}: {info.points_count} points")

print("\n==== Storage test complete! ===")
EOF
```

## Expected Output:

==== Testing Data Storage ====

1. Storing cost optimization decision...

Stored with point\_id: abc12345-...

2. Storing performance pattern...

Stored with point\_id: def67890-...

3. Storing customer context...

Stored with point\_id: ghi24680-...

==== Verification ====

cost\_optimization\_knowledge: 1 points

performance\_patterns: 1 points

customer\_context: 1 points

==== Storage test complete! ====

## Step 12: Test Similarity Search

bash

```
cd ~/optiinfra
```

```
# Test similarity search
python << EOF
from shared.qdrant.client import get_qdrant_client

print("== Testing Similarity Search ==")

client = get_qdrant_client()

# 1. Search for similar cost decisions
print("\n1. Searching for similar cost optimizations...")
results = client.search_similar_cost_decisions(
    query="migrate batch processing to spot instances",
    limit=3,
    filter_outcome="success"
)

print(f" Found {len(results)} similar decisions:")
for i, result in enumerate(results, 1):
    print(f"\n Result {i}:")
    print(f" - Score: {result['score']:.3f}")
    print(f" - Type: {result['payload']['optimization_type']}")
    print(f" - Savings: {result['payload']['savings_percent']}%")
    print(f" - Context: {result['payload']['decision_context'][:100]}")

# 2. Search for similar performance patterns
print("\n2. Searching for similar performance optimizations...")
results = client.search_similar_performance_patterns(
    query="high latency due to KV cache issues",
    limit=3,
    filter_service_type="vllm"
)

print(f" Found {len(results)} similar patterns:")
for i, result in enumerate(results, 1):
    print(f"\n Result {i}:")
    print(f" - Score: {result['score']:.3f}")
    print(f" - Improvement: {result['payload']['improvement_factor']:.1f}x")
    print(f" - Problem: {result['payload']['problem_description'][:80]}")

# 3. Search customer context
print("\n3. Searching customer context...")
```

```
results = client.search_customer_context(  
    customer_id="789e0123-e89b-12d3-a456-426614174000",  
    query="rollout preferences and risk tolerance",  
    limit=3  
)  
  
print(f" Found {len(results)} relevant context items:")  
for i, result in enumerate(results, 1):  
    print(f"\n Result {i}:")  
    print(f" - Score: {result['score']:.3f}")  
    print(f" - Topic: {result['payload']['topic']}")  
    print(f" - Content: {result['payload']['content'][:100]}...")  
  
print("\n==== Similarity search test complete! ===")  
EOF
```

### Expected Output:

==== Testing Similarity Search ====

1. Searching for similar cost optimizations...

Found 1 similar decisions:

Result 1:

- Score: 0.850
- Type: spot\_migration
- Savings: 38.5%
- Context: Migrating batch ETL workload to spot instances. Workload can handle interruptions with ch...

2. Searching for similar performance optimizations...

Found 1 similar patterns:

Result 1:

- Score: 0.920
- Improvement: 2.9x
- Problem: High P95 latency (800ms) due to inefficient KV cache settings...

3. Searching customer context...

Found 1 relevant context items:

Result 1:

- Score: 0.880
- Topic: rollout\_strategy
- Content: Customer prefers slow, cautious rollouts with 24-hour validation periods between stages...

==== Similarity search test complete! ====

## Step 13: Load More Test Data (Optional)

```
bash
```

```
cd ~/optiinfra
```

```
# Load multiple test records for better search testing
python << EOF
from shared.qdrant.client import get_qdrant_client

print("== Loading Additional Test Data ==")

client = get_qdrant_client()

# Add 5 more cost decisions
cost_decisions = [
    {
        "optimization_type": "reserved_instance",
        "decision_context": "Purchased 1-year reserved instances for stable production workload. Predictable usage pattern over",
        "outcome": "success",
        "savings_percent": 42.0,
        "cost_impact": 25000
    },
    {
        "optimization_type": "right_sizing",
        "decision_context": "Downsized from m5.2xlarge to m5.xlarge after observing 40% average CPU utilization for 2 weeks",
        "outcome": "success",
        "savings_percent": 28.5,
        "cost_impact": 8500
    },
    {
        "optimization_type": "spot_migration",
        "decision_context": "Attempted spot migration for real-time API service. Service was too sensitive to interruptions.",
        "outcome": "failed",
        "savings_percent": None,
        "cost_impact": None
    },
    {
        "optimization_type": "reserved_instance",
        "decision_context": "3-year reserved instance commitment for core database servers. Long-term stable workload.",
        "outcome": "success",
        "savings_percent": 58.0,
        "cost_impact": 45000
    },
    {
        "optimization_type": "right_sizing",
        "decision_context": "Increased from m5.large to m5.xlarge due to memory constraints during peak hours."
    }
]
```

```
"outcome": "success",
"savings_percent": -15.0, # Negative = cost increase
"cost_impact": -4500
}
]

print(f"Adding {len(cost_decisions)} cost decisions...")
for decision in cost_decisions:
    client.store_cost_decision(
        optimization_id=f"test-{decision['optimization_type']}-001",
        customer_id="789e0123-e89b-12d3-a456-426614174000",
        **decision
    )
print(f"✓ Added {len(cost_decisions)} decisions")

# Verify total count
info = client.client.get_collection("cost_optimization_knowledge")
print(f"✓ Total cost decisions: {info.points_count}")

print("\n==== Test data loaded! ====")
EOF
```

## ✓ COMPREHENSIVE VERIFICATION

Run this complete verification script:

```
bash
```

```
cd ~/optiinfra
```

```
python << 'EOF'
from shared.qdrant.client import get_qdrant_client
from shared.qdrant.schemas.collections import ALL_COLLECTIONS

print("=" * 70)
print("FOUNDATION-0.4 COMPREHENSIVE VERIFICATION")
print("=" * 70)

client = get_qdrant_client()

# 1. Connection Test
print("\n1. CONNECTION TEST")
if client.ping():
    print("  ✓ Qdrant connected")
else:
    print("  ✗ Qdrant connection failed")
    exit(1)

# 2. Collections Verification
print("\n2. COLLECTIONS VERIFICATION")
existing_collections = {c.name for c in client.client.get_collections().collections}

all_present = True
for config in ALL_COLLECTIONS:
    if config.name in existing_collections:
        info = client.client.get_collection(config.name)
        print(f"  ✓ {config.name} ({info.points_count} points)")
    else:
        print(f"  ✗ {config.name} missing!")
        all_present = False

if not all_present:
    print("  ✗ Some collections are missing")
    exit(1)

# 3. Storage Test
print("\n3. STORAGE TEST")
try:
    point_id = client.store_cost_decision(
        optimization_id="verify-test-001",
        customer_id="verify-customer-001",
```

```

        optimization_type="spot_migration",
        decision_context="Verification test decision",
        outcome="success",
        savings_percent=35.0,
        cost_impact=15000
    )
    print(f" ✅ Stored point: {point_id[:8]}...")

except Exception as e:
    print(f" ❌ Storage failed: {e}")
    exit(1)

# 4. Search Test
print("\n4. SEARCH TEST")
try:
    results = client.search_similar_cost_decisions(
        query="spot instance migration",
        limit=3
    )
    print(f" ✅ Search returned {len(results)} result(s)")

    if results:
        print(f" ✅ Top result score: {results[0]['score']:.3f}")

except Exception as e:
    print(f" ❌ Search failed: {e}")
    exit(1)

# 5. Collection Stats
print("\n5. COLLECTION STATISTICS")
for config in ALL_COLLECTIONS:
    info = client.client.get_collection(config.name)
    print(f" {config.name}:")
    print(f" - Points: {info.points_count}")
    print(f" - Vector size: {config.vector_size}")
    print(f" - Distance: {config.distance}")

print("\n" + "=" * 70)
print(" ✅ ALL VERIFICATION TESTS PASSED!")
print("=" * 70)
print("\nQdrant is ready for production use!")
print("Agents can now learn from past decisions and retrieve context!")
EOF

```

## Expected Output:

---

---

## FOUNDATION-0.4 COMPREHENSIVE VERIFICATION

---

---

### 1. CONNECTION TEST

- Qdrant connected

### 2. COLLECTIONS VERIFICATION

- cost\_optimization\_knowledge (6 points)
- performance\_patterns (1 points)
- customer\_context (1 points)

### 3. STORAGE TEST

- Stored point: abc12345...

### 4. SEARCH TEST

- Search returned 3 result(s)
- Top result score: 0.892

### 5. COLLECTION STATISTICS

cost\_optimization\_knowledge:

- Points: 7
- Vector size: 1536
- Distance: COSINE

performance\_patterns:

- Points: 1
- Vector size: 1536
- Distance: COSINE

customer\_context:

- Points: 1
- Vector size: 1536
- Distance: COSINE

---

---

ALL VERIFICATION TESTS PASSED!

---

---

Qdrant is ready for production use!

Agents can now learn from past decisions and retrieve context!



## TROUBLESHOOTING

## Issue 1: Qdrant Container Not Running

### Symptoms:

```
bash

curl http://localhost:6333/collections
# Connection refused
```

### Solution:

```
bash

# Check if container exists
docker ps -a | grep qdrant

# If not running, start all services
cd ~/optinfra
docker-compose up -d

# Wait for startup
sleep 5

# Test again
curl http://localhost:6333/collections
```

---

## Issue 2: Collections Not Created

### Symptoms:

```
python

client.client.get_collection("cost_optimization_knowledge")
# Collection not found
```

### Solution:

```
bash
```

```
# Re-initialize collections
python << EOF
from shared.qdrant import get_qdrant_client

client = get_qdrant_client()
client.initialize_collections()
print("✅ Collections initialized")
EOF
```

---

## Issue 3: Python Import Fails

### Symptoms:

```
python

from shared.qdrant import get_qdrant_client
# ModuleNotFoundError: No module named 'qdrant_client'
```

### Solution:

```
bash

# Install dependency
cd ~/optiinfra/shared
pip install qdrant-client==1.7.0

# Verify
python -c "import qdrant_client; print('OK')"
```

---

## Issue 4: Search Returns No Results

### Symptoms:

```
python

results = client.search_similar_cost_decisions(...)
# len(results) == 0
```

### Solution:

```
bash

# Check if collection has points
python<<EOF
from shared.qdrant import get_qdrant_client

client = get_qdrant_client()
info = client.client.get_collection("cost_optimization_knowledge")
print(f"Points in collection: {info.points_count}")

# If 0, add some test data first
if info.points_count == 0:
    print("No points found. Add data using store_cost_decision().")
EOF
```

---

## Issue 5: Embedding Function Not Working

### Symptoms:

```
python

# Search returns random/incorrect results
```

### Solution:

The default embedding function uses random vectors for testing.

For production, replace with actual embedding model:

1. Install OpenAI:

```
pip install openai
```

2. Update client.py:

```
import openai
```

```
def embed_text(text: str) -> List[float]:
```

```
    response = openai.Embedding.create(
```

```
        model="text-embedding-ada-002",
```

```
        input=text
```

```
)
```

```
    return response['data'][0]['embedding']
```

```
# In __init__:
```

```
self.embedding_function = embed_text
```

3. Set API key:

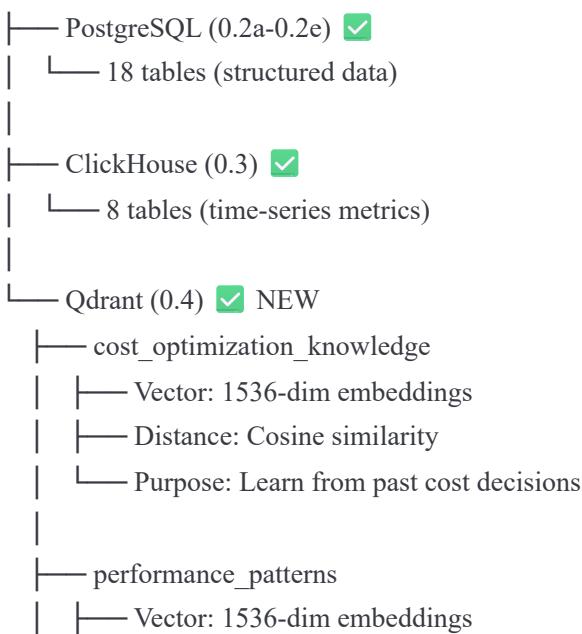
```
export OPENAI_API_KEY="your-key-here"
```



## WHAT YOU HAVE NOW

### Complete Qdrant Setup:

OptiInfra Data Architecture:



```
|   └── Distance: Cosine similarity
|   └── Purpose: Find successful optimizations
|
└── customer_context
    ├── Vector: 1536-dim embeddings
    ├── Distance: Cosine similarity
    └── Purpose: Understand customer preferences
```

Total Databases: 3

Total Collections: 3

Total Code: ~750 lines (Python)

## Capabilities Unlocked:

### Semantic Memory

- Agents remember past decisions
- Retrieve similar scenarios
- Learn from outcomes

### Context-Aware Decisions

- "What worked for similar customers?"
- "What optimizations succeeded before?"
- "What are this customer's preferences?"

### Intelligent Search

- Similarity matching (not exact match)
- Ranked by relevance
- Filtered by criteria

---

## MILESTONE ACHIEVED

### FOUNDATION-0.4 COMPLETE!

#### You now have:

-  3 Qdrant collections operational
-  Python client with semantic search

- Vector storage for agent memory
- Similarity search for past decisions
- Context retrieval for intelligent actions
- Learning from optimization outcomes
- Comprehensive documentation

## Foundation Phase Progress:

Week 1 Progress: 7/15 prompts (47%)

- └── 0.2a: Core Schema
- └── 0.2b: Agent State
- └── 0.2c: Workflow History
- └── 0.2d: Resource Schema
- └── 0.2e: Application Schema
- └── 0.3: ClickHouse
- └── 0.4: Qdrant  JUST COMPLETED!

Remaining (8 prompts):

- └── 0.5: Orchestrator (duplicate of P-02, skip)
- └── 0.6: Agent Registry
- └── 0.7