

# FOUNDATION-0.6: Agent Registry - PART 1 (Code)

## 🎯 CONTEXT

**Phase:** FOUNDATION (Week 1 - Day 3 Morning)

**Component:** Orchestrator Agent Registry (Go)

**Estimated Time:** 15 min AI execution + 10 min verification

**Complexity:** MEDIUM

**Risk Level:** LOW

**Files:** Part 1 of 2 (Code implementation)

**MILESTONE:** Agent registration and discovery system! 🎉

---

## 📦 DEPENDENCIES

### Must Complete First:

- **P-02:** Orchestrator Skeleton (Go)  COMPLETED
- **FOUNDATION-0.2a-0.2e:** PostgreSQL schemas  COMPLETED
- **FOUNDATION-0.3:** ClickHouse  COMPLETED
- **FOUNDATION-0.4:** Qdrant  COMPLETED

### Required Services Running:

```
bash

# Verify all services are healthy
cd ~/optiinfra
make verify

# Expected output:
# PostgreSQL... ✓ HEALTHY
# ClickHouse... ✓ HEALTHY
# Qdrant... ✓ HEALTHY
# Redis... ✓ HEALTHY
```

## 🎯 OBJECTIVE

Build **Agent Registry** system that enables:

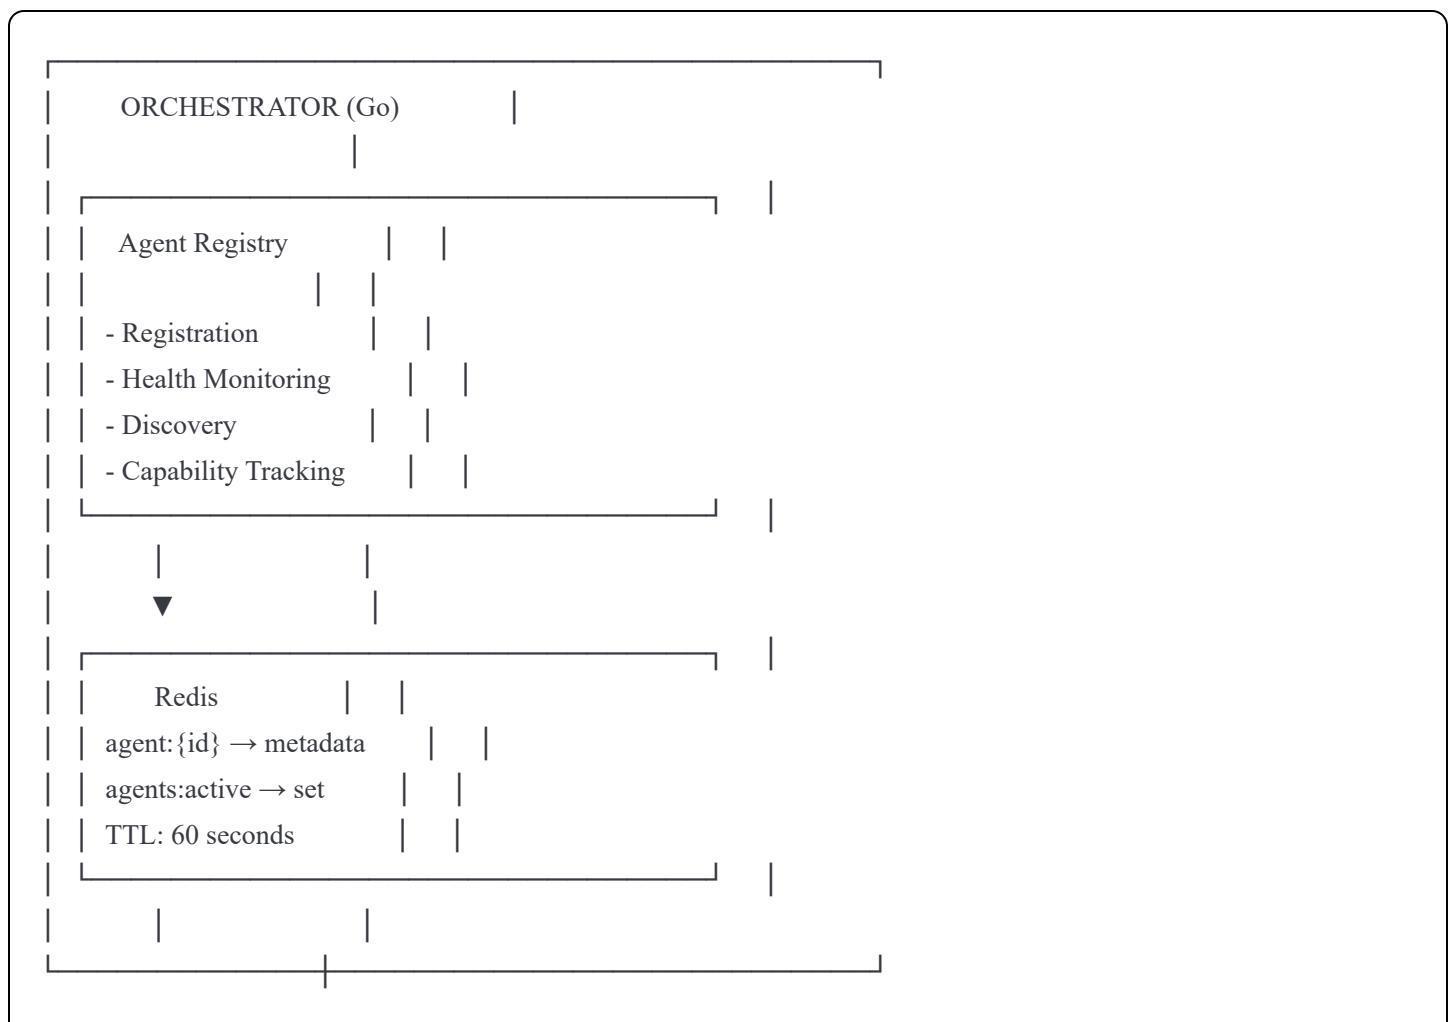
- Agents register themselves on startup
- Health monitoring and heartbeats
- Capability discovery
- Agent status tracking
- Automatic cleanup of dead agents

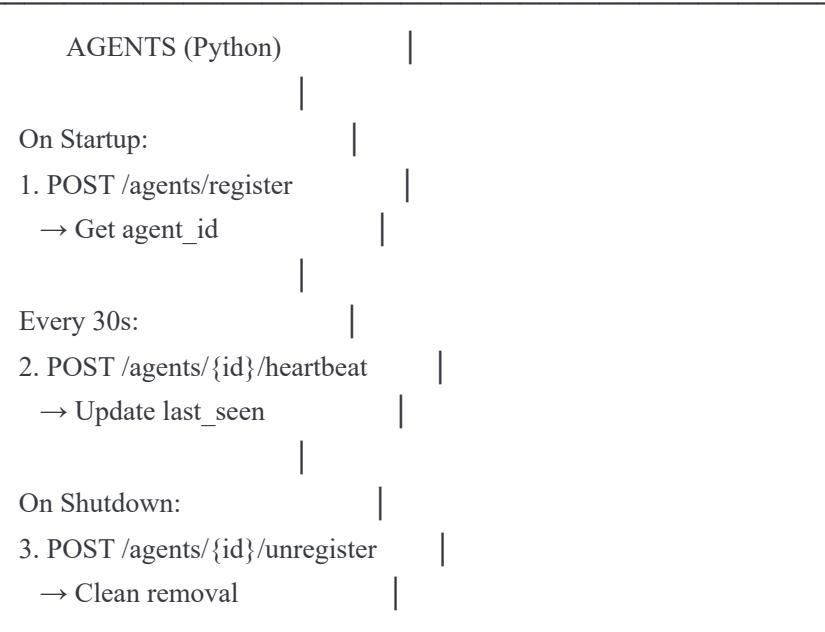
## What We're Building:

### Agent Registry Components:

- 1. Registration API** - POST /agents/register
- 2. Heartbeat API** - POST /agents/{id}/heartbeat
- 3. Discovery API** - GET /agents, GET /agents/{id}
- 4. Health Monitor** - Background goroutine checking agent health
- 5. Redis Storage** - Fast agent lookup with TTL

### Architecture:





## Use Cases:

### Agent Startup:

```

go

// Agent sends:
POST /agents/register
{
  "name": "cost-agent-1",
  "type": "cost",
  "capabilities": ["spot_migration", "reserved_instances"],
  "host": "cost-agent:8001",
  "status": "healthy"
}

// Orchestrator responds:
{
  "agent_id": "550e8400-e29b-41d4-a716-446655440000",
  "registered_at": "2025-10-20T10:30:00Z"
}
  
```

### Request Routing (Future 0.7):

```

go
  
```

```
// Orchestrator needs a cost agent
agents := registry.GetAgentsByType("cost")
// Returns: [cost-agent-1, cost-agent-2]

// Pick healthy agent with capability
agent := registry.GetAgentWithCapability("cost", "spot_migration")
// Returns: cost-agent-1
```

## FILE 1: Agent Registry Models

**Location:** `~/optiinfra/services/orchestrator/internal/registry/models.go`

```
go
```

```
package registry

import (
    "time"
)

// AgentType represents the type of agent
type AgentType string

const (
    AgentTypeCost      AgentType = "cost"
    AgentTypePerformance AgentType = "performance"
    AgentTypeResource   AgentType = "resource"
    AgentTypeApplication AgentType = "application"
)

// AgentStatus represents the current status of an agent
type AgentStatus string

const (
    AgentStatusHealthy    AgentStatus = "healthy"
    AgentStatusDegraded   AgentStatus = "degraded"
    AgentStatusUnhealthy  AgentStatus = "unhealthy"
    AgentStatusUnreachable AgentStatus = "unreachable"
)

// Agent represents a registered agent
type Agent struct {
    ID      string      `json:"id"`
    Name    string      `json:"name"`
    Type    AgentType   `json:"type"`
    Host    string      `json:"host"`
    Port    int         `json:"port"`
    Capabilities []string `json:"capabilities"`
    Status   AgentStatus `json:"status"`
    Version  string      `json:"version"`
    RegisteredAt time.Time `json:"registered_at"`
    LastSeen   time.Time   `json:"last_seen"`
    Metadata  map[string]interface{} `json:"metadata,omitempty"`
}

// RegistrationRequest is sent by agents to register
type RegistrationRequest struct {
```

```

Name      string      `json:"name" binding:"required"`
Type     AgentType    `json:"type" binding:"required"`
Host      string      `json:"host" binding:"required"`
Port      int         `json:"port" binding:"required"`
Capabilities []string   `json:"capabilities"`
Version    string      `json:"version"`
Metadata   map[string]interface{} `json:"metadata,omitempty"`
}

// RegistrationResponse is returned after successful registration
type RegistrationResponse struct {
AgentID    string    `json:"agent_id"`
RegisteredAt time.Time `json:"registered_at"`
HeartbeatURL string   `json:"heartbeat_url"`
Interval    int       `json:"heartbeat_interval_seconds"`
}

// HeartbeatRequest is sent by agents periodically
type HeartbeatRequest struct {
Status  AgentStatus   `json:"status"`
Metadata map[string]interface{} `json:"metadata,omitempty"`
}

// HeartbeatResponse confirms heartbeat received
type HeartbeatResponse struct {
Received  bool     `json:"received"`
NextInterval int     `json:"next_interval_seconds"`
Timestamp  time.Time `json:"timestamp"`
}

// AgentListResponse returns list of agents
type AgentListResponse struct {
Agents []Agent `json:"agents"`
Count  int     `json:"count"`
}

```

## FILE 2: Agent Registry Core Logic

**Location:** `~/optiinfra/services/orchestrator/internal/registry/registry.go`

go

```
package registry

import (
    "context"
    "encoding/json"
    "fmt"
    "log"
    "sync"
    "time"

    "github.com/go-redis/redis/v8"
    "github.com/google/uuid"
)

const (
    // Redis keys
    agentKeyPrefix = "agent:"
    activeAgentsSetKey = "agents:active"

    // TTL for agent entries in Redis
    agentTTL = 60 * time.Second

    // Health check interval
    healthCheckInterval = 30 * time.Second

    // Heartbeat timeout (if no heartbeat for this long, mark unhealthy)
    heartbeatTimeout = 45 * time.Second
)

// Registry manages agent registration and discovery
type Registry struct {
    redis *redis.Client
    ctx   context.Context
    mu    sync.RWMutex
    stopCh chan struct{}
}

// NewRegistry creates a new agent registry
func NewRegistry(redisClient *redis.Client) *Registry {
    return &Registry{
        redis: redisClient,
        ctx:   context.Background(),
        stopCh: make(chan struct{}),
    }
}
```

```
}

// Start begins the health monitoring goroutine
func (r *Registry) Start() {
    go r.healthMonitor()
    log.Println("Agent registry started")
}

// Stop stops the health monitoring
func (r *Registry) Stop() {
    close(r.stopCh)
    log.Println("Agent registry stopped")
}

// Register registers a new agent
func (r *Registry) Register(req *RegistrationRequest) (*RegistrationResponse, error) {
    r.mu.Lock()
    defer r.mu.Unlock()

    // Generate agent ID
    agentID := uuid.New().String()

    // Create agent
    agent := &Agent{
        ID:        agentID,
        Name:      req.Name,
        Type:      req.Type,
        Host:      req.Host,
        Port:      req.Port,
        Capabilities: req.Capabilities,
        Status:    AgentStatusHealthy,
        Version:   req.Version,
        RegisteredAt: time.Now(),
        LastSeen:  time.Now(),
        Metadata:  req.Metadata,
    }

    // Store in Redis
    if err := r.storeAgent(agent); err != nil {
        return nil, fmt.Errorf("failed to store agent: %w", err)
    }

    // Add to active agents set
}
```

```

if err := r.redis.SAdd(r.ctx, activeAgentsSetKey, agentID).Err(); err != nil {
    return nil, fmt.Errorf("failed to add to active set: %w", err)
}

log.Printf("Agent registered: %s (%s) - %s", agent.Name, agent.Type, agent.ID)

return &RegistrationResponse{
    AgentID:    agentID,
    RegisteredAt: agent.RegisteredAt,
    HeartbeatURL: fmt.Sprintf("/agents/%s/heartbeat", agentID),
    Interval:   30, // heartbeat every 30 seconds
}, nil
}

// Heartbeat updates agent's last seen time and status
func (r *Registry) Heartbeat(agentID string, req *HeartbeatRequest) (*HeartbeatResponse, error) {
    r.mu.Lock()
    defer r.mu.Unlock()

    // Get existing agent
    agent, err := r.getAgent(agentID)
    if err != nil {
        return nil, fmt.Errorf("agent not found: %w", err)
    }

    // Update status and last seen
    agent.LastSeen = time.Now()
    agent.Status = req.Status

    // Merge metadata
    if req.Metadata != nil {
        if agent.Metadata == nil {
            agent.Metadata = make(map[string]interface{})
        }
        for k, v := range req.Metadata {
            agent.Metadata[k] = v
        }
    }

    // Store updated agent
    if err := r.storeAgent(agent); err != nil {
        return nil, fmt.Errorf("failed to update agent: %w", err)
    }
}

```

```

return &HeartbeatResponse{
    Received: true,
    NextInterval: 30,
    Timestamp: time.Now(),
}, nil
}

// Unregister removes an agent from the registry
func (r *Registry) Unregister(agentID string) error {
    r.mu.Lock()
    defer r.mu.Unlock()

    // Remove from active set
    if err := r.redis.SRem(r.ctx, activeAgentsSetKey, agentID).Err(); err != nil {
        return fmt.Errorf("failed to remove from active set: %w", err)
    }

    // Delete agent key
    if err := r.redis.Del(r.ctx, agentKey(agentID)).Err(); err != nil {
        return fmt.Errorf("failed to delete agent: %w", err)
    }

    log.Printf("Agent unregistered: %s", agentID)
    return nil
}

// GetAgent retrieves a specific agent by ID
func (r *Registry) GetAgent(agentID string) (*Agent, error) {
    r.mu.RLock()
    defer r.mu.RUnlock()

    return r.getAgent(agentID)
}

// GetAllAgents retrieves all registered agents
func (r *Registry) GetAllAgents() ([]*Agent, error) {
    r.mu.RLock()
    defer r.mu.RUnlock()

    // Get all active agent IDs
    agentIDs, err := r.redis.SMembers(r.ctx, activeAgentsSetKey).Result()
    if err != nil {
        return nil, fmt.Errorf("failed to get active agents: %w", err)
    }
}

```

```

agents := make([]*Agent, 0, len(agentIDs))
for _, id := range agentIDs {
    agent, err := r.GetAgent(id)
    if err != nil {
        log.Printf("Warning: failed to get agent %s: %v", id, err)
        continue
    }
    agents = append(agents, agent)
}

return agents, nil
}

// GetAgentsByType retrieves all agents of a specific type
func (r *Registry) GetAgentsByType(agentType AgentType) ([]*Agent, error) {
    allAgents, err := r.GetAllAgents()
    if err != nil {
        return nil, err
    }

    filtered := make([]*Agent, 0)
    for _, agent := range allAgents {
        if agent.Type == agentType {
            filtered = append(filtered, agent)
        }
    }

    return filtered, nil
}

// GetAgentWithCapability finds an agent with specific capability
func (r *Registry) GetAgentWithCapability(agentType AgentType, capability string) (*Agent, error) {
    agents, err := r.GetAgentsByType(agentType)
    if err != nil {
        return nil, err
    }

    // Filter by capability and status
    for _, agent := range agents {
        if agent.Status == AgentStatusHealthy {
            for _, cap := range agent.Capabilities {
                if cap == capability {
                    return agent, nil
                }
            }
        }
    }
}

```

```

        }
    }
}

return nil, fmt.Errorf("no healthy %s agent found with capability: %os", agentType, capability)
}

// GetHealthyAgents returns only healthy agents
func (r *Registry) GetHealthyAgents() ([]*Agent, error) {
    allAgents, err := r.GetAllAgents()
    if err != nil {
        return nil, err
    }

    healthy := make([]*Agent, 0)
    for _, agent := range allAgents {
        if agent.Status == AgentStatusHealthy {
            healthy = append(healthy, agent)
        }
    }

    return healthy, nil
}

// =====
// INTERNAL HELPERS
// =====

func (r *Registry) storeAgent(agent *Agent) error {
    data, err := json.Marshal(agent)
    if err != nil {
        return fmt.Errorf("failed to marshal agent: %w", err)
    }

    // Store with TTL
    if err := r.redis.Set(r.ctx, agentKey(agent.ID), data, agentTTL).Err(); err != nil {
        return fmt.Errorf("failed to store in redis: %w", err)
    }

    return nil
}

func (r *Registry) getAgent(agentID string) (*Agent, error) {

```

```

data, err := r.redis.Get(r.ctx, agentKey(agentID)).Result()
if err == redis.Nil {
    return nil, fmt.Errorf("agent not found")
} else if err != nil {
    return nil, fmt.Errorf("failed to get from redis: %w", err)
}

var agent Agent
if err := json.Unmarshal([]byte(data), &agent); err != nil {
    return nil, fmt.Errorf("failed to unmarshal agent: %w", err)
}

return &agent, nil
}

func agentKey(agentID string) string {
    return agentKeyPrefix + agentID
}

// =====
// HEALTH MONITORING
// =====

func (r *Registry) healthMonitor() {
    ticker := time.NewTicker(healthCheckInterval)
    defer ticker.Stop()

    for {
        select {
        case <-ticker.C:
            r.checkAgentHealth()
        case <-r.stopCh:
            return
        }
    }
}

func (r *Registry) checkAgentHealth() {
    r.mu.Lock()
    defer r.mu.Unlock()

    agents, err := r.GetAllAgents()
    if err != nil {
        log.Printf("Error checking agent health: %v", err)
    }
}

```

```

    return
}

now := time.Now()
for _, agent := range agents {
    timeSinceLastSeen := now.Sub(agent.LastSeen)

    // Mark unhealthy if no heartbeat for too long
    if timeSinceLastSeen > heartbeatTimeout {
        if agent.Status != AgentStatusUnreachable {
            log.Printf("Agent %s (%s) is unreachable - last seen %v ago",
                agent.Name, agent.ID, timeSinceLastSeen)
            agent.Status = AgentStatusUnreachable
            if err := r.storeAgent(agent); err != nil {
                log.Printf("Failed to update agent status: %v", err)
            }
        }
    }
}

```

---

## FILE 3: HTTP Handlers

**Location:** `~/optiinfra/services/orchestrator/internal/registry/handlers.go`

```
go
```

```
package registry

import (
    "net/http"
    "github.com/gin-gonic/gin"
)

// Handler provides HTTP handlers for the registry
type Handler struct {
    registry *Registry
}

// NewHandler creates a new handler
func NewHandler(registry *Registry) *Handler {
    return &Handler{
        registry: registry,
    }
}

// RegisterRoutes registers all registry routes
func (h *Handler) RegisterRoutes(router *gin.Engine) {
    agents := router.Group("/agents")
    {
        agents.POST("/register", h.Register)
        agents.POST("/:id/heartbeat", h.Heartbeat)
        agents.POST("/:id/unregister", h.Unregister)
        agents.GET("", h.List)
        agents.GET("/:id", h.Get)
        agents.GET("/type/:type", h.ListByType)
    }
}

// Register handles agent registration
func (h *Handler) Register(c *gin.Context) {
    var req RegistrationRequest
    if err := c.ShouldBindJSON(&req); err != nil {
        c.JSON(http.StatusBadRequest, gin.H{"error": err.Error()})
        return
    }

    resp, err := h.registry.Register(&req)
    if err != nil {
```

```

c.JSON(http.StatusInternalServerError, gin.H{"error": err.Error()})
return
}

c.JSON(http.StatusCreated, resp)
}

// Heartbeat handles agent heartbeat
func (h *Handler) Heartbeat(c *gin.Context) {
agentID := c.Param("id")

var req HeartbeatRequest
if err := c.ShouldBindJSON(&req); err != nil {
    c.JSON(http.StatusBadRequest, gin.H{"error": err.Error()})
    return
}

resp, err := h.registry.Heartbeat(agentID, &req)
if err != nil {
    c.JSON(http.StatusNotFound, gin.H{"error": err.Error()})
    return
}

c.JSON(http.StatusOK, resp)
}

// Unregister handles agent unregistration
func (h *Handler) Unregister(c *gin.Context) {
agentID := c.Param("id")

if err := h.registry.Unregister(agentID); err != nil {
    c.JSON(http.StatusInternalServerError, gin.H{"error": err.Error()})
    return
}

c.JSON(http.StatusOK, gin.H{"message": "Agent unregistered successfully"})
}

// List returns all registered agents
func (h *Handler) List(c *gin.Context) {
agents, err := h.registry.GetAllAgents()
if err != nil {
    c.JSON(http.StatusInternalServerError, gin.H{"error": err.Error()})
    return
}

```

```

    }

c.JSON(http.StatusOK, AgentListResponse{
    Agents: convertToAgentSlice(agents),
    Count: len(agents),
})
}

// Get returns a specific agent
func (h *Handler) Get(c *gin.Context) {
    agentID := c.Param("id")

    agent, err := h.registry.GetAgent(agentID)
    if err != nil {
        c.JSON(http.StatusNotFound, gin.H{"error": "Agent not found"})
        return
    }

    c.JSON(http.StatusOK, agent)
}

// ListByType returns agents of a specific type
func (h *Handler) ListByType(c *gin.Context) {
    agentType := AgentType(c.Param("type"))

    agents, err := h.registry.GetAgentsByType(agentType)
    if err != nil {
        c.JSON(http.StatusInternalServerError, gin.H{"error": err.Error()})
        return
    }

    c.JSON(http.StatusOK, AgentListResponse{
        Agents: convertToAgentSlice(agents),
        Count: len(agents),
    })
}

func convertToAgentSlice(agents []*Agent) []Agent {
    result := make([]Agent, len(agents))
    for i, agent := range agents {
        result[i] = *agent
    }
}

```

```
    return result  
}
```

---

## FILE 4: Update Main Orchestrator

**Location:** `~/optiinfra/services/orchestrator/cmd/server/main.go`

```
go
```

```
package main

import (
    "context"
    "log"
    "os"
    "os/signal"
    "syscall"
    "time"

    "github.com/gin-gonic/gin"
    "github.com/go-redis/redis/v8"

    "optiinfra/services/orchestrator/internal/registry"
)

func main() {
    // Initialize Redis
    redisClient := redis.NewClient(&redis.Options{
        Addr:   getEnv("REDIS_ADDR", "localhost:6379"),
        Password: getEnv("REDIS_PASSWORD", ""),
        DB:      0,
    })

    // Test Redis connection
    ctx := context.Background()
    if err := redisClient.Ping(ctx).Err(); err != nil {
        log.Fatal("Failed to connect to Redis:", err)
    }
    log.Println("Connected to Redis")

    // Initialize Agent Registry
    agentRegistry := registry.NewRegistry(redisClient)
    agentRegistry.Start()
    defer agentRegistry.Stop()

    // Initialize Gin
    router := gin.Default()

    // Health check endpoint
    router.GET("/health", func(c *gin.Context) {
        c.JSON(200, gin.H{
            "status": "healthy",
        })
    })
}
```

```
"service": "orchestrator",
"timestamp": time.Now(),
})

})

// Register agent registry routes
registryHandler := registry.NewHandler(agentRegistry)
registryHandler.RegisterRoutes(router)

// Start server
port := getEnv("PORT", "8080")
log.Printf("Starting orchestrator on port %s", port)

// Graceful shutdown
srv := &http.Server{
    Addr:   ":" + port,
    Handler: router,
}

go func() {
    if err := srv.ListenAndServe(); err != nil && err != http.ErrServerClosed {
        log.Fatalf("Server failed: %v", err)
    }
}()

// Wait for interrupt signal
quit := make(chan os.Signal, 1)
signal.Notify(quit, syscall.SIGINT, syscall.SIGTERM)
<-quit

log.Println("Shutting down server...")

// Graceful shutdown with timeout
ctx, cancel := context.WithTimeout(context.Background(), 5*time.Second)
defer cancel()

if err := srv.Shutdown(ctx); err != nil {
    log.Fatal("Server forced to shutdown:", err)
}

log.Println("Server exited")
}

func getEnv(key, defaultValue string) string {
```

```
if value := os.Getenv(key); value != "" {
    return value
}
return defaultValue
}
```

---

## 📁 FILE 5: Go Module Dependencies

**Location:** `~/optiinfra/services/orchestrator/go.mod`

```
go

module optiinfra/services/orchestrator

go 1.21

require (
    github.com/gin-gonic/gin v1.9.1
    github.com/go-redis/redis/v8 v8.11.5
    github.com/google/uuid v1.5.0
)
```

---

## 📁 FILE 6: Python Agent Registration Helper

**Location:** `~/optiinfra/shared/orchestrator/registration.py`

```
python
```

....

Agent registration helper for Python agents.

Usage:

```
from shared.orchestrator.registration import AgentRegistration
```

```
registration = AgentRegistration(  
    agent_name="cost-agent-1",  
    agent_type="cost",  
    host="localhost",  
    port=8001,  
    capabilities=["spot_migration", "reserved_instances"]  
)
```

```
# Register on startup  
registration.register()
```

```
# Start heartbeat loop  
registration.start_heartbeat()
```

```
# Unregister on shutdown  
registration.unregister()
```

....

```
import requests  
import threading  
import time  
import logging  
from typing import List, Dict, Any, Optional
```

```
logger = logging.getLogger(__name__)
```

```
class AgentRegistration:
```

```
    """Handles agent registration with the orchestrator."""
```

```
    def __init__(  
        self,  
        agent_name: str,  
        agent_type: str,  
        host: str,  
        port: int,  
        capabilities: List[str],
```

```
orchestrator_url: str = "http://localhost:8080",
version: str = "1.0.0",
metadata: Optional[Dict[str, Any]] = None
):
    self.agent_name = agent_name
    self.agent_type = agent_type
    self.host = host
    self.port = port
    self.capabilities = capabilities
    self.orchestrator_url = orchestrator_url
    self.version = version
    self.metadata = metadata or {}

    self.agent_id: Optional[str] = None
    self.heartbeat_interval: int = 30
    self.heartbeat_thread: Optional[threading.Thread] = None
    self.stop_heartbeat = threading.Event()
```

## def register(self) -> bool:

"""

Register this agent with the orchestrator.

Returns:

bool: True if registration successful

"""

try:

```
    response = requests.post(
        f"{self.orchestrator_url}/agents/register",
        json={
            "name": self.agent_name,
            "type": self.agent_type,
            "host": self.host,
            "port": self.port,
            "capabilities": self.capabilities,
            "version": self.version,
            "metadata": self.metadata
        },
        timeout=10
    )
```

```
    if response.status_code == 201:
        data = response.json()
        self.agent_id = data["agent_id"]
        self.heartbeat_interval = data.get("heartbeat_interval_seconds", 30)
```

```
logger.info(f"Agent registered successfully: {self.agent_id}")
    return True
else:
    logger.error(f"Registration failed: {response.status_code} - {response.text}")
    return False

except Exception as e:
    logger.error(f"Registration error: {e}")
    return False

def start_heartbeat(self):
    """Start the heartbeat thread."""
    if not self.agent_id:
        logger.error("Cannot start heartbeat - agent not registered")
        return

    self.stop_heartbeat.clear()
    self.heartbeat_thread = threading.Thread(target=self._heartbeat_loop, daemon=True)
    self.heartbeat_thread.start()
    logger.info("Heartbeat started")

def stop_heartbeat_loop(self):
    """Stop the heartbeat thread."""
    self.stop_heartbeat.set()
    if self.heartbeat_thread:
        self.heartbeat_thread.join(timeout=5)
    logger.info("Heartbeat stopped")

def unregister(self):
    """Unregister this agent from the orchestrator."""
    if not self.agent_id:
        return

    try:
        # Stop heartbeat first
        self.stop_heartbeat_loop()

        # Unregister
        response = requests.post(
            f"{self.orchestrator_url}/agents/{self.agent_id}/unregister",
            timeout=10
        )

        if response.status_code == 200:
```

```
    logger.info(f"Agent unregistered successfully: {self.agent_id}")

else:
    logger.warning(f"Unregister failed: {response.status_code}")

except Exception as e:
    logger.error(f"Unregister error: {e}")

def _heartbeat_loop(self):
    """Background thread that sends periodic heartbeats."""
    while not self.stop_heartbeat.is_set():
        try:
            self._send_heartbeat()
        except Exception as e:
            logger.error(f"Heartbeat error: {e}")

        # Wait for next interval
        self.stop_heartbeat.wait(self.heartbeat_interval)

def _send_heartbeat(self):
    """Send a single heartbeat to the orchestrator."""
    if not self.agent_id:
        return

    try:
        response = requests.post(
            f"{self.orchestrator_url}/agents/{self.agent_id}/heartbeat",
            json={
                "status": "healthy",
                "metadata": self.metadata
            },
            timeout=5
        )
    
```