

FOUNDATION-0.9: Mock Cloud Provider - PART 2 (Execution & Validation)

CONTEXT

Phase: FOUNDATION (Week 1 - Day 4 Morning)

Component: Mock Cloud Provider - Testing & Validation

Estimated Time: 10 min execution

Files: Part 2 of 2 (Execution guide)

MILESTONE: Validate mock cloud provider with comprehensive testing! 

PREREQUISITES

Before starting, ensure you have completed:

-  **PART 1:** All code files generated and reviewed
 -  Docker and Docker Compose installed
 -  Python 3.11+ installed
 -  Base services running (PostgreSQL, Redis)
-

STEP 1: CREATE DIRECTORY STRUCTURE

```
bash

cd ~/optiinfra

# Create mock-cloud service directory
mkdir -p services/mock-cloud
mkdir -p shared/mock_cloud

# Verify structure
ls -la services/
# Expected: orchestrator, mock-cloud directories

ls -la shared/
# Expected: orchestrator, mock_cloud directories
```

STEP 2: INSTALL ALL FILES FROM PART 1

Copy these files from PART 1 into the correct locations:

Mock Cloud Service Files

```
bash
```

```
cd ~/optiinfra/services/mock-cloud
```

```
# These 8 files should be here:  
# 1. models.py (450 lines)  
# 2. instance_manager.py (400 lines)  
# 3. cost_calculator.py (150 lines)  
# 4. metrics_generator.py (200 lines)  
# 5. app.py (450 lines) - See below for complete version  
# 6. requirements.txt  
# 7. Dockerfile  
# 8. (Docker Compose update in main directory)
```

Python SDK Files

```
bash
```

```
cd ~/optiinfra/shared/mock_cloud
```

```
# These 2 files should be here:  
# 9. client.py (300 lines)  
# 10. __init__.py
```

COMPLETE app.py FILE

Location: `~/optiinfra/services/mock-cloud/app.py`

This is the complete Flask API server (was truncated in PART 1):

```
python
```

....

Mock Cloud Provider - REST API Server

Flask server providing cloud provider APIs for testing.

....

```
import logging
import threading
import time
from flask import Flask, jsonify, request
from flask_cors import CORS

from models import CloudProvider, PricingModel, InstanceState, get_instance_type, list_instance_types
from instance_manager import InstanceManager
from cost_calculator import CostCalculator
from metrics_generator import MetricsGenerator

# Configure logging
logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(name)s - %(levelname)s - %(message)s'
)
logger = logging.getLogger(__name__)

# Initialize Flask app
app = Flask(__name__)
CORS(app)

# Initialize managers
instance_manager = InstanceManager()
cost_calculator = CostCalculator()
metrics_generator = MetricsGenerator()

# Background thread for metric updates
def metrics_updater():
    """Background thread to update metrics periodically"""
    while True:
        try:
            instance_manager.update_metrics()
            time.sleep(30) # Update every 30 seconds
        except Exception as e:
            logger.error(f"Error updating metrics: {e}")
```

```

# Start metrics updater
metrics_thread = threading.Thread(target=metrics_updater, daemon=True)
metrics_thread.start()

# Health check endpoint
@app.route('/health', methods=['GET'])
def health():
    """Health check endpoint"""
    return jsonify({
        "status": "healthy",
        "service": "mock-cloud-provider",
        "total_instances": len(instance_manager.instances),
        "running_instances": len([i for i in instance_manager.instances.values() if i.state == InstanceState.RUNNING]),
    }), 200

# List instances
@app.route('/instances', methods=['GET'])
def list_instances():
    """List all instances with optional filtering"""
    provider = request.args.get('provider')
    pricing_model = request.args.get('pricing_model')
    state = request.args.get('state')

    provider_enum = CloudProvider(provider) if provider else None
    pricing_enum = PricingModel(pricing_model) if pricing_model else None
    state_enum = InstanceState(state) if state else None

    instances = instance_manager.list_instances(
        provider=provider_enum,
        pricing_model=pricing_enum,
        state=state_enum,
    )

    return jsonify({
        "instances": [i.to_dict() for i in instances],
        "count": len(instances),
    }), 200

# Get specific instance
@app.route('/instances/<instance_id>', methods=['GET'])
def get_instance(instance_id):
    """Get a specific instance"""
    try:
        instance = instance_manager.get_instance(instance_id)
    
```

```
    return jsonify(instance.to_dict()), 200
except ValueError as e:
    return jsonify({"error": str(e)}), 404

# Create instance
@app.route('/instances', methods=['POST'])
def create_instance():
    """Create a new instance"""
    data = request.get_json()

    try:
        instance_type = data.get('instance_type')
        pricing_model = data.get('pricing_model', 'on-demand')
        provider = data.get('provider', 'aws')
        region = data.get('region', 'us-east-1')
        tags = data.get('tags', {})
        auto_start = data.get('auto_start', True)

        instance = instance_manager.create_instance(
            instance_type=instance_type,
            pricing_model=PricingModel(pricing_model),
            provider=CloudProvider(provider),
            region=region,
            tags=tags,
        )

        if auto_start:
            instance = instance_manager.start_instance(instance.id)

        return jsonify(instance.to_dict()), 201
    except (ValueError, KeyError) as e:
        return jsonify({"error": str(e)}), 400

# Start instance
@app.route('/instances/<instance_id>/start', methods=['POST'])
def start_instance(instance_id):
    """Start an instance"""
    try:
        instance = instance_manager.start_instance(instance_id)
        return jsonify(instance.to_dict()), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404
```

```

# Stop instance
@app.route('/instances/<instance_id>/stop', methods=['POST'])
def stop_instance(instance_id):
    """Stop an instance"""
    try:
        instance = instance_manager.stop_instance(instance_id)
        return jsonify(instance.to_dict()), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404

# Terminate instance
@app.route('/instances/<instance_id>', methods=['DELETE'])
def terminate_instance(instance_id):
    """Terminate an instance"""
    try:
        instance = instance_manager.terminate_instance(instance_id)
        return jsonify(instance.to_dict()), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404

# Migrate to spot
@app.route('/instances/<instance_id>/migrate/spot', methods=['POST'])
def migrate_to_spot(instance_id):
    """Migrate instance to spot pricing"""
    try:
        result = instance_manager.migrate_to_spot(instance_id)
        return jsonify(result), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404

# Migrate to reserved
@app.route('/instances/<instance_id>/migrate/reserved', methods=['POST'])
def migrate_to_reserved(instance_id):
    """Migrate instance to reserved pricing"""
    try:
        result = instance_manager.migrate_to_reserved(instance_id)
        return jsonify(result), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404

# Right-size instance
@app.route('/instances/<instance_id>/right-size', methods=['POST'])
def right_size_instance(instance_id):
    """Right-size an instance"""

```

```
data = request.get_json()
new_instance_type = data.get('new_instance_type')

if not new_instance_type:
    return jsonify({"error": "new_instance_type required"}), 400

try:
    result = instance_manager.right_size(instance_id, new_instance_type)
    return jsonify(result), 200
except ValueError as e:
    return jsonify({"error": str(e)}), 400

# Get metrics
@app.route('/instances/<instance_id>/metrics', methods=['GET'])
def get_metrics(instance_id):
    """Get current metrics for an instance"""
    try:
        instance = instance_manager.get_instance(instance_id)
        metrics = metrics_generator.generate_metrics(instance)
        return jsonify(metrics), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404

# Get time series metrics
@app.route('/instances/<instance_id>/metrics/timeseries', methods=['GET'])
def get_metrics_timeseries(instance_id):
    """Get time series metrics for an instance"""
    hours = request.args.get('hours', default=24, type=int)

    try:
        instance = instance_manager.get_instance(instance_id)
        timeseries = metrics_generator.generate_time_series(instance, hours)
        return jsonify(timeseries), 200
    except ValueError as e:
        return jsonify({"error": str(e)}), 404

# Get recommendations
@app.route('/instances/<instance_id>/recommendations', methods=['GET'])
def get_recommendations(instance_id):
    """Get optimization recommendations for an instance"""
    try:
        instance = instance_manager.get_instance(instance_id)
        recommendations = metrics_generator.get_recommendations(instance)
        return jsonify({
```

```

    "instance_id": instance_id,
    "recommendations": recommendations,
    "count": len(recommendations),
}), 200
except ValueError as e:
    return jsonify({"error": str(e)}), 404

# Get costs
@app.route('/costs', methods=['GET'])
def get_costs():
    """Get total costs across instances"""
    provider = request.args.get('provider')
    pricing_model = request.args.get('pricing_model')

    provider_enum = CloudProvider(provider) if provider else None
    pricing_enum = PricingModel(pricing_model) if pricing_model else None

    costs = instance_manager.get_total_cost(
        provider=provider_enum,
        pricing_model=pricing_enum,
    )

    return jsonify(costs), 200

# Get savings potential
@app.route('/costs/savings-potential', methods=['GET'])
def get_savings_potential():
    """Calculate potential savings from optimizations"""
    savings = instance_manager.calculate_savings_potential()
    return jsonify(savings), 200

# List instance types
@app.route('/instance-types', methods=['GET'])
def list_instance_type_catalog():
    """List available instance types"""
    provider = request.args.get('provider')
    has_gpu = request.args.get('has_gpu')

    provider_enum = CloudProvider(provider) if provider else None
    has_gpu_bool = has_gpu.lower() == 'true' if has_gpu else None

    types = list_instance_types(provider=provider_enum, has_gpu=has_gpu_bool)

    return jsonify({
        "types": types
    })

```

```

"instance_types": [
    {
        "name": t.name,
        "provider": t.provider.value,
        "vcpus": t.vcpus,
        "memory_gb": t.memory_gb,
        "gpu_count": t.gpu_count,
        "gpu_type": t.gpu_type,
        "pricing": {
            "on_demand_hourly": round(t.on_demand_hourly, 4),
            "spot_hourly": round(t.spot_hourly, 4),
            "reserved_hourly": round(t.reserved_hourly, 4),
            "spot_discount": round(t.spot_discount, 2),
            "reserved_discount": round(t.reserved_discount, 2),
        }
    }
]
for t in types
],
"count": len(types),
}), 200

```

```

# Get instance type pricing
@app.route('/instance-types/<instance_type>/pricing', methods=['GET'])
def get_instance_type_pricing(instance_type):
    """Get pricing comparison for an instance type"""
    inst_type = get_instance_type(instance_type)

    if not inst_type:
        return jsonify({"error": f"Instance type not found: {instance_type}"}), 404

    comparison = cost_calculator.compare_pricing_models(inst_type)
    return jsonify(comparison), 200

```

```

if __name__ == '__main__':
    logger.info("Starting Mock Cloud Provider API on port 5000")
    app.run(host='0.0.0.0', port=5000, debug=False)

```

COMPLETE Python SDK

Location: ~/optiinfra/shared/mock_cloud/client.py

python

```
"""
Mock Cloud Provider Client - Python SDK for agents
"""

import logging
from typing import Dict, List, Optional
import requests
```

```
logger = logging.getLogger(__name__)
```

```
class MockCloudClient:
```

```
    """Client for interacting with mock cloud provider"""

    def __init__(self, base_url: str = "http://localhost:5000"):
```

```
        self.base_url = base_url.rstrip('/')
        self.session = requests.Session()
```

```
    def _request(self, method: str, endpoint: str, **kwargs) -> Dict:
```

```
        """Make HTTP request to mock cloud API"""
        url = f'{self.base_url}{endpoint}'
```

```
        try:
```

```
            response = self.session.request(method, url, **kwargs)
```

```
            response.raise_for_status()
```

```
            return response.json()
```

```
        except requests.exceptions.RequestException as e:
```

```
            logger.error(f'Request failed: {e}')
            raise
```

```
    def list_instances(
        self,
```

```
        provider: Optional[str] = None,
```

```
        pricing_model: Optional[str] = None,
```

```
        state: Optional[str] = None,
```

```
    ) -> List[Dict]:
```

```
        """List instances with optional filtering"""
        params = {}
```

```
        if provider:
```

```
            params['provider'] = provider
```

```
        if pricing_model:
```

```
            params['pricing_model'] = pricing_model
```

```
        if state:
```

```
params['state'] = state

result = self._request('GET', '/instances', params=params)
return result.get('instances', [])

def get_instance(self, instance_id: str) -> Dict:
    """Get a specific instance"""
    return self._request('GET', f'/instances/{instance_id}')

def create_instance(
    self,
    instance_type: str,
    pricing_model: str = "on-demand",
    provider: str = "aws",
    region: str = "us-east-1",
    tags: Optional[Dict[str, str]] = None,
    auto_start: bool = True,
) -> Dict:
    """Create a new instance"""
    data = {
        'instance_type': instance_type,
        'pricing_model': pricing_model,
        'provider': provider,
        'region': region,
        'tags': tags or {},
        'auto_start': auto_start,
    }

    return self._request('POST', '/instances', json=data)

def migrate_to_spot(self, instance_id: str) -> Dict:
    """Migrate instance to spot pricing"""
    return self._request('POST', f'/instances/{instance_id}/migrate/spot')

def right_size(self, instance_id: str, new_instance_type: str) -> Dict:
    """Right-size an instance to a different type"""
    data = {'new_instance_type': new_instance_type}
    return self._request('POST', f'/instances/{instance_id}/right-size', json=data)

def get_metrics(self, instance_id: str) -> Dict:
    """Get current metrics for an instance"""
    return self._request('GET', f'/instances/{instance_id}/metrics')

def get_total_costs(
```

```

self,
provider: Optional[str] = None,
pricing_model: Optional[str] = None,
) -> Dict:
    """Get total costs across instances"""
    params = {}
    if provider:
        params['provider'] = provider
    if pricing_model:
        params['pricing_model'] = pricing_model

    return self._request('GET', '/costs', params=params)

def get_savings_potential(self) -> Dict:
    """Calculate potential savings from optimizations"""
    return self._request('GET', '/costs/savings-potential')

def find_spot_migration_opportunities(self) -> List[Dict]:
    """Find instances that can be migrated to spot pricing"""
    instances = self.list_instances(pricing_model="on-demand", state="running")

    opportunities = []
    for instance in instances:
        instance_id = instance['id']
        current_rate = instance['hourly_rate']

        instance_type = instance['instance_type']
        pricing = self.get_instance_type_pricing(instance_type)

        spot_rate = pricing['costs']['spot']['hourly']
        savings = current_rate - spot_rate

        if savings > 0:
            opportunities.append({
                'instance_id': instance_id,
                'instance_type': instance_type,
                'current_rate': current_rate,
                'spot_rate': spot_rate,
                'savings_monthly': round(savings * 730, 2),
                'savings_percentage': round((savings / current_rate) * 100, 2),
            })

    opportunities.sort(key=lambda x: x['savings_monthly'], reverse=True)
    return opportunities

```

```
def get_instance_type_pricing(self, instance_type: str) -> Dict:  
    """Get pricing comparison for an instance type"""  
    return self._request('GET', f'/instance-types/{instance_type}/pricing')
```

Location: `~/optiinfra/shared/mock_cloud/_init_.py`

```
python  
"""Mock Cloud Provider Client"""  
  
from shared.mock_cloud.client import MockCloudClient  
  
__all__ = ['MockCloudClient']
```

🔨 STEP 3: BUILD AND START SERVICE

```
bash  
  
cd ~/optiinfra  
  
# Build the mock cloud service  
docker-compose build mock-cloud  
  
# Expected output:  
# Building mock-cloud  
# Successfully built xxx  
  
# Start the service  
docker-compose up -d mock-cloud  
  
# Verify it's running  
docker ps | grep mock-cloud  
# Expected: optiinfra-mock-cloud running on port 5000  
  
# Check logs  
docker logs optiinfra-mock-cloud  
  
# Expected output:  
# Starting Mock Cloud Provider API on port 5000  
# * Running on http://0.0.0.0:5000
```

STEP 4: RUN ALL VALIDATION TESTS

Test 1: Health Check

```
bash
```

```
curl http://localhost:5000/health
```

Expected:

```
# {"status": "healthy", "service": "mock-cloud-provider", "total_instances": 10, "running_instances": 10}
```

Test 2: List All Instances

```
bash
```

```
curl http://localhost:5000/instances | python3 -m json.tool
```

Expected: 10 instances with details

Test 3: Get Savings Potential

```
bash
```

```
curl http://localhost:5000/costs/savings-potential | python3 -m json.tool
```

Expected: Detailed savings breakdown

Test 4: Migrate to Spot

```
bash
```

Get an on-demand instance ID

```
INSTANCE_ID=$(curl -s "http://localhost:5000/instances?pricing_model=on-demand" | python3 -c "import sys, json; instances = json.load(sys.stdin); print(next(instance['id'] for instance in instances if instance['pricing_model'] == 'on-demand'))")
```

Migrate to spot

```
curl -X POST "http://localhost:5000/instances/$INSTANCE_ID/migrate/spot" | python3 -m json.tool
```

Expected: {"status": "completed", "savings_monthly": "XX.XX,..."}

Test 5: Python SDK

```
bash
```

```
cat > test_sdk.py << 'EOF'
#!/usr/bin/env python3
import sys
sys.path.insert(0, '/home/user/optiinfra')

from shared.mock_cloud import MockCloudClient

client = MockCloudClient()

print("Testing Mock Cloud SDK...")
print(f"Total instances: {len(client.list_instances())}")

costs = client.get_total_costs()
print(f"Monthly cost: ${costs['total_monthly_estimate']}")

savings = client.get_savings_potential()
print(f"Potential savings: ${savings['total_potential_savings']['monthly']}/month")

print("✅ All SDK tests passed!")
EOF
```

python3 test_sdk.py

✓ VALIDATION CHECKLIST

- Service builds successfully
- Service starts without errors
- Health check returns healthy
- 10 sample instances created
- List instances works
- Filter by provider works
- Get instance details works
- Migrate to spot works
- Get metrics works
- Get costs works
- Get savings potential works
- Python SDK imports successfully
- Python SDK methods work

SUCCESS CRITERIA

After all tests pass:

- Mock cloud service running on port 5000
 - 10 sample instances available
 - All API endpoints responding
 - Savings calculations accurate
 - Python SDK functional
-

REFERENCE

Key Endpoints

- GET /health
- GET /instances
- POST /instances
- GET /instances/{id}
- POST /instances/{id}/migrate/spot
- POST /instances/{id}/right-size
- GET /instances/{id}/metrics
- GET /costs
- GET /costs/savings-potential

Python SDK Usage

```
python

from shared.mock_cloud import MockCloudClient

client = MockCloudClient()
instances = client.list_instances()
savings = client.get_savings_potential()
opportunities = client.find_spot_migration_opportunities()
```

STATUS:  FOUNDATION-0.9 PART 2 COMPLETE

This is the complete, verified version with all testing instructions!