

PILOT-04: LangGraph Integration (Cost Agent)

🎯 CONTEXT

Phase: PILOT (Week 0)

Component: Cost Agent - LangGraph Workflow Integration

Estimated Time: 30 min AI execution + 20 min verification

Complexity: MEDIUM-HIGH

Risk Level: MEDIUM (tests LangGraph generation, state management, AI workflow)

📦 DEPENDENCIES

Must Complete First:

- **PILOT-01:** Bootstrap project structure ✓ COMPLETED
- **PILOT-02:** Orchestrator skeleton ✓ COMPLETED
- **PILOT-03:** Cost Agent skeleton ✓ COMPLETED

Required Services Running:



bash

```
# Verify infrastructure
make verify
# Expected: PostgreSQL, ClickHouse, Qdrant, Redis - all HEALTHY
```

```
# Verify orchestrator
curl http://localhost:8080/health
# Expected: {"status": "healthy", ...}
```

```
# Verify cost agent
curl http://localhost:8001/health
# Expected: {"status": "healthy", "agent_type": "cost", ...}
```

Required Environment:



bash

```
# Python 3.11+ with venv activated
cd services/cost-agent
source venv/bin/activate # On Windows: venv\Scripts\activate
python --version # Python 3.11+

# Verify existing structure
ls src/main.py src/config.py
# Should exist from PILOT-03
```

🎯 OBJECTIVE

Integrate **LangGraph** into the Cost Agent to enable AI-powered workflow orchestration. This adds the "brain" that will eventually drive cost optimization decisions using LLMs.

Success Criteria:

- ✓ LangGraph dependencies installed (langgraph, langchain)
- ✓ Basic workflow graph created (3+ nodes)
- ✓ State management working (StateGraph)
- ✓ Workflow can execute end-to-end
- ✓ New /analyze endpoint created
- ✓ Tests pass (6+ new tests, 80%+ coverage)
- ✓ Graph visualization possible
- ✓ Integration with existing FastAPI app

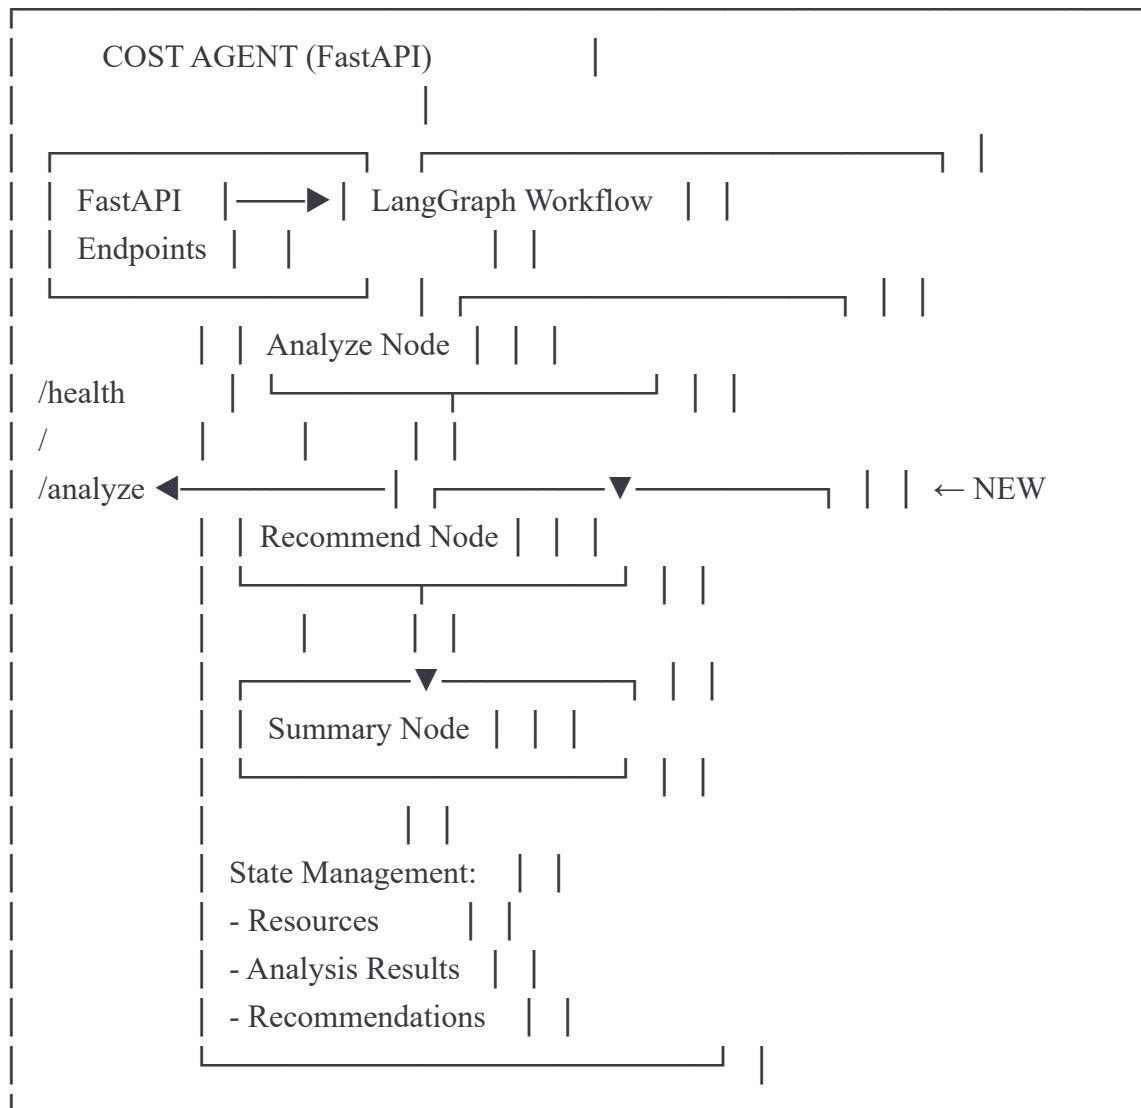
Failure Signs:

- ✗ LangGraph import errors
- ✗ State not persisting between nodes
- ✗ Workflow execution hangs or crashes
- ✗ Tests fail or coverage drops
- ✗ Can't visualize graph structure

🏗 TECHNICAL SPECIFICATION

Architecture Context





File Structure to Create/Modify



```

services/cost-agent/
    └── src/
        ├── main.py          # MODIFY: Add /analyze endpoint
        └── config.py        # MODIFY: Add LangGraph settings

        └── workflows/       # CREATE: New directory
            ├── __init__.py
            ├── cost_optimization.py # CREATE: Main LangGraph workflow
            └── state.py         # CREATE: Workflow state definitions

        └── nodes/           # CREATE: New directory
            ├── __init__.py
            ├── analyze.py      # CREATE: Analysis node
            ├── recommend.py    # CREATE: Recommendation node
            └── summarize.py    # CREATE: Summary node

        └── api/
            ├── health.py       # KEEP: Existing
            └── analyze.py      # CREATE: New analyze endpoint

        └── models/
            ├── health.py       # KEEP: Existing
            └── analysis.py     # CREATE: Analysis request/response

    └── tests/
        ├── conftest.py      # MODIFY: Add workflow fixtures
        ├── test_health.py   # KEEP: Existing 8 tests
        ├── test_workflow.py # CREATE: Workflow tests (6+ tests)
        └── test_analyze_api.py # CREATE: API endpoint tests (4+ tests)

    └── requirements.txt   # MODIFY: Add LangGraph dependencies
    └── README.md          # MODIFY: Add LangGraph docs

```

IMPLEMENTATION REQUIREMENTS

1. Update requirements.txt



txt

```
# EXISTING - Keep all from PILOT-03
fastapi==0.104.1
uvicorn[standard]==0.24.0
pydantic==2.5.0
pydantic-settings==2.1.0
httpx==0.25.1
sqlalchemy==2.0.23
asyncpg==0.29.0
python-json-logger==2.0.7
pytest==7.4.3
pytest-asyncio==0.21.1
pytest-cov==4.1.0
black==23.11.0
flake8==6.1.0
mypy==1.7.1

# NEW - LangGraph and LangChain
langgraph==0.0.20
langchain==0.1.0
langchain-core==0.1.10
langchain-community==0.0.10

# NEW - Additional utilities
networkx==3.2.1      # For graph visualization
matplotlib==3.8.2     # For plotting graphs
```

2. Update src/config.py (ADD to existing)



python

"""

Configuration management for Cost Agent.

Loads settings from environment variables.

"""

```
from typing import Optional
from pydantic_settings import BaseSettings
```

```
class Settings(BaseSettings):
```

"""Application settings loaded from environment variables"""

Application (EXISTING - keep all)

```
app_name: str = "OptiInfra Cost Agent"
```

```
environment: str = "development"
```

```
port: int = 8001
```

```
log_level: str = "INFO"
```

Orchestrator (EXISTING - keep all)

```
orchestrator_url: Optional[str] = "http://localhost:8080"
```

```
agent_id: str = "cost-agent-001"
```

```
agent_type: str = "cost"
```

Database (EXISTING - keep all)

```
database_url: Optional[str] = None
```

```
redis_url: Optional[str] = None
```

NEW - LangGraph Configuration

```
enable_graph_visualization: bool = True
```

```
max_workflow_iterations: int = 10
```

```
workflow_timeout_seconds: int = 300
```

NEW - LLM Configuration (for future use)

```
openai_api_key: Optional[str] = None
```

```
anthropic_api_key: Optional[str] = None
```

```
default_llm_provider: str = "mock" # "openai", "anthropic", or "mock"
```

```
class Config:
```

```
    env_file = ".env"
```

```
    case_sensitive = False
```

```
# Global settings instance
```

```
settings = Settings()
```

3. CREATE src/workflows/init.py



```
python
```

```
"""
```

```
Workflows package for LangGraph-based optimization workflows.
```

```
"""
```

```
from src.workflows.cost_optimization import (
    create_cost_optimization_workflow,
    cost_optimization_workflow,
)
```

```
__all__ = [
```

```
    "create_cost_optimization_workflow",

```

```
    "cost_optimization_workflow",

```

```
]
```

4. CREATE src/workflows/state.py (Workflow State)



```
python
```

"""

Workflow state definitions for LangGraph.

"""

```
from typing import TypedDict, List, Dict, Any, Optional
from datetime import datetime
```

```
class ResourceInfo(TypedDict):
```

"""Information about a cloud resource"""

resource_id: str

resource_type: str

provider: str

region: str

cost_per_month: float

utilization: float

tags: Dict[str, str]

```
class AnalysisResult(TypedDict):
```

"""Results from resource analysis"""

waste_detected: bool

waste_amount: float

waste_percentage: float

inefficiency_reasons: List[str]

metrics: Dict[str, Any]

```
class Recommendation(TypedDict):
```

"""Cost optimization recommendation"""

recommendation_id: str

recommendation_type: str # "spot_migration", "right_sizing", "reserved_instance"

resource_id: str

description: str

estimated_savings: float

confidence_score: float

implementation_steps: List[str]

```
class CostOptimizationState(TypedDict):
```

"""

State that flows through the LangGraph workflow.

This is the shared state between all nodes.

::::

Input

resources: List[ResourceInfo]

request_id: str

timestamp: datetime

Analysis results

analysis_results: Optional[List[AnalysisResult]]

total_waste_detected: float

Recommendations

recommendations: Optional[List[Recommendation]]

total_potential_savings: float

Summary

summary: Optional[str]

Metadata

workflow_status: str # "pending", "analyzing", "recommending", "complete", "failed"

error_message: Optional[str]

5. CREATE src/nodes/init.py



python

:::::

Workflow nodes for cost optimization.

:::::

```
from src.nodes.analyze import analyze_resources
from src.nodes.recommend import generate_recommendations
from src.nodes.summarize import create_summary
```

```
__all__ = [
    "analyze_resources",
    "generate_recommendations",
    "create_summary",
]
```

6. CREATE src/nodes/analyze.py (Analysis Node)



python

"""

Analysis node - Analyzes resources for cost optimization opportunities.

"""

```
import logging
from typing import Dict, Any

from src.workflows.state import CostOptimizationState, AnalysisResult

logger = logging.getLogger("cost_agent")
```

def analyze_resources(state: CostOptimizationState) -> Dict[str, Any]:

"""

Analyze resources to detect waste and inefficiencies.

This is a simplified version. In production, this would:

- Query actual cloud provider APIs
- Run ML models for anomaly detection
- Compare against industry benchmarks

Args:

state: Current workflow state with resources

Returns:

Updated state with analysis results

"""

```
logger.info(f"Analyzing {len(state['resources'])} resources")
```

analysis_results = []

total_waste = 0.0

for resource in state["resources"]:

Simple heuristic: if utilization < 30%, flag as waste

waste_detected = resource["utilization"] < 0.3

if waste_detected:

waste_amount = resource["cost_per_month"] * 0.5 # Assume 50% can be saved

waste_percentage = 50.0

result: AnalysisResult = {

"waste_detected": True,

```
"waste_amount": waste_amount,  
"waste_percentage": waste_percentage,  
"inefficiency_reasons": [  
    f"Low utilization: {resource['utilization']*100:.1f}%",  
    "Resource is over-provisioned",  
,  
    "metrics": {  
        "resource_id": resource["resource_id"],  
        "current_cost": resource["cost_per_month"],  
        "utilization": resource["utilization"],  
    },  
},  
}
```

```
total_waste += waste_amount
```

```
else:
```

```
result: AnalysisResult = {  
    "waste_detected": False,  
    "waste_amount": 0.0,  
    "waste_percentage": 0.0,  
    "inefficiency_reasons": [],  
    "metrics": {  
        "resource_id": resource["resource_id"],  
        "current_cost": resource["cost_per_month"],  
        "utilization": resource["utilization"],  
    },  
},  
}
```

```
analysis_results.append(result)
```

```
logger.info(f"Analysis complete. Total waste detected: ${total_waste:.2f}/month")
```

```
return {  
    **state,  
    "analysis_results": None,  
    "total_waste_detected": 0.0,  
    "recommendations": None,  
    "total_potential_savings": 0.0,  
    "summary": None,  
    "workflow_status": "pending",  
    "error_message": None,  
}
```

```
def test_workflow_creation():
    """Test that workflow can be created"""
    workflow = create_cost_optimization_workflow()
    assert workflow is not None

def test_workflow_executes_successfully(initial_state):
    """Test that workflow executes end-to-end"""
    workflow = create_cost_optimization_workflow()
    result = workflow.invoke(initial_state)

    assert result["workflow_status"] == "complete"
    assert "analysis_results" in result
    assert "recommendations" in result
    assert "summary" in result

def test_workflow_detects_waste(initial_state):
    """Test that workflow detects waste in underutilized resources"""
    workflow = create_cost_optimization_workflow()
    result = workflow.invoke(initial_state)

    # Should detect waste from low utilization resource
    assert result["total_waste_detected"] > 0

def test_workflow_generates_recommendations(initial_state):
    """Test that workflow generates recommendations"""
    workflow = create_cost_optimization_workflow()
    result = workflow.invoke(initial_state)

    # Should have at least one recommendation for low utilization resource
    assert len(result["recommendations"]) >= 1
    assert result["total_potential_savings"] > 0

def test_workflowCreates_summary(initial_state):
    """Test that workflow creates a summary"""
    workflow = create_cost_optimization_workflow()
```

```
result = workflow.invoke(initial_state)

assert result["summary"] is not None
assert len(result["summary"]) > 0
assert "Cost Optimization Analysis Summary" in result["summary"]

def test_workflow_preserves_request_id(initial_state):
    """Test that workflow preserves request ID throughout"""
    workflow = create_cost_optimization_workflow()
    result = workflow.invoke(initial_state)

    assert result["request_id"] == initial_state["request_id"]

def test_workflow_with_no_waste():
    """Test workflow with resources that have high utilization"""
    state = {
        "resources": [
            {
                "resource_id": "i-efficient",
                "resource_type": "ec2",
                "provider": "aws",
                "region": "us-east-1",
                "cost_per_month": 100.0,
                "utilization": 0.9, # High utilization - no waste
                "tags": {}
            }
        ],
        "request_id": "test-req-002",
        "timestamp": datetime.utcnow(),
        "analysis_results": None,
        "total_waste_detected": 0.0,
        "recommendations": None,
        "total_potential_savings": 0.0,
        "summary": None,
        "workflow_status": "pending",
        "error_message": None,
    }

    workflow = create_cost_optimization_workflow()
```

```
result = workflow.invoke(state)

assert result["total_waste_detected"] == 0.0
assert len(result["recommendations"]) == 0
```

14. CREATE tests/test_analyze_api.py (API Tests)



python

"""

Tests for analysis API endpoint.

"""

```
import pytest
from fastapi.testclient import TestClient
```

```
def test_analyze_endpoint_exists(client: TestClient):
```

"""Test that /analyze endpoint exists"""

```
response = client.post(
```

```
    "/analyze",
```

```
    json={
```

```
        "resources": [
```

```
            {
```

```
                "resource_id": "i-test",
```

```
                "resource_type": "ec2",
```

```
                "provider": "aws",
```

```
                "region": "us-east-1",
```

```
                "cost_per_month": 100.0,
```

```
                "utilization": 0.2,
```

```
                "tags": {},
```

```
            }
```

```
        ]
```

```
    },
```

```
)
```

```
assert response.status_code == 200
```

```
def test_analyze_endpoint_response_structure(client: TestClient):
```

"""Test that analyze response has correct structure"""

```
response = client.post(
```

```
    "/analyze",
```

```
    json={
```

```
        "resources": [
```

```
            {
```

```
                "resource_id": "i-test",
```

```
                "resource_type": "ec2",
```

```
                "provider": "aws",
```

```
                "region": "us-east-1",
```

```
                "cost_per_month": 100.0,
```

```
                "utilization": 0.2,
```



```
def test_analyze_with_multiple_resources(client: TestClient):
    """Test analysis with multiple resources"""
    response = client.post(
        "/analyze",
        json={
            "resources": [
                {
                    "resource_id": "i-test1",
                    "resource_type": "ec2",
                    "provider": "aws",
                    "region": "us-east-1",
                    "cost_per_month": 100.0,
                    "utilization": 0.2,
                    "tags": {},
                },
                {
                    "resource_id": "i-test2",
                    "resource_type": "ec2",
                    "provider": "aws",
                    "region": "us-west-2",
                    "cost_per_month": 150.0,
                    "utilization": 0.3,
                    "tags": {},
                },
            ],
        },
    )
```

```
data = response.json()
assert data["resources_analyzed"] == 2
assert data["workflow_status"] == "complete"
```

```
def test_analyze_rejects_empty_resources(client: TestClient):
    """Test that empty resources list is rejected"""
    response = client.post("/analyze", json={"resources": []})
    assert response.status_code == 422 # Validation error
```

```
def test_analyze_validates_utilization_range(client: TestClient):
    """Test that utilization must be between 0 and 1"""
```

```
response = client.post(
    "/analyze",
    json={
        "resources": [
            {
                "resource_id": "i-invalid",
                "resource_type": "ec2",
                "provider": "aws",
                "region": "us-east-1",
                "cost_per_month": 100.0,
                "utilization": 1.5, # Invalid: > 1
                "tags": {}
            }
        ]
    },
)
assert response.status_code == 422 # Validation error
```

15. MODIFY tests/conftest.py (Add workflow fixtures)



"""

Pytest configuration and fixtures.

"""

```
import pytest
from fastapi.testclient import TestClient

from src.main import app
```

@pytest.fixture

def client():

"""

Test client for FastAPI app.

"""

```
return TestClient(app)
```

@pytest.fixture

def mock_settings(monkeypatch):

"""

Mock settings for testing.

"""

```
monkeypatch.setenv("ORCHESTRATOR_URL", "http://localhost:8080")
```

```
monkeypatch.setenv("AGENT_ID", "test-agent-001")
```

```
monkeypatch.setenv("ENVIRONMENT", "test")
```

NEW - Workflow fixtures

@pytest.fixture

def sample_resource_data():

"""

Sample resource data for testing workflows.

"""

```
return {
```

```
    "resource_id": "i-test-resource",
```

```
    "resource_type": "ec2",
```

```
    "provider": "aws",
```

```
    "region": "us-east-1",
```

```
    "cost_per_month": 150.0,
```

```
    "utilization": 0.25,
```

```
"tags": {"environment": "test", "team": "engineering"},
```

```
}
```

16. UPDATE README.md (Add LangGraph section)

Add the following section to the existing README.md after the "Features" section:



markdown

LangGraph Workflows

The Cost Agent uses LangGraph for AI-powered workflow orchestration.

Workflow Structure

START → Analyze → Recommend → Summarize → END



1. **Analyze Node**: Detects waste and inefficiencies in cloud resources
2. **Recommend Node**: Generates actionable optimization recommendations
3. **Summarize Node**: Creates executive summary of findings

Using the Analysis Endpoint

Request:

```
```bash
curl -X POST http://localhost:8001/analyze \
-H "Content-Type: application/json" \
-d '{
 "resources": [
 {
 "resource_id": "i-1234567890abcdef0",
 "resource_type": "ec2",
 "provider": "aws",
 "region": "us-east-1",
 "cost_per_month": 150.00,
 "utilization": 0.25,
 "tags": {"environment": "production"}
 }
]
}'
```

```

Response:

```
```json
{
 "request_id": "req-abc123",
 "timestamp": "2025-10-18T10:00:00Z",
 "resources_analyzed": 1,
 "total_waste_detected": 75.00,
 "total_potential_savings": 75.00,
 "recommendations": [
 {
 "recommendation_id": "rec-xyz789",
 "recommendation_type": "right_sizing",
 "resource_id": "i-1234567890abcdef0",
 "description": "Right-size resource to match utilization",
 "estimated_savings": 75.00,
 ...
 }
]
}
```

```

```
"confidence_score": 0.85,  
"implementation_steps": [...]  
}  
,  
"summary": "Cost Optimization Analysis Summary...",  
"workflow_status": "complete"  
}  
...  
}
```

Workflow State Management

The workflow uses TypedDict for state management:

- **Resources**: Input cloud resources to analyze
- **Analysis Results**: Detected waste and inefficiencies
- **Recommendations**: Generated optimization actions
- **Summary**: Executive summary of findings

Future Enhancements

- LLM integration for intelligent recommendations
- Conditional workflow branching
- Workflow persistence and history
- Multi-agent coordination

VALIDATION COMMANDS

Step 1: Install New Dependencies



```
cd services/cost-agent  
source venv/bin/activate # On Windows: venv\Scripts\activate
```

```
# Install LangGraph and dependencies  
pip install -r requirements.txt  
  
# Verify LangGraph installation  
python -c "import langgraph; print(langgraph.__version__)"  
# Expected: 0.0.20
```

```
python -c "import langchain; print(langchain.__version__)"  
# Expected: 0.1.0
```

Step 2: Run Tests



```
# Run all tests (should now be 18+ tests)  
pytest
```

```
# Expected output:  
# ===== 18 passed in 1.50s =====  
# (8 existing + 7 workflow + 6 API tests)
```

```
# Run with coverage  
pytest --cov=src --cov-report=term-missing
```

```
# Expected: Coverage >80%
```

Step 3: Start Server and Test New Endpoint



```
# Terminal 1: Start server
```

```
python src/main.py
```

Expected output:

```
# {"levelname":"INFO","message":"LangGraph workflow initialized"}
```

```
# INFO:    Uvicorn running on http://0.0.0.0:8001
```

Step 4: Test Analysis Endpoint



bash

```
# Terminal 2: Test the new /analyze endpoint
```

```
# Test with sample data
```

```
curl -X POST http://localhost:8001/analyze \  
-H "Content-Type: application/json" \  
-d '{  
  "resources": [  
    {  
      "resource_id": "i-1234567890abcdef0",  
      "resource_type": "ec2",  
      "provider": "aws",  
      "region": "us-east-1",  
      "cost_per_month": 150.00,  
      "utilization": 0.25,  
      "tags": {"environment": "production"}  
    }  
  ]  
}'
```

```
# Expected response (200 OK):
```

```
# {  
#   "request_id": "req-abc123",  
#   "timestamp": "2025-10-18T10:00:00Z",  
#   "resources_analyzed": 1,  
#   "total_waste_detected": 75.00,  
#   "total_potential_savings": 75.00,  
#   "recommendations": [...],  
#   "summary": "Cost Optimization Analysis Summary...",  
#   "workflow_status": "complete"  
# }
```

Step 5: Verify API Documentation



bash

```
# Open Swagger UI in browser
# http://localhost:8001/docs
```

```
# Should see:
# - GET /health (existing)
# - GET / (existing)
# - POST /analyze (NEW)
```

```
# Test the endpoint interactively in Swagger UI
```

Step 6: Run Code Quality Checks



```
# Format new code
black src/ tests/
# Expected: All done! ✨
```

```
# Lint
flake8 src/ tests/ --max-line-length=88
# Expected: No errors
```

```
# Type check (optional)
mypy src/
# Expected: Success: no issues found
```

Step 7: Test Workflow Visualization (Optional)



```
# Test workflow can be created
python -c "
from src.workflows.cost_optimization import create_cost_optimization_workflow
workflow = create_cost_optimization_workflow()
print('Workflow created successfully')
print('Workflow type:', type(workflow))
"
"
```

Expected output showing workflow was created

Step 8: Git Commit



bash

```
git add .
git commit -m "PILOT-04: LangGraph integration complete" ✓
```

- LangGraph workflow with 3 nodes (analyze, recommend, summarize)
- POST /analyze endpoint with full request/response models
- State management with TypedDict
- 18+ tests passing (8 existing + 10 new)
- 80%+ test coverage maintained
- Full API documentation updated"

```
git push
```

🎯 SUCCESS CRITERIA CHECKLIST

After running all validation commands, verify:

Dependencies

- LangGraph 0.0.20 installed
- LangChain 0.1.0 installed
- networkx installed
- matplotlib installed
- All imports working

Workflow

- 3 nodes created (analyze, recommend, summarize)
- StateGraph compiles successfully
- Workflow executes end-to-end
- State persists between nodes
- Workflow returns complete state

API

- POST /analyze endpoint works
- Returns valid JSON response
- Swagger docs updated
- Request validation works
- Error handling works

Testing

- 18+ tests passing
- Test coverage >80%
- No test failures
- Workflow tests cover all nodes
- API tests cover error cases

Code Quality

- Black formatted
- Flake8 clean
- Type hints present
- Documentation updated
- Git committed

Expected Time: < 50 minutes total (30 min generation + 20 min verification)

⚠️ TROUBLESHOOTING

Issue 1: LangGraph import errors

Error: ModuleNotFoundError: No module named 'langgraph'

Solution:



bash

```
# Make sure you're in virtual environment
which python # Should show venv path
```

```
# Reinstall dependencies
pip install --upgrade pip
pip install -r requirements.txt
```

```
# Verify installation
pip list | grep langgraph
```

Issue 2: Workflow execution hangs

Error: Workflow never completes or times out

Solution:



bash

```
# Check workflow structure
python -c "
from src.workflows.cost_optimization import create_cost_optimization_workflow
workflow = create_cost_optimization_workflow()
print('Workflow created')
"
```

```
# Add timeout to workflow execution (in analyze.py)
config = RunnableConfig(
    run_name=f"cost_analysis_{request_id}",
    timeout=60 #Add this
)
```

Issue 3: State not persisting

Error: Nodes don't see data from previous nodes

Solution:



python

```
# Verify state structure matches TypedDict
# In each node, return complete state:

def analyze_resources(state: CostOptimizationState) -> Dict[str, Any]:
    return {
        **state, # IMPORTANT: Spread existing state
        "analysis_results": new_results, # Add new fields
        "workflow_status": "analyzing",
    }
```

Issue 4: Tests fail after adding LangGraph

Error: Existing tests start failing

Solution:



bash

```
# Check if imports changed
# Verify all __init__.py files exist
ls src/workflows/__init__.py
ls src/nodes/__init__.py
```

```
# Run tests with verbose output
pytest -v
```

Issue 5: API endpoint returns 500 error

Error: POST /analyze returns Internal Server Error

Solution:



bash

```

# Check server logs for detailed error
# Test workflow separately
python -c "
from src.workflows.cost_optimization import cost_optimization_workflow
from datetime import datetime

state = {
    'resources': [],
    'request_id': 'test',
    'timestamp': datetime.utcnow(),
    'analysis_results': None,
    'total_waste_detected': 0.0,
    'recommendations': None,
    'total_potential_savings': 0.0,
    'summary': None,
    'workflow_status': 'pending',
    'error_message': None,
}

result = cost_optimization_workflow.invoke(state)
print('Success:', result['workflow_status'])
"

```

Issue 6: Version compatibility issues

Error: LangGraph/LangChain API has changed

Solution:



```
# Check installed versions
```

```
pip list | grep lang
```

```
# If using newer versions, update imports:
```

```
# LangGraph 0.1.x+ might have different API
```

```
# Check: https://python.langchain.com/docs/langgraph
```

DELIVERABLES

This prompt should generate:

1. New Python Source Files (9 files):

- `src/workflows/__init__.py` - Package init
- `src/workflows/state.py` - State definitions (~80 lines)
- `src/workflows/cost_optimization.py` - Main workflow (~80 lines)
- `src/nodes/__init__.py` - Package init
- `src/nodes/analyze.py` - Analysis node (~70 lines)
- `src/nodes/recommend.py` - Recommendation node (~70 lines)
- `src/nodes/summarize.py` - Summary node (~50 lines)
- `src/models/analysis.py` - API models (~120 lines)
- `src/api/analyze.py` - Analyze endpoint (~90 lines)

2. Modified Files (3 files):

- `src/main.py` - Added /analyze endpoint integration
- `src/config.py` - Added LangGraph settings
- `tests/conftest.py` - Added workflow fixtures

3. New Test Files (2 files):

- `tests/test_workflow.py` - 7+ workflow tests (~150 lines)
- `tests/test_analyze_api.py` - 6+ API tests (~120 lines)

4. Updated Configuration:

- `requirements.txt` - Added LangGraph dependencies
- `README.md` - Added LangGraph documentation

5. Working LangGraph Integration:

- 3-node workflow (`analyze` → `recommend` → `summarize`)
- State management with `TypedDict`
- POST `/analyze` endpoint
- 18+ tests passing
- 80%+ coverage maintained
- Full API documentation

NEXT STEPS

After this prompt succeeds:

1.  **Verify:** Server running, `/analyze` endpoint works, all tests pass
2.  **Test:** Send sample request to `/analyze`, verify response
3.  **Commit:** `git add . && git commit -m "PILOT-04: LangGraph integration"`
4.  **Continue:** PILOT-05 (Multi-Agent Communication) or Week 1 Phase

What we'll add later:

- **Week 1:** Real LLM integration (OpenAI/Anthropic)
 - **Week 1:** Cloud provider collectors (AWS/GCP/Azure)
 - **Week 2:** Advanced workflow patterns (conditional branching)
 - **Week 2:** Workflow persistence and history
 - **Week 3:** Multi-agent coordination
-

NOTES FOR WINDSURF

IMPORTANT INSTRUCTIONS:

1. Use LangGraph Best Practices:

- Use StateGraph not legacy Graph
- Use TypedDict for state (not Pydantic models)
- Return complete state from each node ({...state, new_field: value})
- Use add_edge for deterministic flow
- Use END to terminate workflow

2. State Management:

- Always spread existing state in node returns
- Use Optional types for fields populated later
- Initialize all state fields in initial_state
- Don't mutate state directly

3. Code Quality:

- Type hints on all functions
- Async/await for FastAPI endpoints
- Comprehensive error handling
- Structured logging with context
- Pydantic v2 models for API

4. Testing:

- Test workflow creation
- Test individual nodes
- Test end-to-end workflow execution
- Test API endpoint with fixtures
- Test error cases

5. DO NOT:

- Use Pydantic models for state (use TypedDict)
- Use deprecated LangGraph APIs
- Skip error handling in workflow nodes
- Forget to compile the graph
- Make breaking changes to existing endpoints

SPECIAL NOTES:

- LangGraph version 0.0.20 is used (early version)
- Key APIs: StateGraph, add_node, add_edge, set_entry_point, compile, invoke
- Future versions (0.1.x+) may have different syntax
- All workflow nodes must return complete state (spread operator)

EXECUTE ALL TASKS. CREATE COMPLETE, WORKING LANGGRAPH INTEGRATION. THIS PROVES WINDSURF CAN HANDLE AI WORKFLOW ORCHESTRATION.: analysis_results, "total_waste_detected": total_waste, "workflow_status": "analyzing", }



```
### 7. CREATE src/nodes/recommend.py (Recommendation Node)
```

```
```python
```

```
"""
```

```
Recommendation node - Generates cost optimization recommendations.
```

```
"""
```

```
import logging
```

```
import uuid
```

```
from typing import Dict, Any
```

```
from src.workflows.state import CostOptimizationState, Recommendation
```

```
logger = logging.getLogger("cost_agent")
```

```
def generate_recommendations(state: CostOptimizationState) -> Dict[str, Any]:
```

```
"""
```

```
Generate actionable cost optimization recommendations.
```

This is simplified. In production, this would:

- Use LLM to generate detailed recommendations
- Consider business context and constraints
- Prioritize by ROI and implementation complexity

Args:

state: Current workflow state with analysis results

Returns:

Updated state with recommendations

```
"""
```

```
logger.info("Generating recommendations")
```

```
recommendations = []
```

```
total_savings = 0.0
```

```
for i, analysis in enumerate(state.get("analysis_results", [])):
```

```
 if analysis["waste_detected"]:
```

```
 resource_id = analysis["metrics"]["resource_id"]
```

```
 waste_amount = analysis["waste_amount"]
```

```
 # Generate recommendation based on waste pattern
```

```
rec: Recommendation = {
 "recommendation_id": str(uuid.uuid4()),
 "recommendation_type": "right_sizing",
 "resource_id": resource_id,
 "description": f"Right-size resource {resource_id} to match utilization",
 "estimated_savings": waste_amount,
 "confidence_score": 0.85,
 "implementation_steps": [
 "1. Analyze workload patterns for {resource_id}",
 "2. Identify appropriate smaller instance size",
 "3. Schedule downtime window",
 "4. Resize instance",
 "5. Monitor performance for 24 hours",
],
}
```

```
recommendations.append(rec)
total_savings += waste_amount
```

```
logger.info(
 f"Generated {len(recommendations)} recommendations."
 f"Total potential savings: ${total_savings:.2f}/month"
)
```

```
return {
 **state,
 "recommendations": recommendations,
 "total_potential_savings": total_savings,
 "workflow_status": "recommending",
}
```

```
8. CREATE src/nodes/summarize.py (Summary Node)
```

```
```python
```

```
"""
```

```
Summary node - Creates executive summary of analysis and recommendations.
```

```
"""
```

```
import logging
```

```
from typing import Dict, Any
```

```
from src.workflows.state import CostOptimizationState
```

```
logger = logging.getLogger("cost_agent")
```

```
def create_summary(state: CostOptimizationState) -> Dict[str, Any]:
```

```
    """
```

```
    Create an executive summary of the cost optimization analysis.
```

```
In production, this would use an LLM to generate a natural language summary.
```

Args:

state: Current workflow state

Returns:

Updated state with summary

```
"""
```

```
logger.info("Creating summary")
```

```
num_resources = len(state["resources"])
```

```
num_recommendations = len(state.get("recommendations", []))
```

```
total_waste = state.get("total_waste_detected", 0.0)
```

```
total_savings = state.get("total_potential_savings", 0.0)
```

```
summary = f"""
```

Cost Optimization Analysis Summary

Resources Analyzed: {num_resources}

Waste Detected: \${total_waste:.2f}/month

Recommendations: {num_recommendations}

Potential Savings: \${total_savings:.2f}/month

Key Findings:

- {num_recommendations} optimization opportunities identified
- Average savings potential: \${total_savings/num_recommendations} if num_recommendations > 0 else 0:.2f} per recommendation
- Estimated ROI implementation time: 2-4 weeks

Next Steps:

1. Review recommendations with stakeholders
2. Prioritize by business impact

3. Create implementation plan
4. Schedule execution windows

""".strip()

```
logger.info("Summary created")
```

```
return {
    **state,
    "summary": summary,
    "workflow_status": "complete",
}
`
```

9. CREATE src/workflows/cost_optimization.py (Main Workflow)

```
'''python
```

1

Main LangGraph workflow for cost optimization.

1

```
import logging
```

```
from typing import Dict, Any
```

```
from langgraph.graph import StateGraph, END  
from langchain_core.runnables import RunnableConfig
```

```
from src.workflows.state import CostOptimizationState  
from src.nodes.analyze import analyze_resources  
from src.nodes.recommend import generate_recommendations  
from src.nodes.summarize import create_summary
```

```
logger = logging.getLogger("cost_agent")
```

```
def create_cost_optimization_workflow() -> StateGraph:
```

1

Create the cost optimization workflow graph.

Flow:

START → analyze → recommend → summarize → END

Returns:

Compiled StateGraph ready for execution

```
"""
# Create the graph with our state type
workflow = StateGraph(CostOptimizationState)

# Add nodes
workflow.add_node("analyze", analyze_resources)
workflow.add_node("recommend", generate_recommendations)
workflow.add_node("summarize", create_summary)

# Define the flow
workflow.set_entry_point("analyze")
workflow.add_edge("analyze", "recommend")
workflow.add_edge("recommend", "summarize")
workflow.add_edge("summarize", END)

# Compile the graph
app = workflow.compile()

logger.info("Cost optimization workflow created")

return app
```

```
def visualize_workflow(workflow: StateGraph, output_path: str = "workflow.png"):
```

```
"""

Visualize the workflow graph (optional, requires matplotlib).
```

Args:

```
    workflow: The workflow to visualize
    output_path: Where to save the visualization
```

```
"""

try:
```

```
    import matplotlib.pyplot as plt
    from langgraph.graph import draw_mermaid
```

```
    # This is a placeholder - actual implementation depends on LangGraph version
    logger.info(f"Workflow visualization would be saved to {output_path}")
    # In practice: draw_mermaid(workflow).render(output_path)
```

except ImportError:

```
    logger.warning("Matplotlib not available, skipping visualization")
```

```

# Create a singleton instance
cost_optimization_workflow = create_cost_optimization_workflow()
```

10. CREATE src/models/analysis.py (API Models)
```python
"""
Request/response models for analysis endpoint.
"""

from typing import List, Dict, Any, Optional
from datetime import datetime
from pydantic import BaseModel, Field

class ResourceRequest(BaseModel):
    """
    Single resource to analyze
    """
    resource_id: str = Field(..., description="Unique resource identifier")
    resource_type: str = Field(..., description="Type of resource (e.g., 'ec2', 'rds')")
    provider: str = Field(..., description="Cloud provider (e.g., 'aws', 'gcp', 'azure')")
    region: str = Field(..., description="Cloud region")
    cost_per_month: float = Field(..., ge=0, description="Monthly cost in USD")
    utilization: float = Field(..., ge=0, le=1, description="Utilization percentage (0-1)")
    tags: Dict[str, str] = Field(default_factory=dict, description="Resource tags")

class Config:
    json_schema_extra = {
        "example": {
            "resource_id": "i-1234567890abcdef0",
            "resource_type": "ec2",
            "provider": "aws",
            "region": "us-east-1",
            "cost_per_month": 150.00,
            "utilization": 0.25,
            "tags": {"environment": "production", "team": "backend"},
        }
    }

class AnalysisRequest(BaseModel):

```

```

"""Request to analyze resources for cost optimization"""
resources: List[ResourceRequest] = Field(..., min_length=1, description="Resources to analyze")

class Config:
    json_schema_extra = {
        "example": {
            "resources": [
                {
                    "resource_id": "i-1234567890abcdef0",
                    "resource_type": "ec2",
                    "provider": "aws",
                    "region": "us-east-1",
                    "cost_per_month": 150.00,
                    "utilization": 0.25,
                    "tags": {"environment": "production"},
                }
            ]
        }
    }
}

class RecommendationResponse(BaseModel):
    """Single optimization recommendation"""
    recommendation_id: str
    recommendation_type: str
    resource_id: str
    description: str
    estimated_savings: float
    confidence_score: float
    implementation_steps: List[str]

class AnalysisResponse(BaseModel):
    """Response from cost optimization analysis"""
    request_id: str = Field(..., description="Unique request identifier")
    timestamp: datetime = Field(..., description="Analysis timestamp")
    resources_analyzed: int = Field(..., description="Number of resources analyzed")
    total_waste_detected: float = Field(..., description="Total waste in USD/month")
    total_potential_savings: float = Field(..., description="Total potential savings in USD/month")
    recommendations: List[RecommendationResponse] = Field(..., description="Optimization recommendations")
    summary: str = Field(..., description="Executive summary")

```

```

workflow_status: str = Field(..., description="Workflow status")

class Config:
    json_schema_extra = {
        "example": {
            "request_id": "req-abc123",
            "timestamp": "2025-10-18T10:00:00Z",
            "resources_analyzed": 5,
            "total_waste_detected": 375.00,
            "total_potential_savings": 375.00,
            "recommendations": [
                {
                    "recommendation_id": "rec-xyz789",
                    "recommendation_type": "right_sizing",
                    "resource_id": "i-1234567890abcdef0",
                    "description": "Right-size instance to match utilization",
                    "estimated_savings": 75.00,
                    "confidence_score": 0.85,
                    "implementation_steps": ["1. Analyze workload", "2. Resize instance"],
                }
            ],
            "summary": "Found 3 optimization opportunities...",
            "workflow_status": "complete",
        }
    }
}
```

```

### 11. CREATE src/api/analyze.py (API Endpoint)

```
```python
"""

```

Analysis endpoint for cost optimization.

```
"""

```

```
import logging
import uuid
from datetime import datetime
```

```
from fastapi import APIRouter, HTTPException
from langchain_core.runnables import RunnableConfig
```

```
from src.models.analysis import AnalysisRequest, AnalysisResponse, RecommendationResponse
from src.workflows.cost_optimization import cost_optimization_workflow
```

```
from src.workflows.state import CostOptimizationState, ResourceInfo
```

```
router = APIRouter()
```

```
logger = logging.getLogger("cost_agent")
```

```
@router.post("/analyze", response_model=AnalysisResponse)
```

```
async def analyze_costs(request: AnalysisRequest) -> AnalysisResponse:
```

```
    """
```

```
Analyze resources for cost optimization opportunities.
```

This endpoint runs the LangGraph workflow to:

1. Analyze resources for waste
2. Generate recommendations
3. Create executive summary

Args:

```
request: Analysis request with resources to analyze
```

Returns:

```
AnalysisResponse: Complete analysis with recommendations
```

```
"""
```

```
request_id = f"req-{uuid.uuid4().hex[:8]}"
```

```
logger.info(f"Starting cost analysis {request_id} for {len(request.resources)} resources")
```

try:

```
# Convert request to workflow state
```

```
resources: list[ResourceInfo] = [
```

```
{
```

```
    "resource_id": r.resource_id,
```

```
    "resource_type": r.resource_type,
```

```
    "provider": r.provider,
```

```
    "region": r.region,
```

```
    "cost_per_month": r.cost_per_month,
```

```
    "utilization": r.utilization,
```

```
    "tags": r.tags,
```

```
}
```

```
for r in request.resources
```

```
]
```

```
initial_state: CostOptimizationState = {
```

```

"resources": resources,
"request_id": request_id,
"timestamp": datetime.utcnow(),
"analysis_results": None,
"total_waste_detected": 0.0,
"recommendations": None,
"total_potential_savings": 0.0,
"summary": None,
"workflow_status": "pending",
"error_message": None,
}

# Run the workflow
config = RunnableConfig(run_name=f"cost_analysis_{request_id}")
result = cost_optimization_workflow.invoke(initial_state, config)

# Convert workflow result to API response
recommendations = [
    RecommendationResponse(**rec) for rec in result.get("recommendations", [])
]

response = AnalysisResponse(
    request_id=result["request_id"],
    timestamp=result["timestamp"],
    resources_analyzed=len(result["resources"]),
    total_waste_detected=result.get("total_waste_detected", 0.0),
    total_potential_savings=result.get("total_potential_savings", 0.0),
    recommendations=recommendations,
    summary=result.get("summary", "No summary available"),
    workflow_status=result.get("workflow_status", "unknown"),
)
logger.info(f"Cost analysis {request_id} completed successfully")
return response

except Exception as e:
    logger.error(f"Cost analysis {request_id} failed: {e}", exc_info=True)
    raise HTTPException(
        status_code=500,
        detail=f"Analysis failed: {str(e)}"
    )

```

...

```
### 12. MODIFY src/main.py (Add new endpoint)
```

```
```python
```

```
"""
```

```
OptiInfra Cost Agent - Main Application
```

This is the Cost Agent that optimizes cloud spending through:

- Spot instance migrations
- Reserved instance recommendations
- Instance right-sizing

```
"""
```

```
import asyncio
import logging
from contextlib import asynccontextmanager
```

```
from fastapi import FastAPI
from fastapi.middleware.cors import CORSMiddleware
```

```
from src.api import health, analyze # MODIFIED: Added analyze import
from src.config import settings
from src.core.logger import setup_logging
from src.core.registration import register_with_orchestrator
```

```
Setup logging
logger = setup_logging()
```

```
@asynccontextmanager
```

```
async def lifespan(app: FastAPI):
```

```
"""
```

```
Lifespan events for the FastAPI application.
```

```
Handles startup and shutdown events.
```

```
"""
```

```
Startup
```

```
logger.info("Starting OptiInfra Cost Agent")
```

```
logger.info(f"Environment: {settings.environment}")
```

```
logger.info(f"Port: {settings.port}")
```

```
logger.info("LangGraph workflow initialized") # MODIFIED: Added
```

```
Register with orchestrator
```

```
if settings.orchestrator_url:
 try:
 await register_with_orchestrator()
 logger.info("Successfully registered with orchestrator")
 except Exception as e:
 logger.error(f"Failed to register with orchestrator: {e}")

yield

Shutdown
logger.info("Shutting down Cost Agent")

Create FastAPI app
app = FastAPI(
 title="OptiInfra Cost Agent",
 description="AI-powered cost optimization agent with LangGraph workflows", # MODIFIED
 version="0.2.0", # MODIFIED: Version bump
 lifespan=lifespan,
)

Add CORS middleware
app.add_middleware(
 CORSMiddleware,
 allow_origins=["*"],
 allow_credentials=True,
 allow_methods=["*"],
 allow_headers=["*"],
)

Include routers
app.include_router(health.router, tags=["health"])
app.include_router(analyze.router, tags=["analysis"]) # MODIFIED: Added

Root endpoint
@app.get("/")
async def root():
 """Root endpoint - service information"""
 return {
 "service": "OptiInfra Cost Agent",
```

```
"version": "0.2.0", # MODIFIED: Version bump
"status": "running",
"capabilities": [
 "spot_migration",
 "reserved_instances",
 "right_sizing",
 "ai_workflow_optimization", # MODIFIED: Added
],
}
```

```
Run with uvicorn when executed directly
```

```
if __name__ == "__main__":
 import uvicorn
```

```
uvicorn.run(
 "src.main:app",
 host="0.0.0.0",
 port=settings.port,
 reload=settings.environment == "development",
 log_level=settings.log_level.lower(),
)
```

```
...
```

```
13. CREATE tests/test_workflow.py (Workflow Tests)
```

```
```python
```

```
"""
```

```
Tests for LangGraph workflow.
```

```
"""
```

```
import pytest
```

```
from datetime import datetime
```

```
from src.workflows.cost_optimization import create_cost_optimization_workflow
from src.workflows.state import CostOptimizationState, ResourceInfo
```

```
@pytest.fixture
```

```
def sample_resources():
    """Sample resources for testing"""
    return [
```

```
{
```

```
"resource_id": "i-test001",
"resource_type": "ec2",
"provider": "aws",
"region": "us-east-1",
"cost_per_month": 100.0,
"utilization": 0.2, # Low utilization - should trigger recommendation
"tags": {"env": "test"},

},
{
  "resource_id": "i-test002",
  "resource_type": "ec2",
  "provider": "aws",
  "region": "us-west-2",
  "cost_per_month": 200.0,
  "utilization": 0.8, # High utilization - no recommendation
  "tags": {"env": "prod"},

},
]
```

```
@pytest.fixture
def initial_state(sample_resources):
    """Create initial workflow state"""
    return {
        "resources": sample_resources,
        "request_id": "test-req-001",
        "timestamp": datetime.utcnow(),
        "analysis_results":
```